

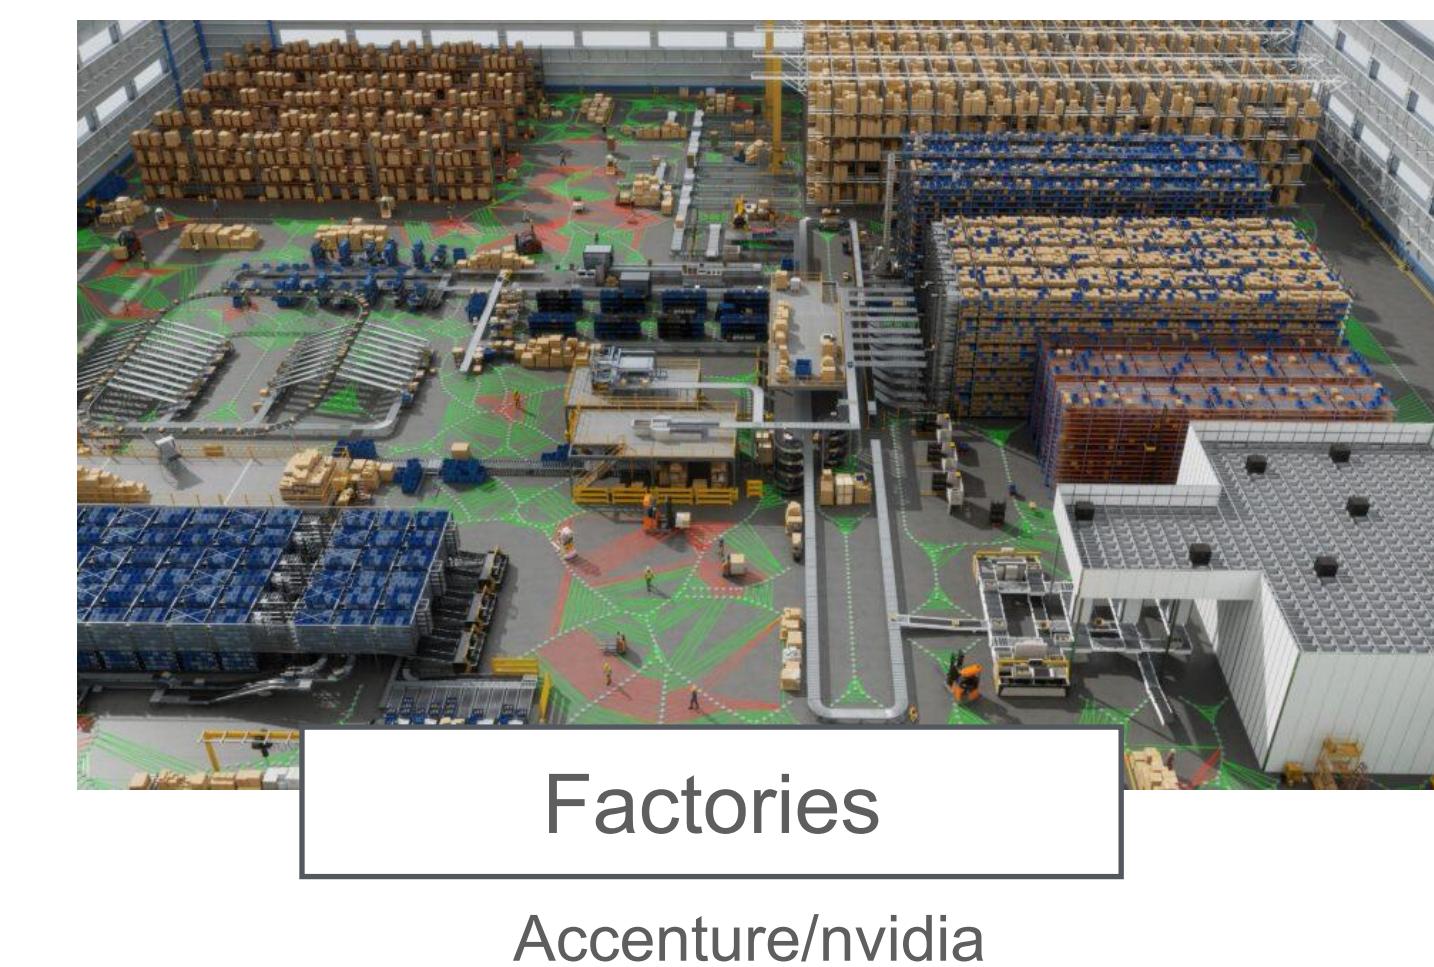
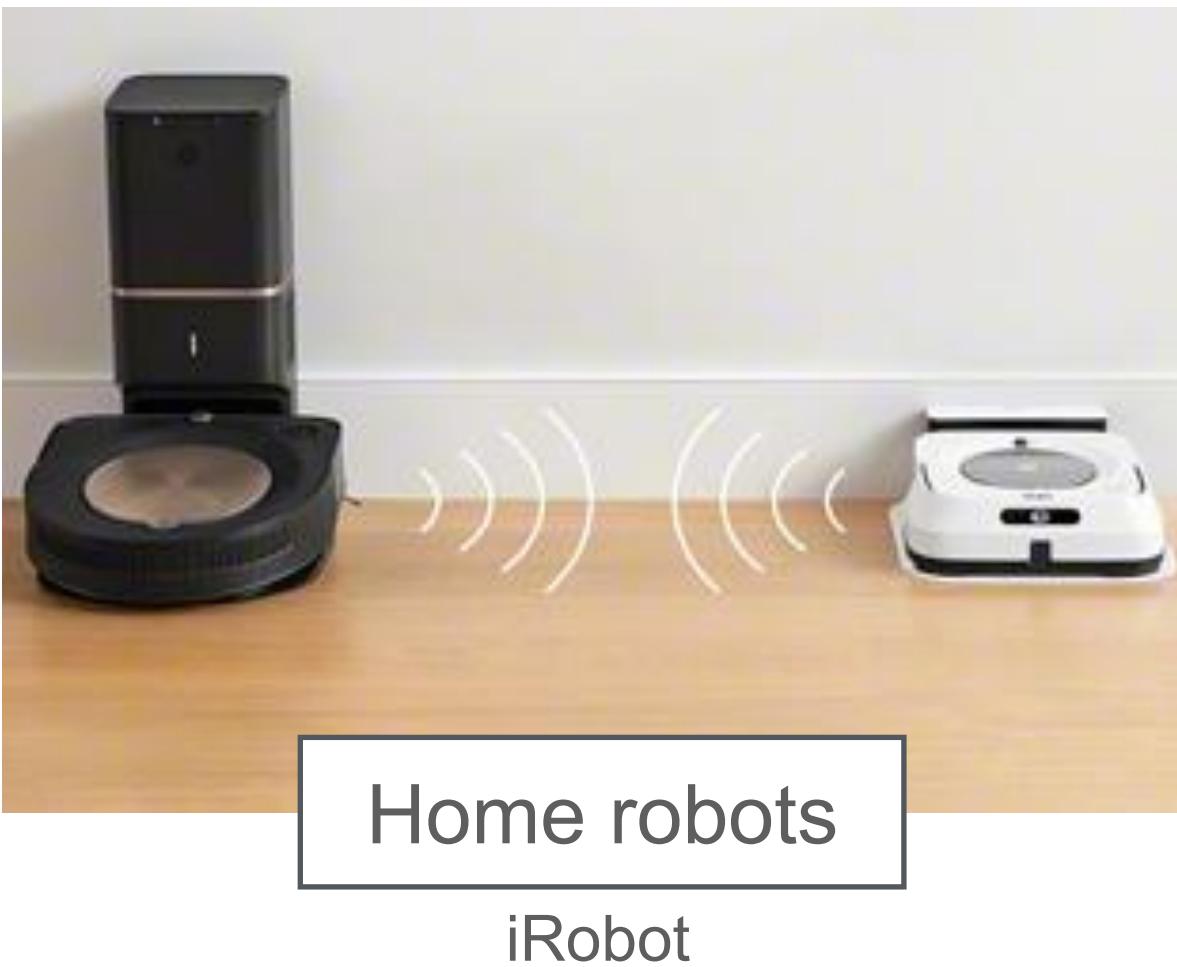
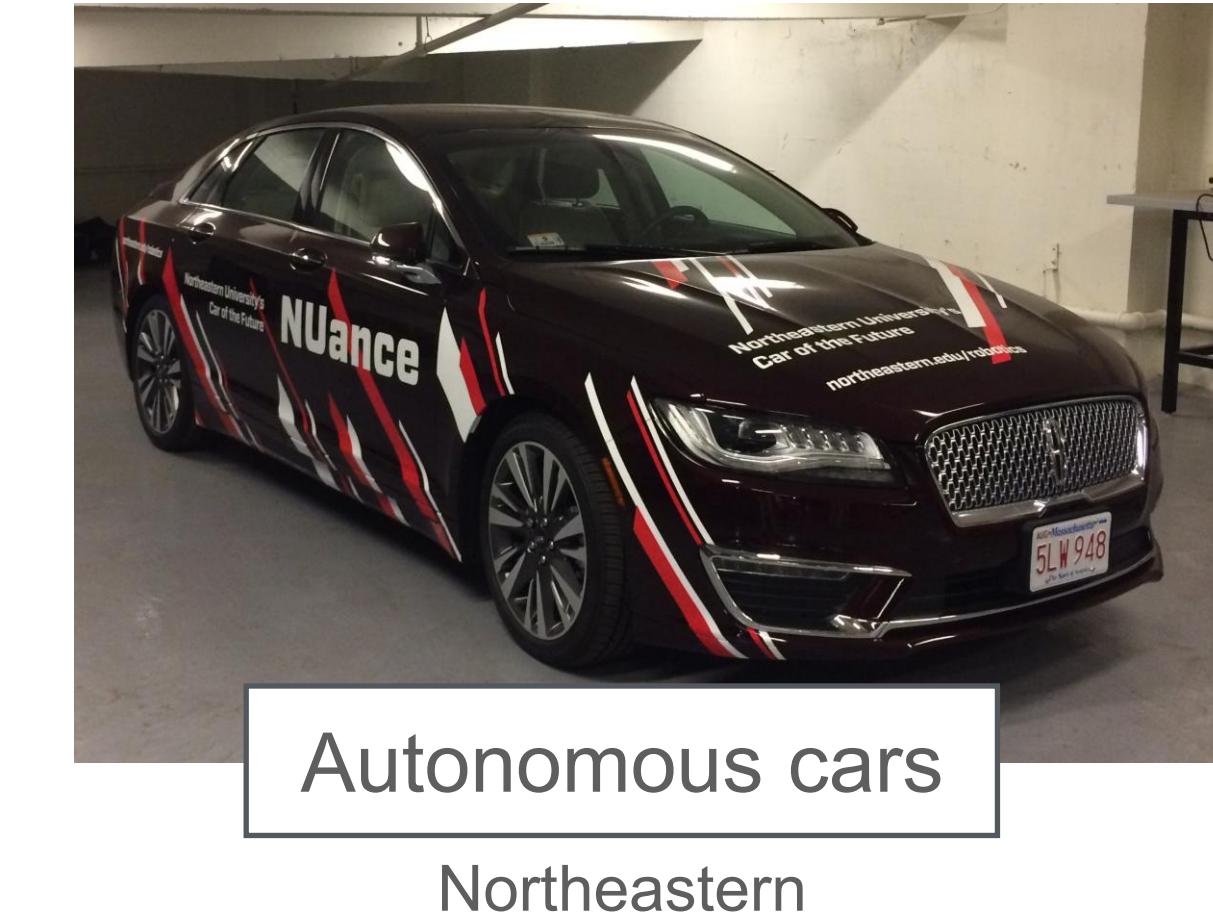


Northeastern University
Khoury College of
Computer Sciences

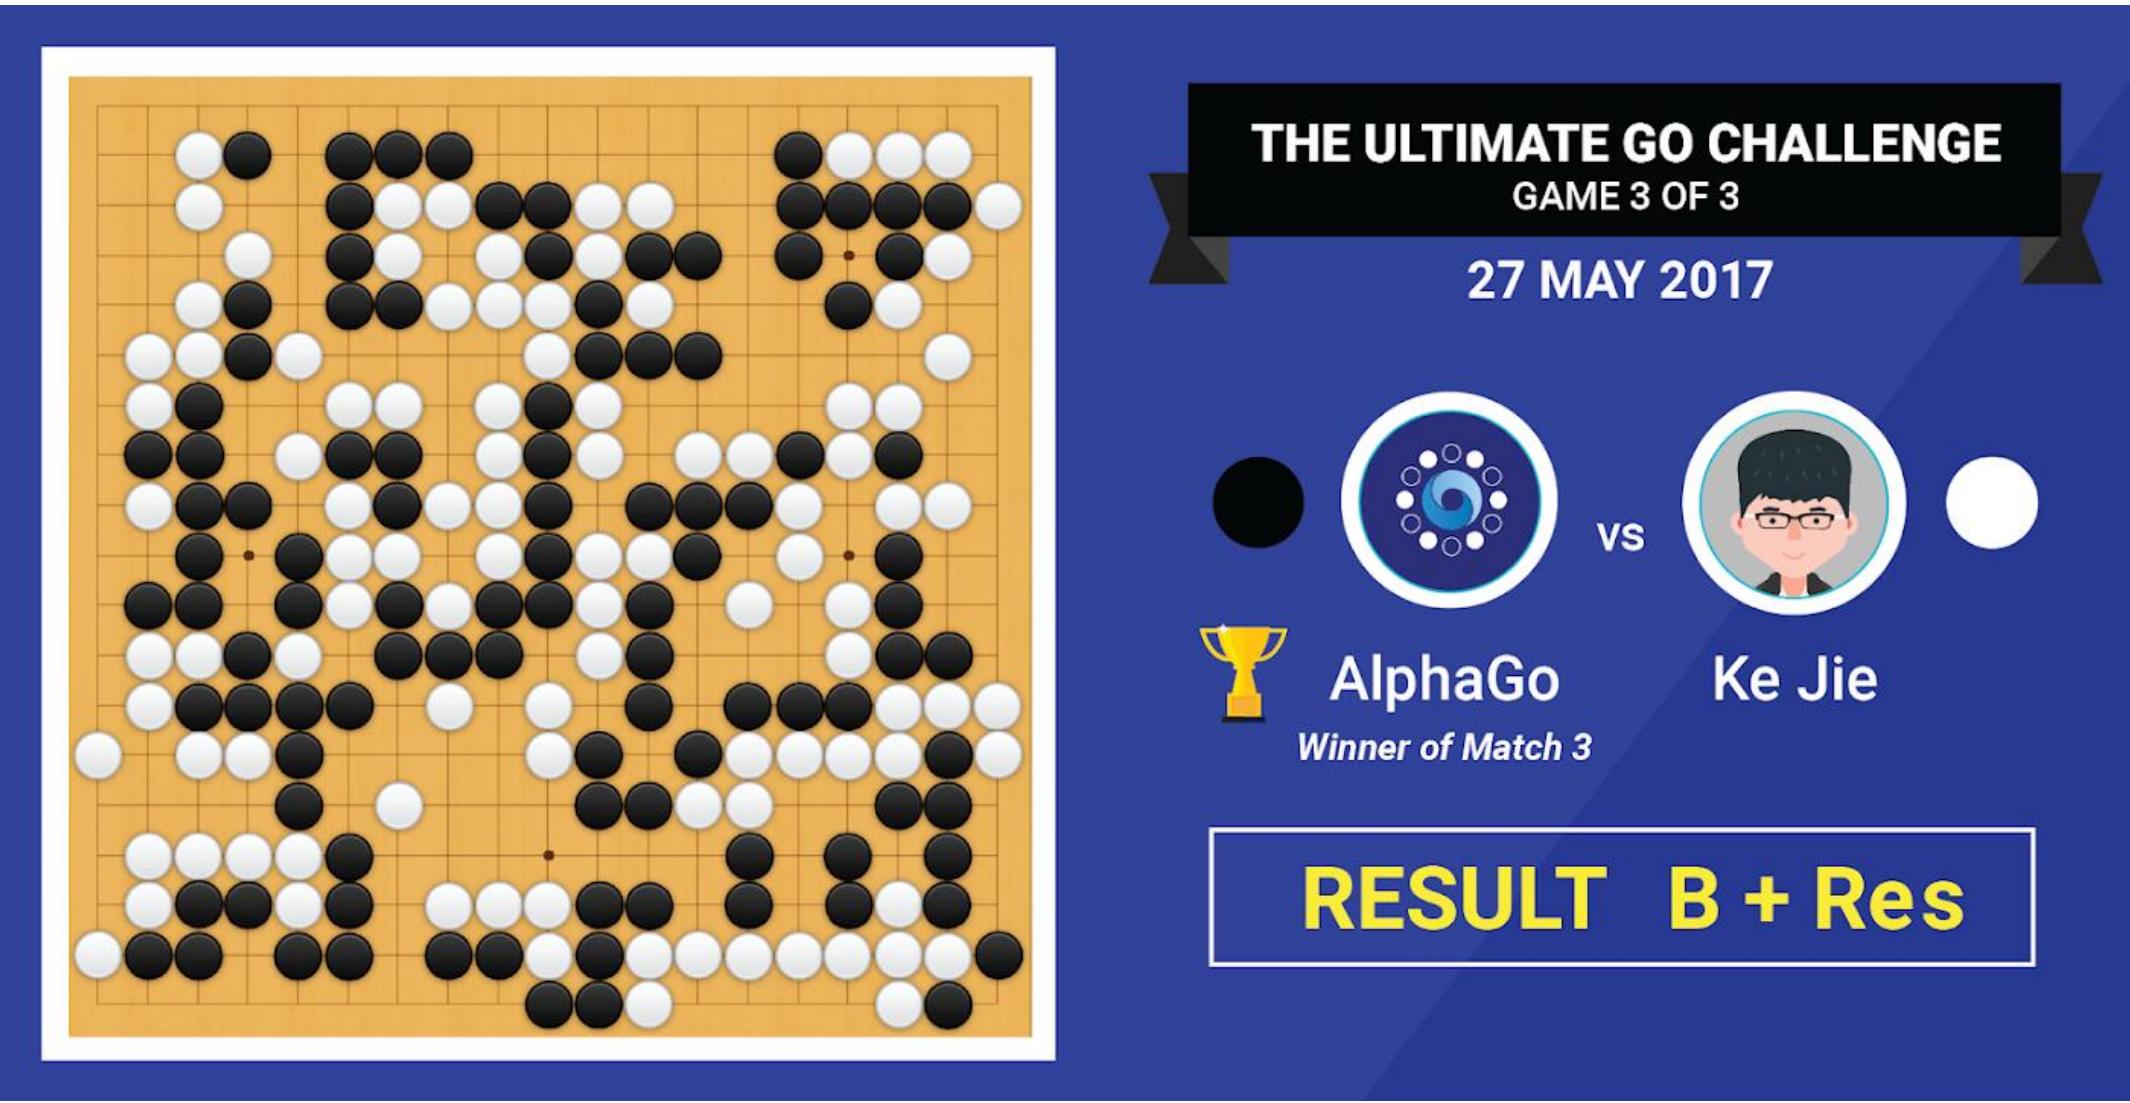
A Short Introduction to Cooperative Multi-Agent Reinforcement Learning

Chris Amato

Multi-agent systems are (going to be) everywhere



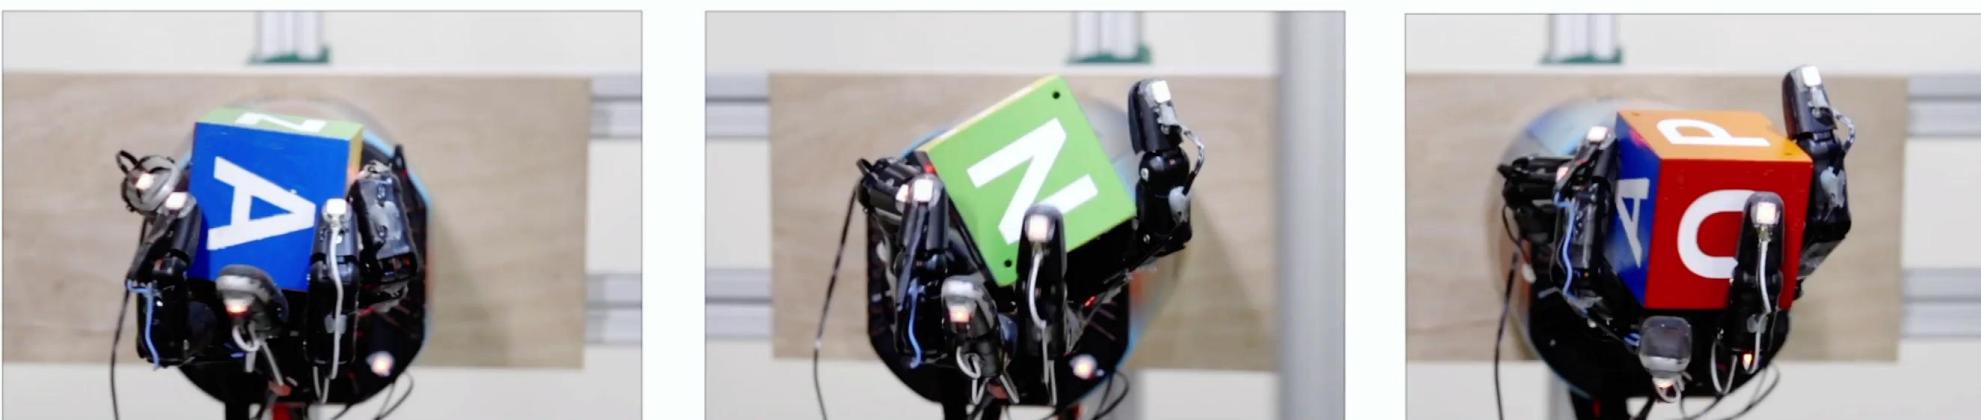
Reinforcement learning has a number of successes



AlphaGo (Google DeepMind)



Atari (Google DeepMind)



Robot manipulation (OpenAI)



ChatGPT (OpenAI)

Multi-agent RL has had some successes



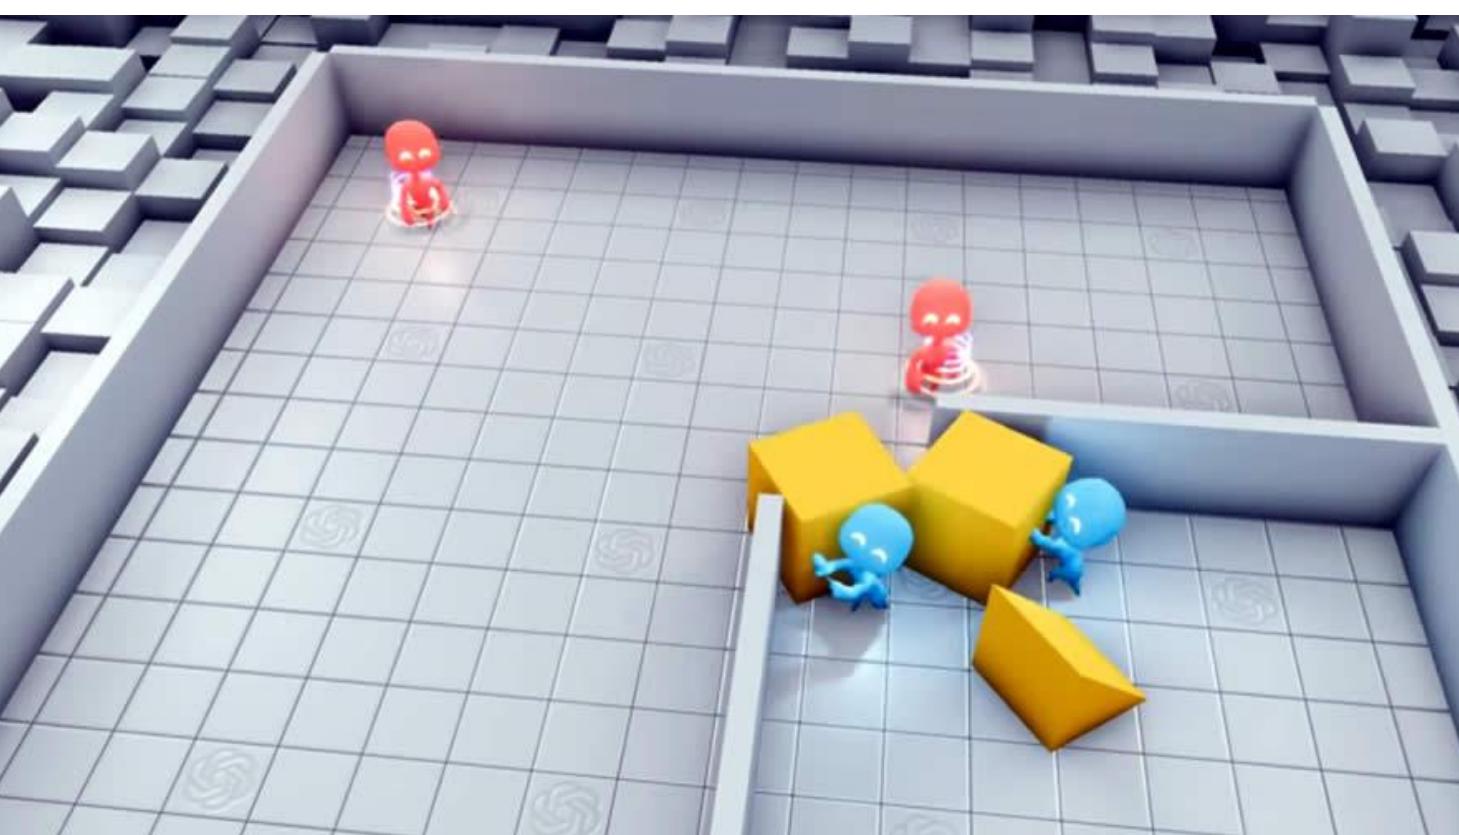
Project Malmo (Microsoft Research)



Humanoid Soccer (Google DeepMind)



OpenAI Five Dota 2 (OpenAI)



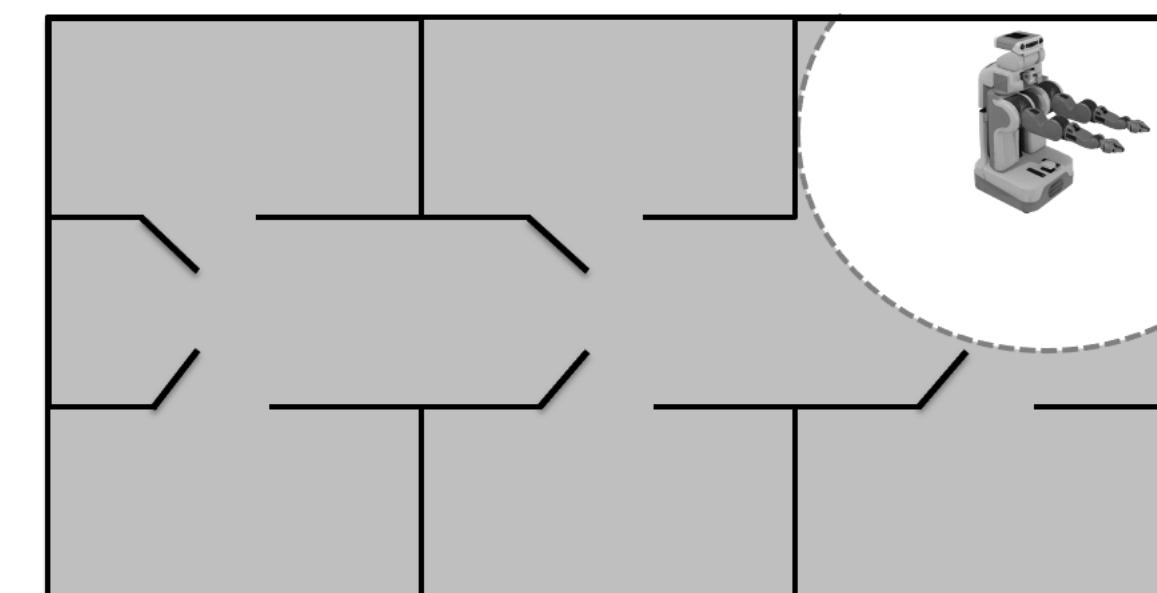
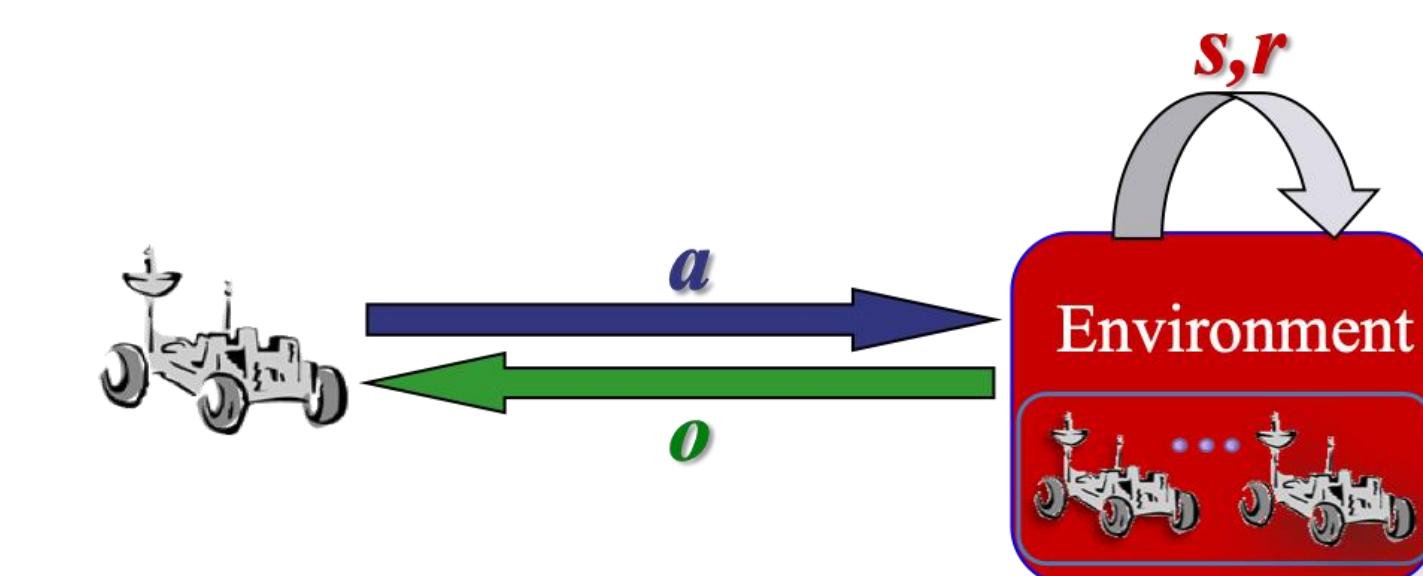
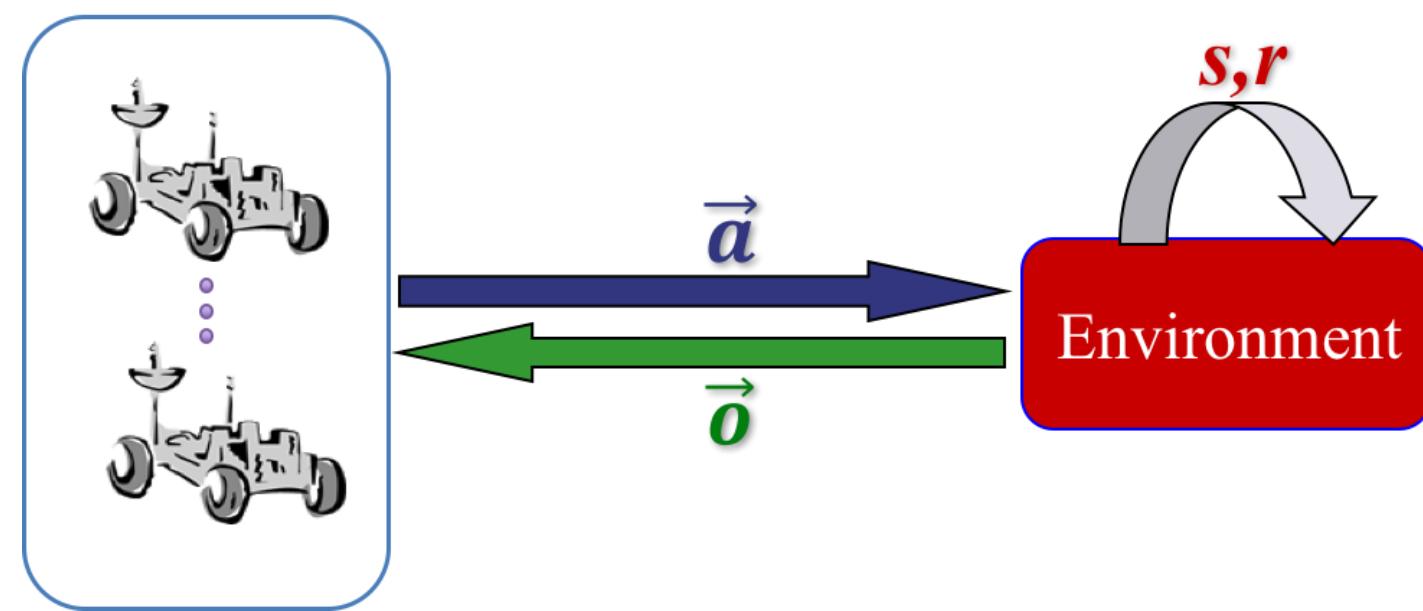
Emergent tool use (OpenAI)



AlphaStar (Google DeepMind)

Multi-agent RL is hard

- How can we apply RL here?
 - *Centralized learning and control?* Need fast, perfect communication (cooperative case)
 - *Decentralized learning and control?* Limited knowledge of other agents and environmental nonstationarity
 - *Centralized training for decentralized execution?* Use centralized information offline but still execute in a decentralized way
 - Almost always have partial observability

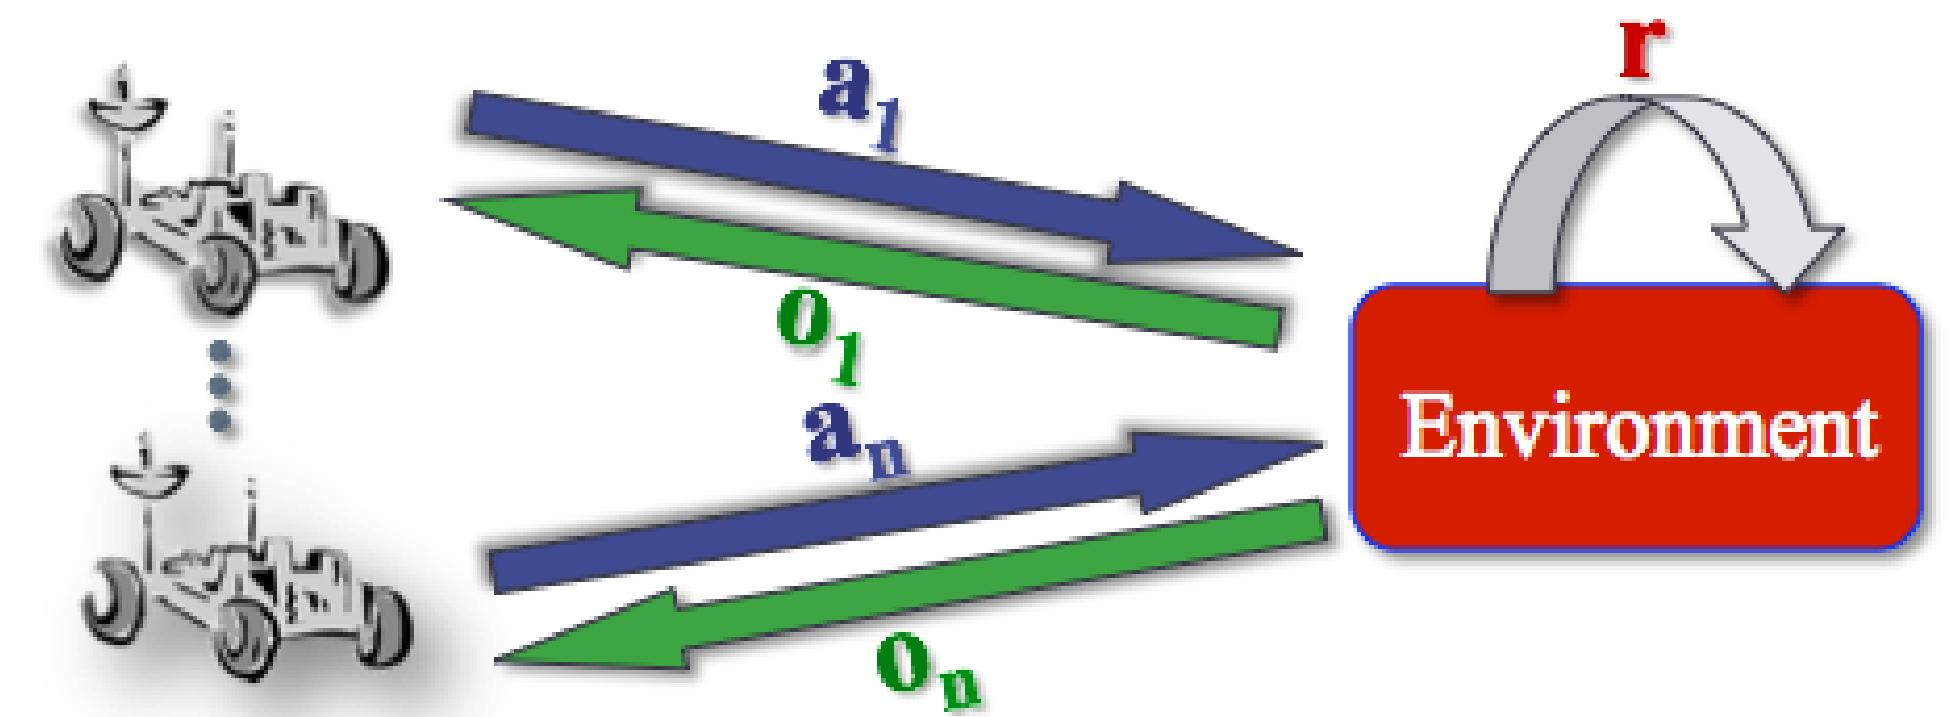


Overview

- Define the cooperative multi-agent RL (MARL) problem
- Discuss the current state-of-the art for the different classes of solutions
 - Centralized training and execution
 - Decentralized training and execution: IQL, decentralized REINFORCE, deep extensions
 - CTDE: VDN, QMIX, QPLEX, MADDPG, MAPPO
- Identify misconceptions/issues with current methods
- Applications, code, other topics, and the future (LLMs?)

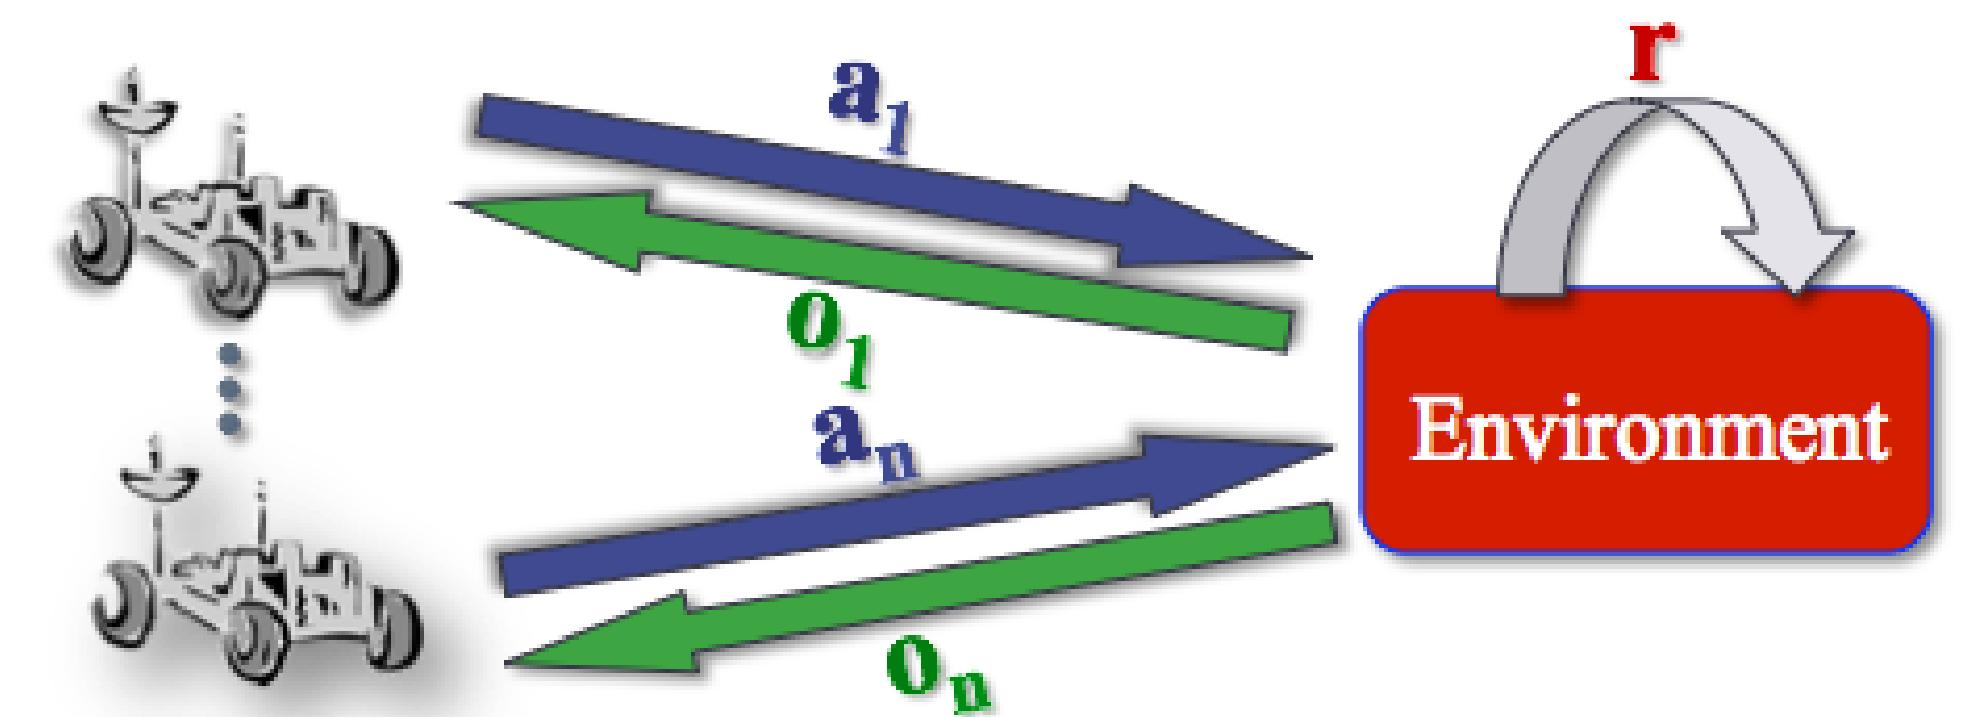
Cooperative MARL

- Cooperative case represented as Decentralized POMDP: $\langle I, S, \{A_i\}, T, R, \{\Omega_i\}, O, \mathbb{P} \rangle$
 - I , a finite set of agents
 - S , a set of states
 - A_i , each agent's set of actions
 - T , the state transition model: $P(s'|s, a)$
 - R , the reward model: $R(s, a)$
 - Ω_i , each agent's finite set of observations
 - O , the observation model: $P(o|s', a)$
 - h , horizon or discount \mathbb{P}



Cooperative MARL

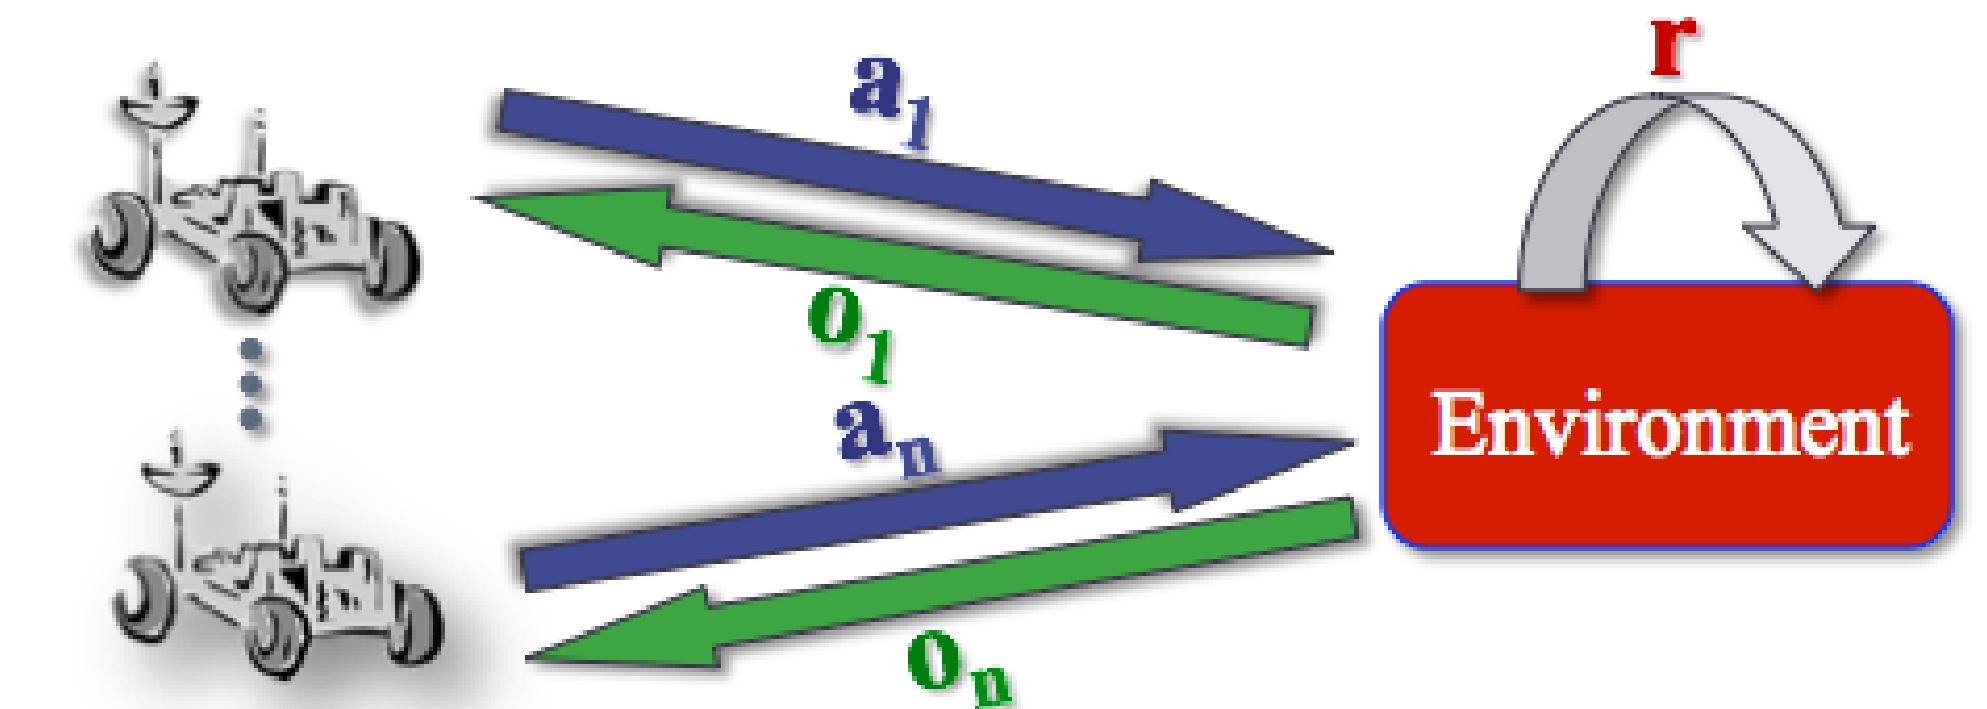
- Cooperative case represented as Decentralized POMDP: $\langle I, S, \{A_i\}, T, R, \{\Omega_i\}, O, \mathbb{P} \rangle$
 - I , a finite set of agents
 - S , a set of states
 - A_i , each agent's set of actions
 - T , the state transition model: $P(s'|s, a)$
 - R , the reward model: $R(s, a)$
 - Ω_i , each agent's finite set of observations
 - O , the observation model: $P(o|s', a)$
 - h , horizon or discount \mathbb{P}



Objective: Maximize the (discounted) sum of future **(joint) rewards** Cooperative

Cooperative MARL

- Cooperative case represented as Decentralized POMDP: $\langle I, S, \{A_i\}, T, R, \{\Omega_i\}, O, \mathbb{P} \rangle$
 - I , a finite set of agents
 - S , a set of states
 - A_i , each agent's set of actions
 - T , the state transition model: $P(s'|s, a)$
 - R , the reward model: $R(s, a)$
 - Ω_i , each agent's finite set of observations
 - O , the observation model: $P(o|s', a)$
 - h , horizon or discount \mathbb{P}



Objective: Maximize the (discounted) sum of future (joint) rewards

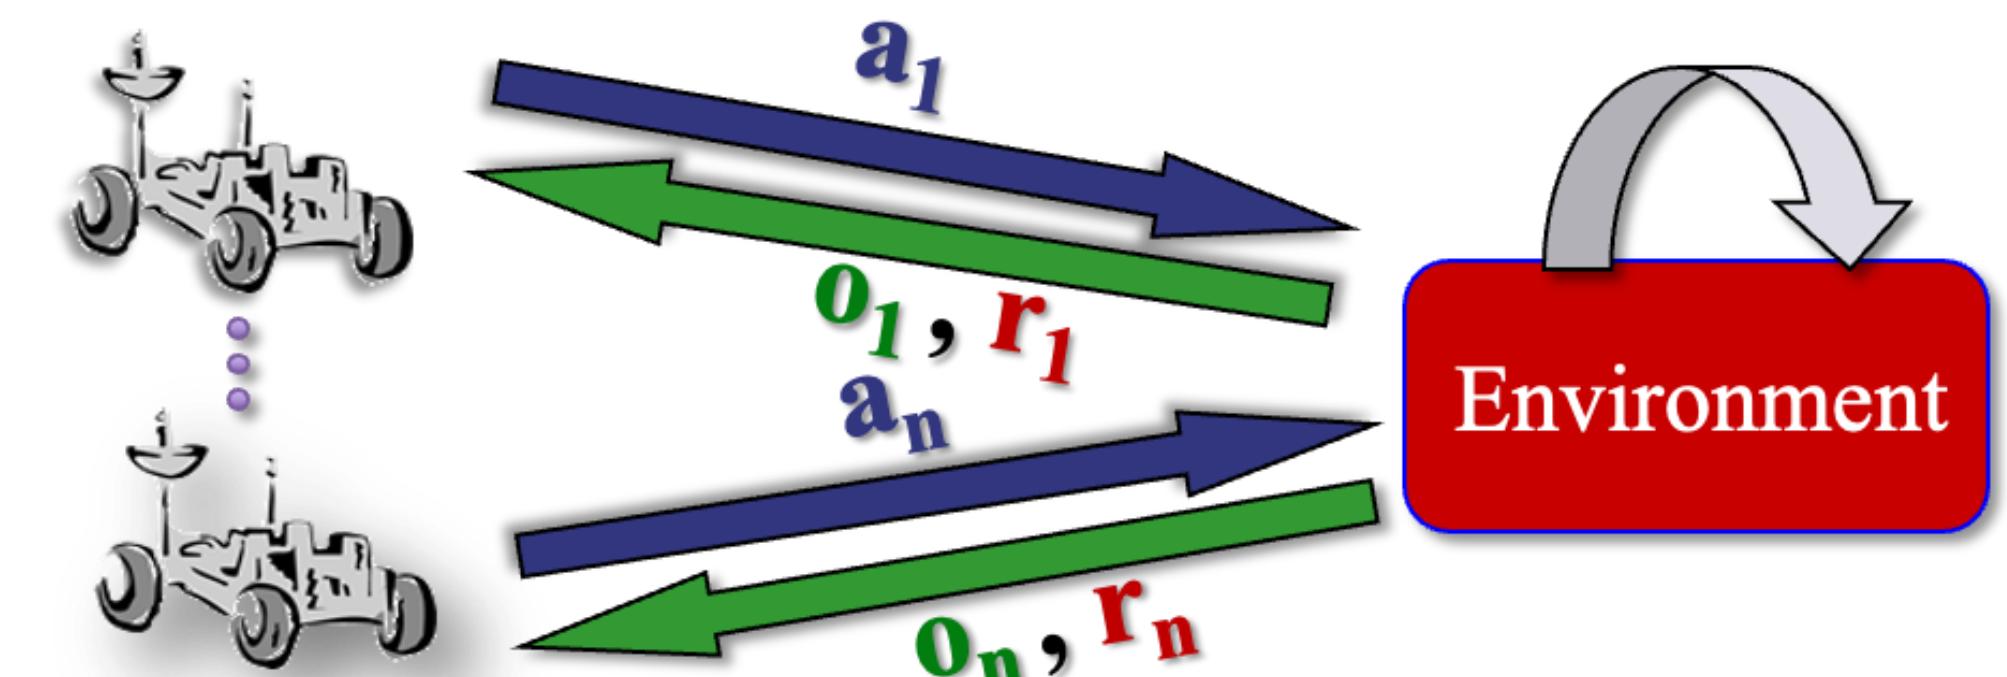
Calculate a set of optimal policies for each agent $\pi_i^*: H_i \rightarrow A_i$ that maximize joint objective

Decentralized partially observable execution

General MARL

- General case as Partially Observable Stochastic Game (POSG): $\langle I, S, \{A_i\}, T, \{R_i\}, \{\Omega_i\}, O, \mathbb{P} \rangle$
 - I , a finite set of agents
 - S , a set of states
 - A_i , each agent's set of actions
 - T , the state transition model: $P(s'|s, a)$
 - R_i , the reward model: $R_i(s, a)$
 - Ω_i , each agent's finite set of observations
 - O , the observation model: $P(o|s', a)$
 - h , horizon or discount \mathbb{P}

Objective unclear: Some form of each agent maximizing the (discounted) sum of future **individual rewards**
Mixed/competitive



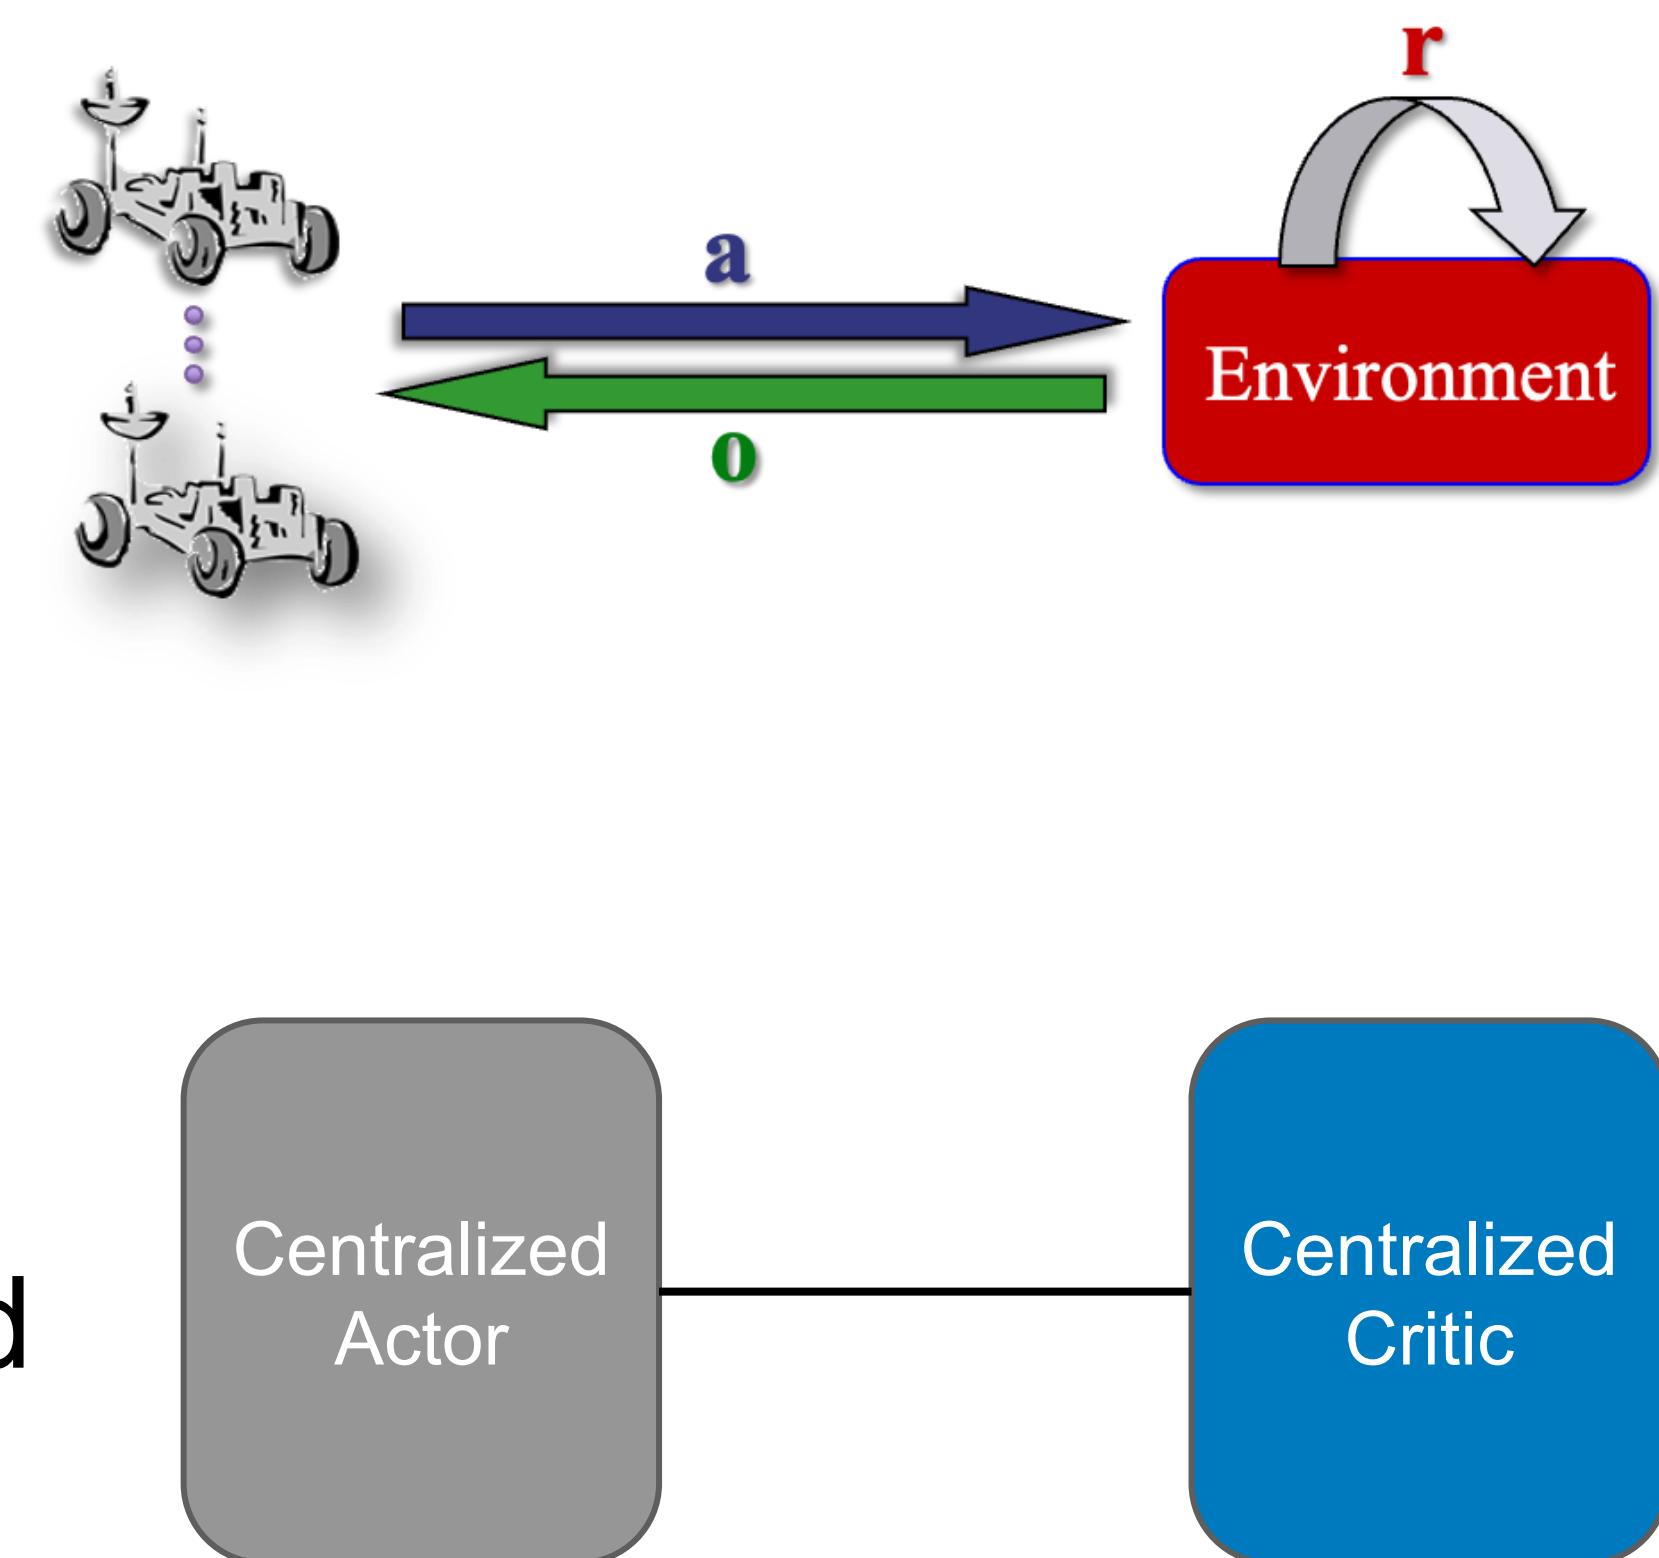
Centralized MARL

Models and methods

Centralized MARL

Assumptions:

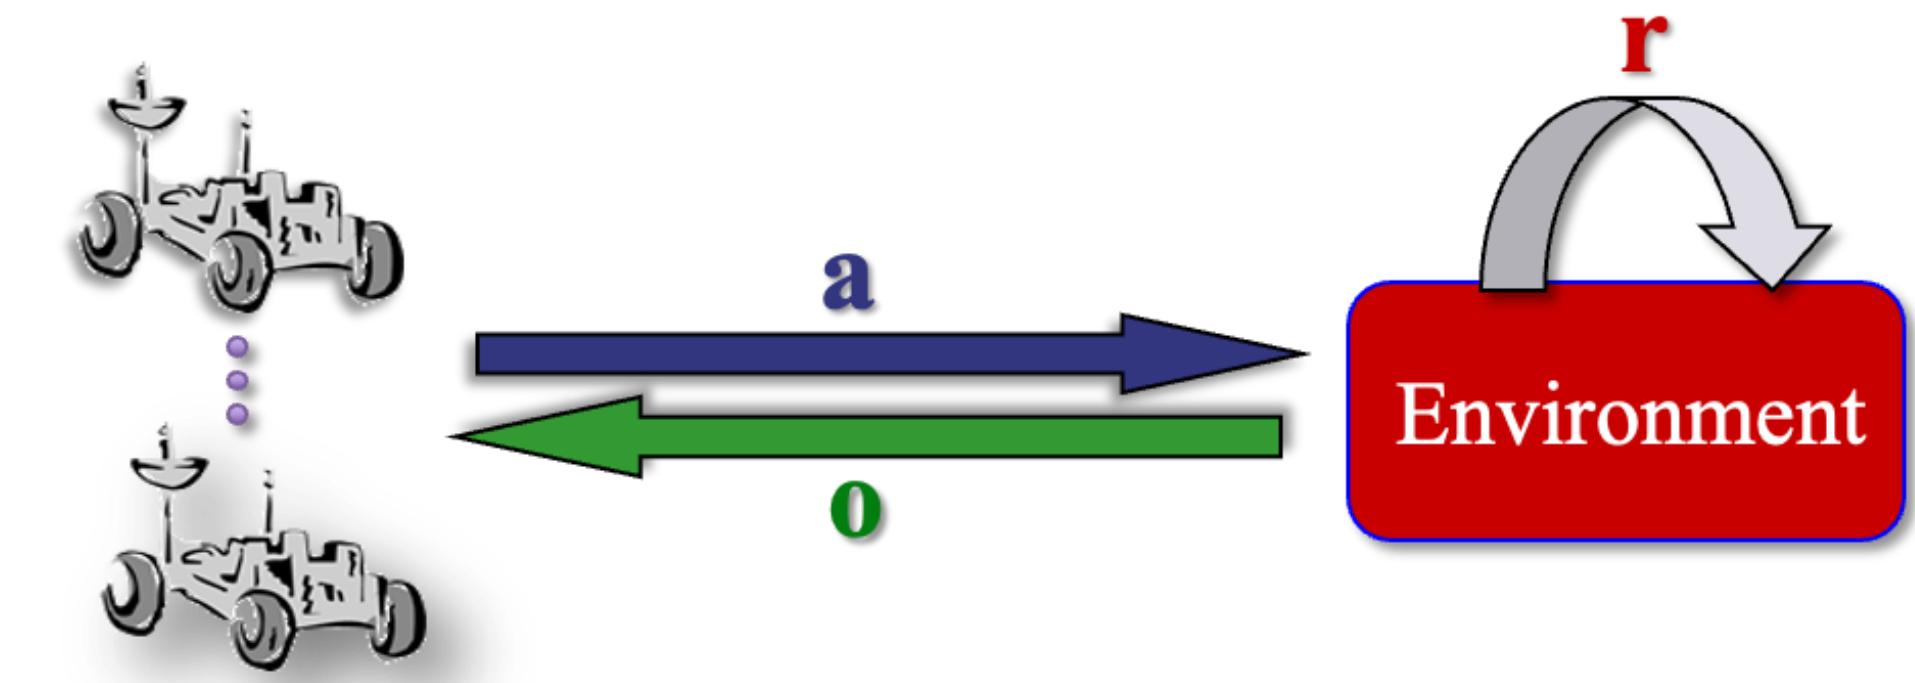
- a centralized controller chooses actions for each agent, a
- each agent takes the chosen actions $a = \langle a_1, \dots, a_n \rangle$,
- the centralized controller observes the resulting observations $o = \langle o_1, \dots, o_n \rangle$
- the (centralized) algorithm/controller observes o (and a) and the joint reward r



Note: Not a Dec-POMDP (or POSG) anymore since execution is centralized

Centralized MARL (partially observable)

- Cooperative case represented as MPOMDP: $\langle I, S, \{A_i\}, T, R, \{\Omega_i\}, O, \mathbb{H} \rangle$
 - I , a finite set of agents
 - S , a set of states
 - A_i , each agent's set of actions
 - T , the state transition model: $P(s'|s, a)$
 - R , the reward model: $R(s, a)$
 - Ω_i , each agent's finite set of observations
 - O , the observation model: $P(o|s', a)$
 - h , horizon or discount \mathbb{H}

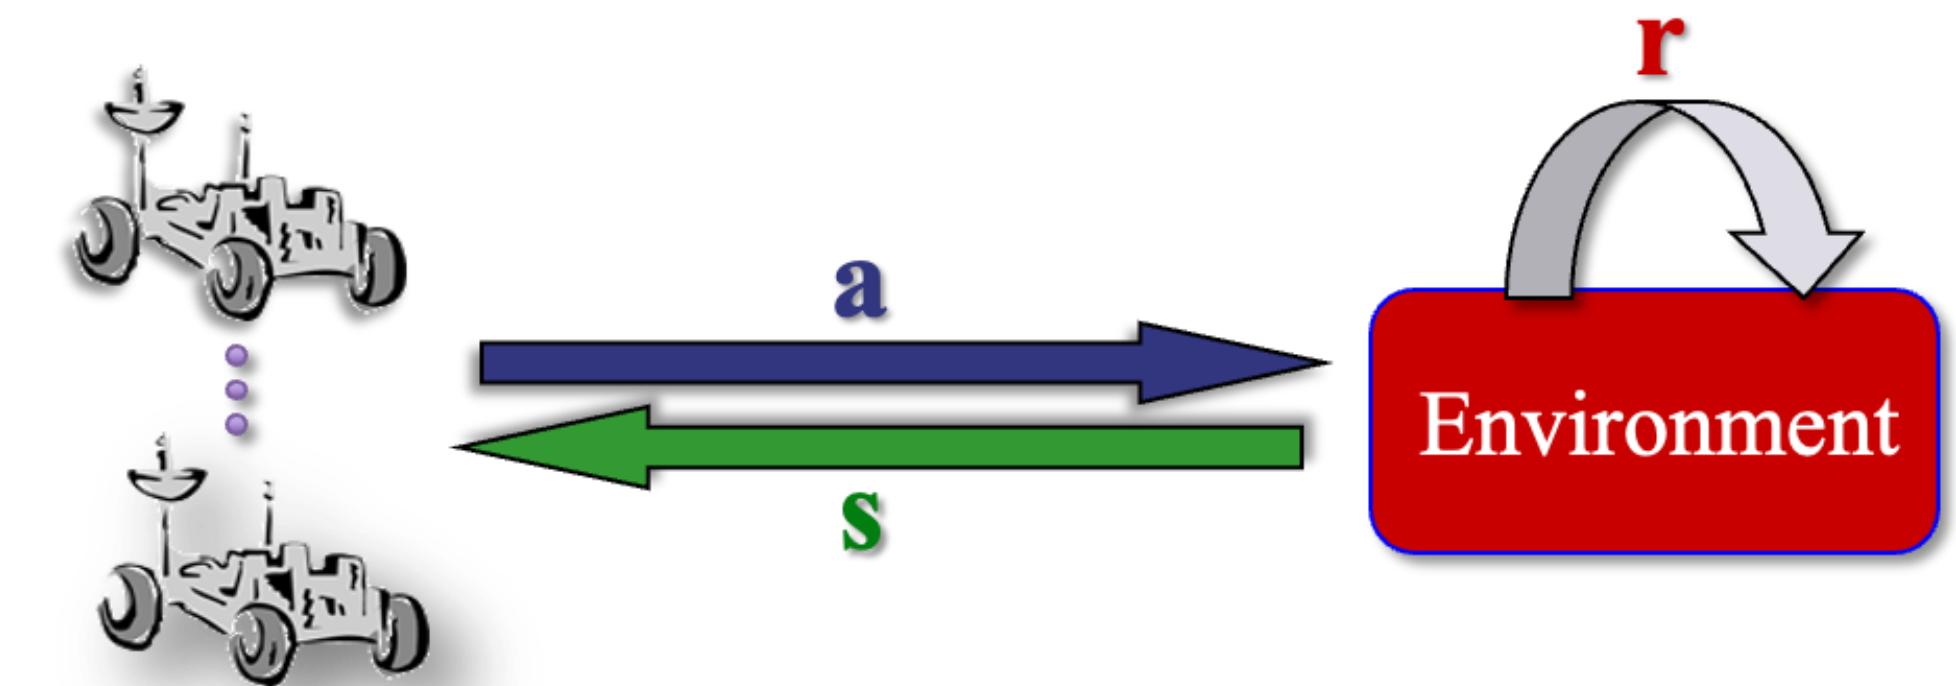


Objective: Maximize the (discounted) sum of future (joint) rewards

Calculate a single optimal policy for all agents $\pi^*: H \rightarrow A$ that maximizes centralized objective

Centralized MARL (fully observable)

- Cooperative case represented as MMDP: $\langle I, S, \{A_i\}, T, R, \mathbb{D} \rangle$
 - I , a finite set of agents
 - S , a set of states
 - A_i , each agent's set of actions
 - T , the state transition model: $P(s'|s, a)$
 - R , the reward model: $R(s, a)$
 - h , horizon or discount \mathbb{D}



Objective: Maximize the (discounted) sum of future (joint) rewards

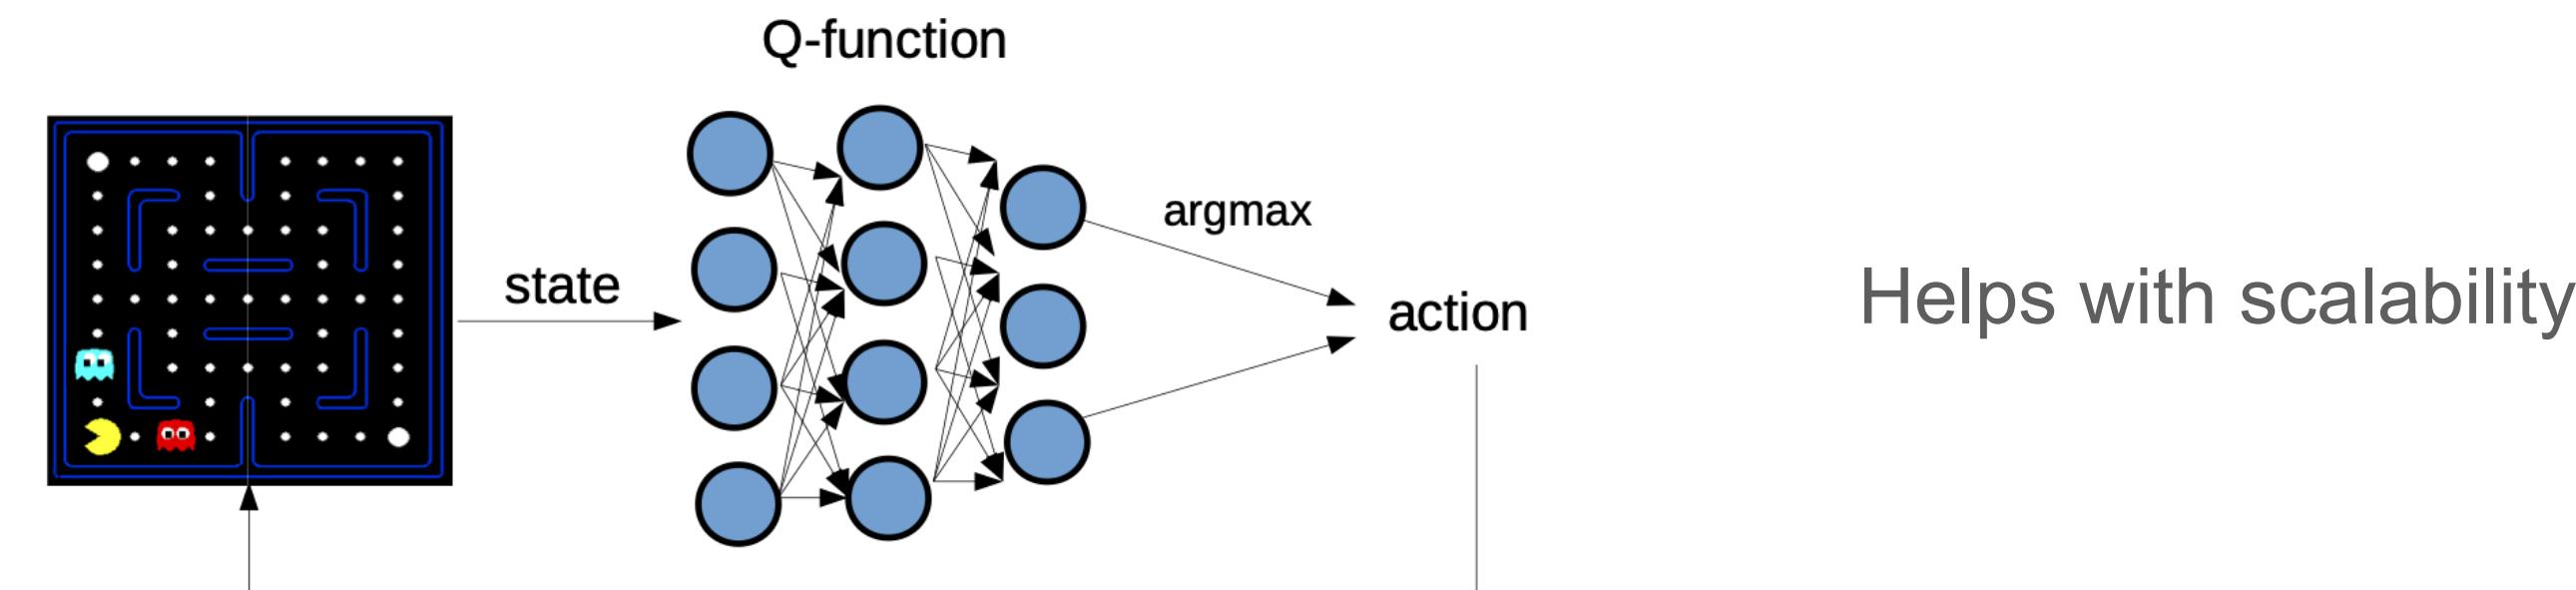
Calculate a single optimal policy for all agents $\pi^*: S \rightarrow A$ that maximizes centralized objective

Centralized MARL (DRQN version)

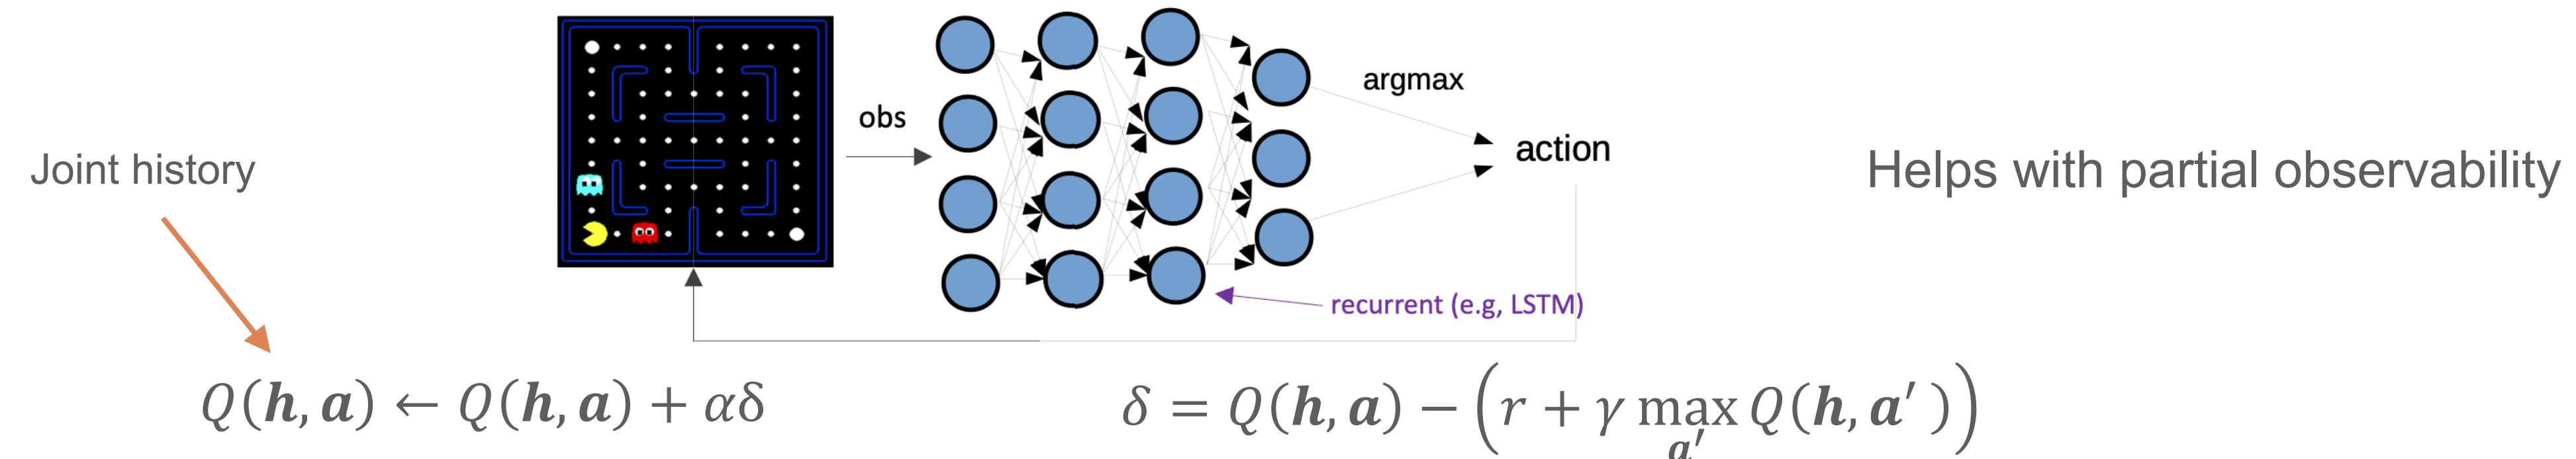
- Traditional Q-learning: estimate Q-value with (x can be state, observation or history)

$$Q(x, a) \leftarrow Q(x, a) + \alpha \delta \quad \text{For learning rate } \alpha$$
$$\delta = Q(x, a) - (r + \gamma \max_{a'} Q(x', a'))$$

- Deep Q-Networks (DQN) (Mnih et al., Nature 15) uses a neural net for function approximation

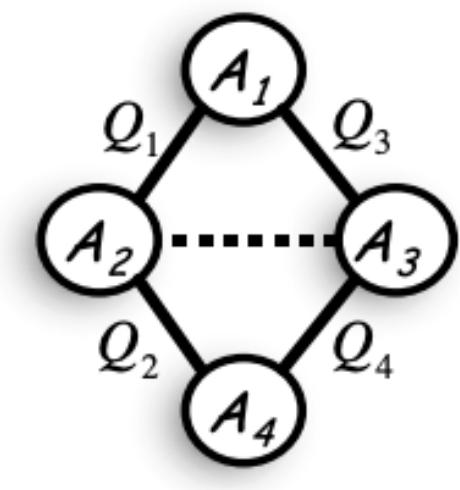
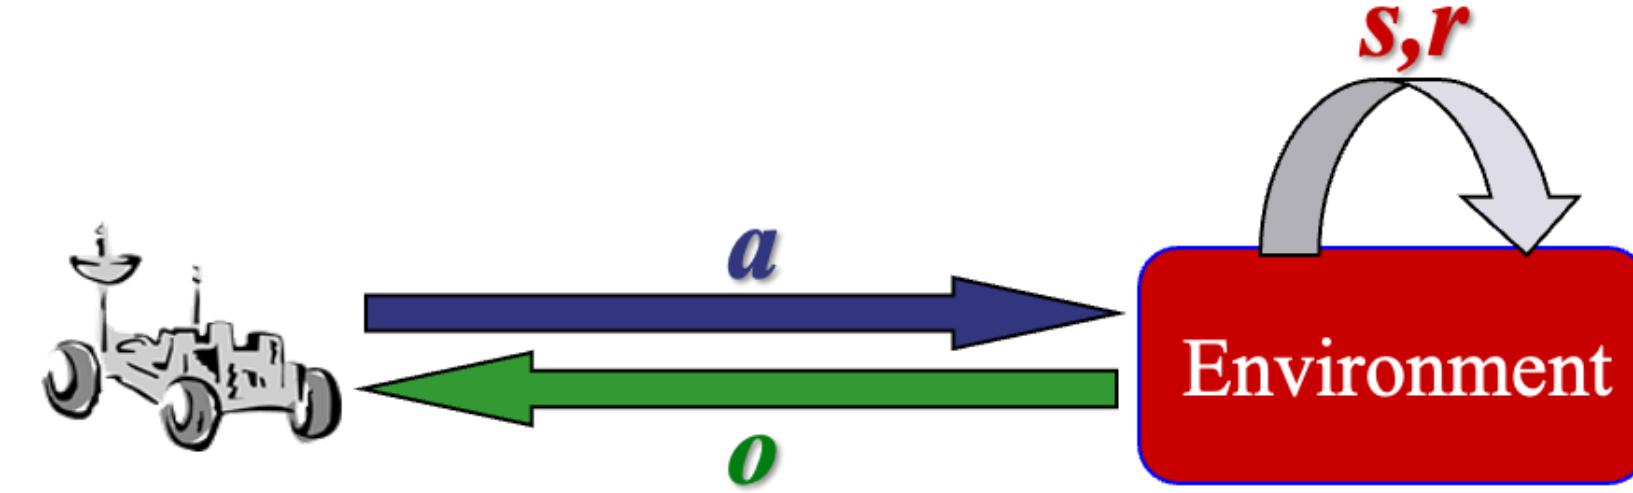


- DRQN (Hausknecht and Stone, arXiv 15) adds a recurrent layer for memory



Centralized MARL methods

- Now just a (factored) single-agent problem
 - Multi-agent MDP or POMDP (not Dec-POMDP/POSG)
 - Can use any single-agent RL method
 - But it doesn't scale well
 - And assumes centralized information and control
 - Some methods exploit multi-agent factorization but not very active
 - Coordination graphs [Guestrin et al., 2001]
 - AlphaStar [Vinyals et al., 2019]



Decentralizing centralized solutions

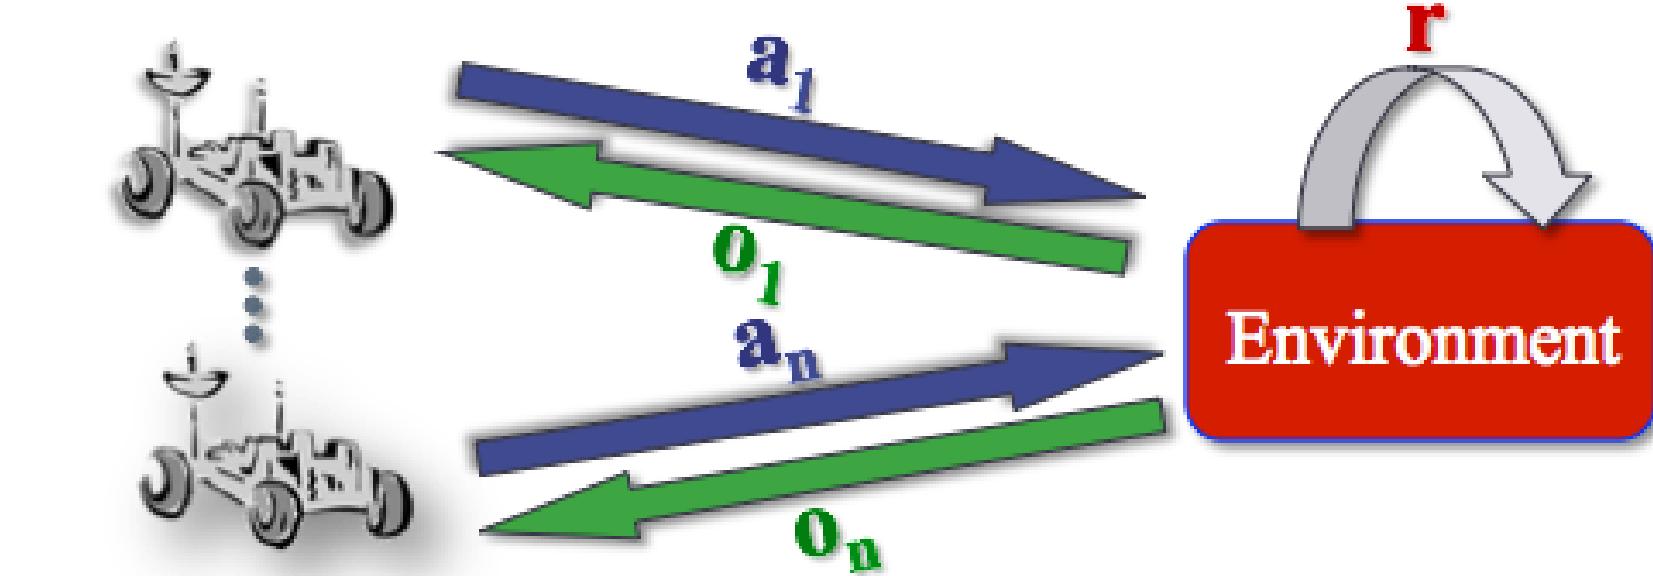
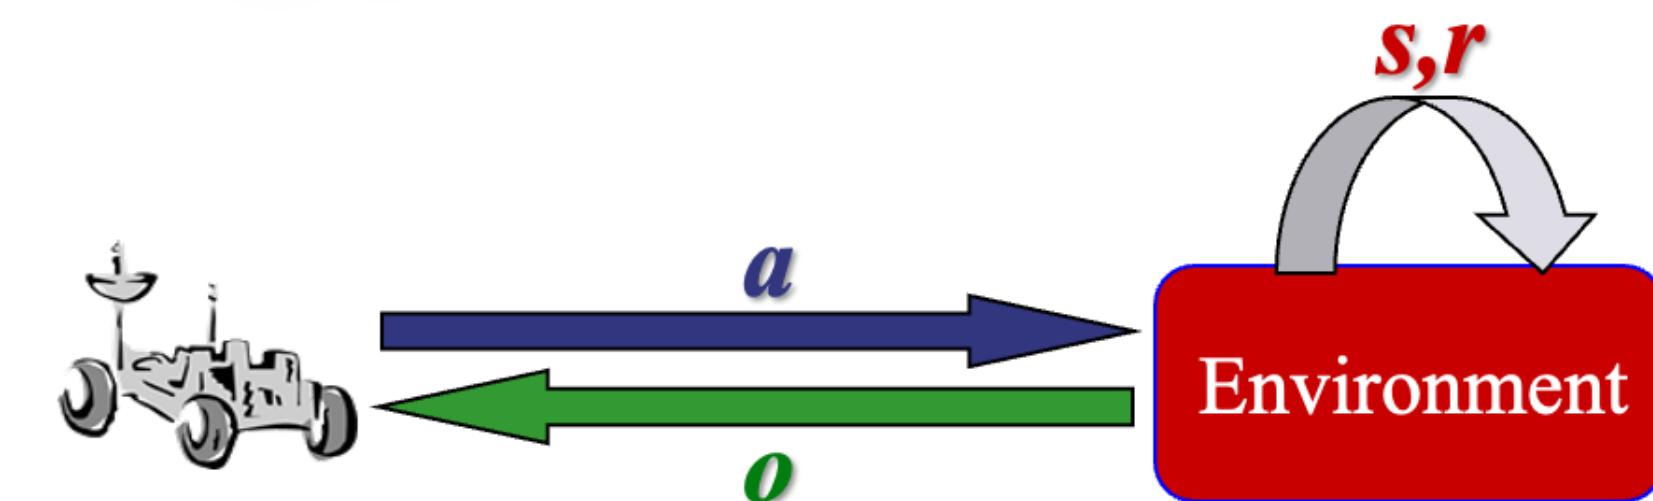
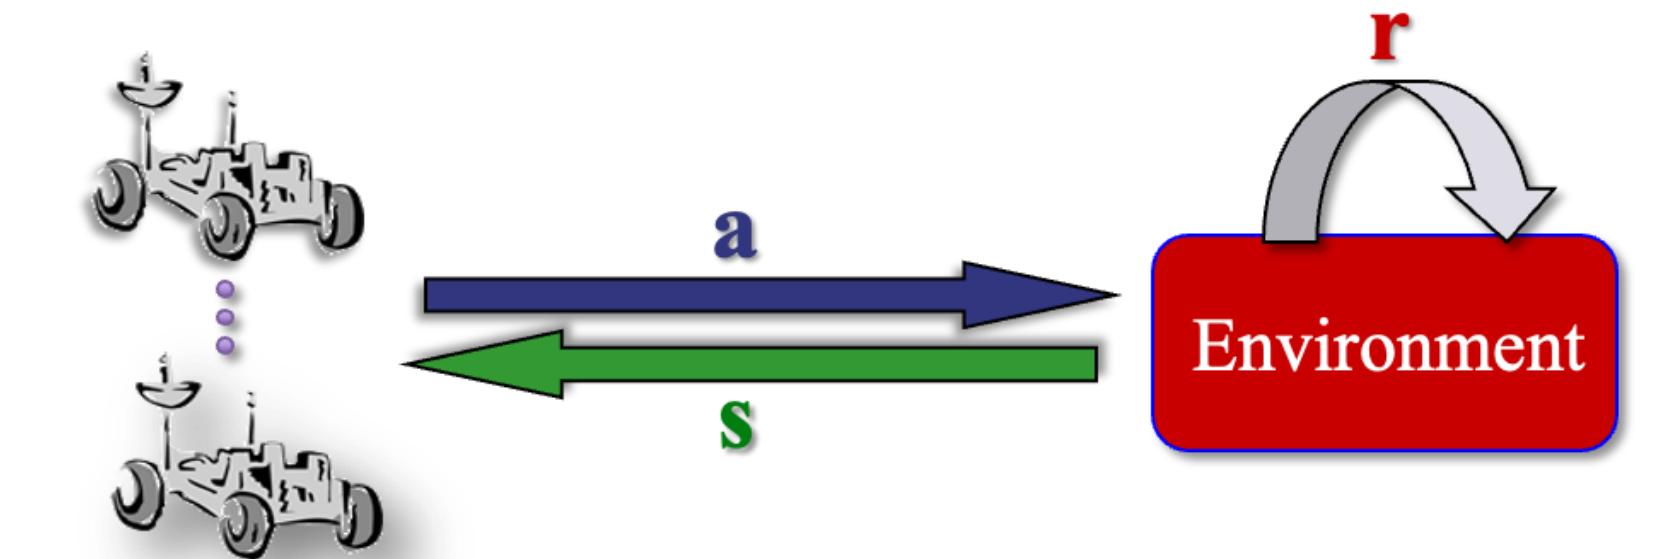
Easy to ‘decentralize’ in a MMDP or MPOMDP

- MMDP
 - $S \rightarrow A$ or $S \rightarrow A_i$
- MPOMDP
 - $H \rightarrow A$ or $H \rightarrow A_i$

Hard in a Dec-POMDP

Once you have $H \rightarrow A$ how do you get $H_i \rightarrow A_i$?

Easy to decentralize control but not information



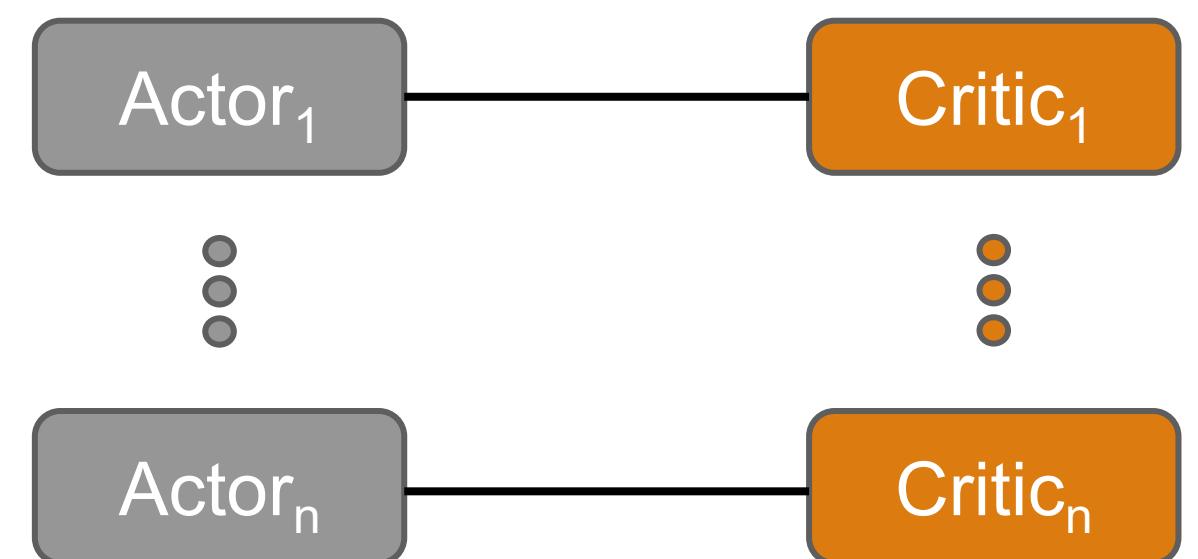
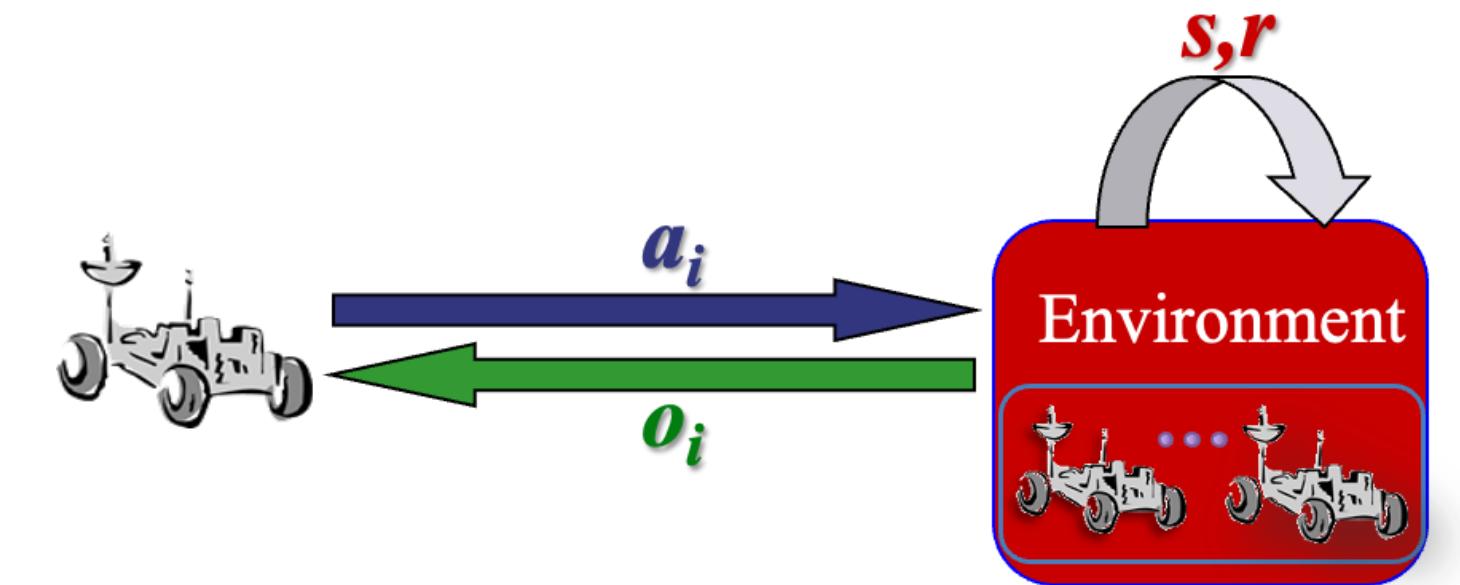
Decentralized MARL

Models and methods

Decentralized MARL

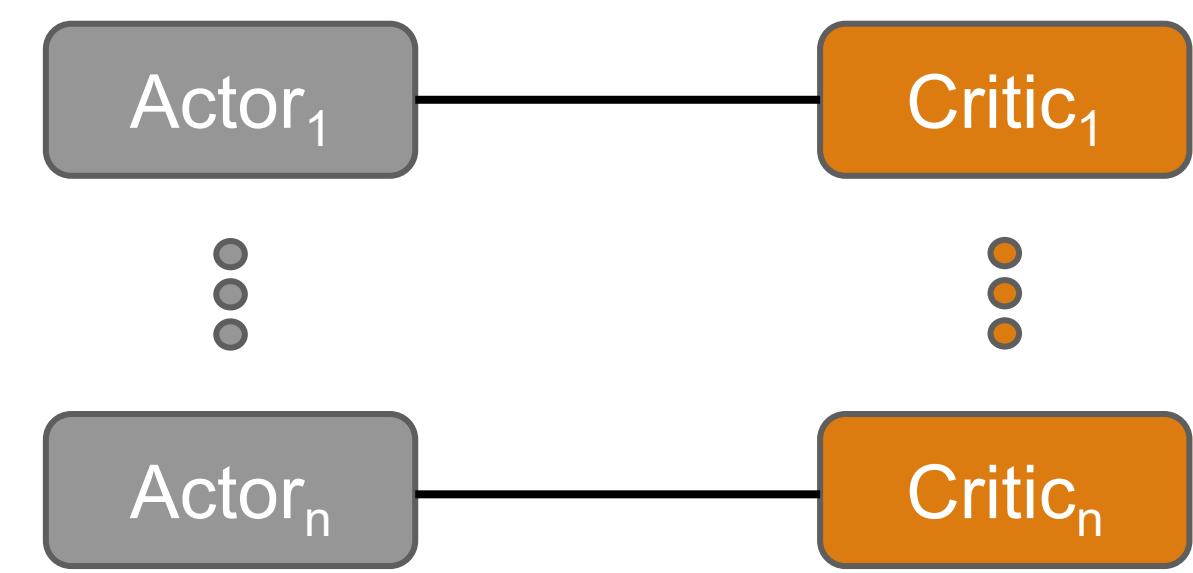
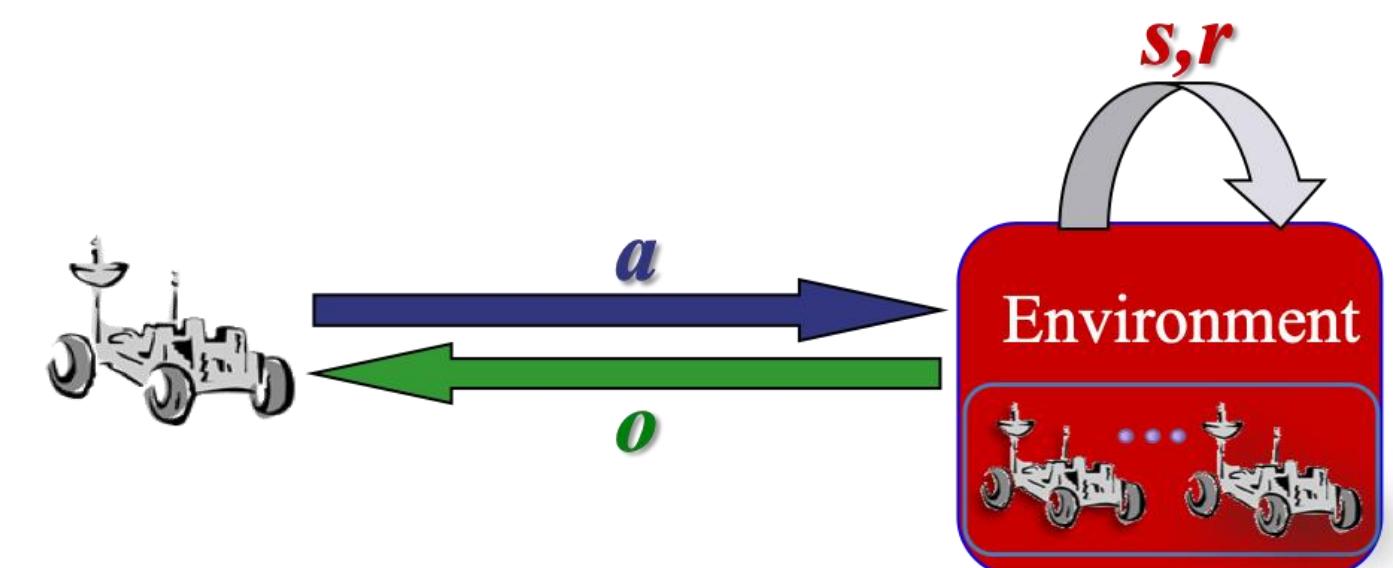
Assumptions:

- each agent, i , observes its current observation, o_i , and takes action a_i at the resulting history, h_i ,
- the (decentralized) algorithm/controller sees the same information (o_i and a_i) as well as the joint reward r .



Decentralized MARL

- Agents each learn separately
 - Assumes training and execution are decentralized (e.g., lack of communication)
 - Is more scalable
 - The realistic case for POSGs and online learning in Dec-POMDPs
- Each agent i learns a policy that maps from *local histories* to *local actions* π_i :
 $H_i \rightarrow A_i$
- Can also use any single-agent method here
 - May be nonstationarity but there are many methods for dealing with that
 - Many improvements: Distributed Q, ICML-00; Hysteretic Q, IROS-07, ICML-17; Lenient Q JMLR-08, AAMAS-18; Likelihood Q, AAMAS-20; IPPO arxiv-20



Decentralized Action-Value Methods

**IQL, Distributed Q, Hysteretic Q, Lenient Q
Deep extensions**

Independent Q-Learning (IQL)

[Tan – ICML 93](#)

- Just apply Q-learning pretending the other agents don't exist

Algorithm 1 Independent Q-Learning for agent i (finite-horizon)

```
1: set  $\alpha$  and  $\epsilon$  (learning rate, exploration)
2: Initialize  $Q_i$  for all  $h_i \in \mathbb{H}_i, a_i \in \mathbb{A}_i$ 
3: for all episodes do
4:    $h_i \leftarrow \emptyset$                                      {Empty initial history}
5:   for  $t = 1$  to  $\mathcal{H}$  do
6:     Choose  $a_i$  at  $h_i$  from  $Q_i(h_i, \cdot)$  with exploration (e.g.,  $\epsilon$ -greedy)
7:     See joint reward  $r$ , local observation  $o_i$                       {Depends on joint action a}
8:      $h'_i \leftarrow h_i a_i o_i$ 
9:      $Q_i(h_i, a_i) \leftarrow Q_i(h_i, a_i) + \alpha [r + \gamma \max_{a'_i} Q_i(h'_i, a'_i) - Q_i(h_i, a_i)]$ 
10:     $h_i \leftarrow h'_i$ 
11:  end for
12: end for
13: return  $Q_i$ 
```

Independent Q-Learning (IQL)

Tan – ICML 93

- Just apply Q-learning pretending the other agents don't exist
- Where do the observations and joint rewards come from?

Algorithm 1 Independent Q-Learning for agent i (finite-horizon)

```
1: set  $\alpha$  and  $\epsilon$  (learning rate, exploration)
2: Initialize  $Q_i$  for all  $h_i \in \mathbb{H}_i, a_i \in \mathbb{A}_i$            $P(o|s', a)$        $P(s'|s, a)$ 
3: for all episodes do
4:    $h_i \leftarrow \emptyset$                                          {Empty initial history}
5:   for  $t = 1$  to  $\mathcal{H}$  do
6:     Choose  $a_i$  at  $h_i$  from  $Q_i(h_i, \cdot)$  with exploration (e.g.,  $\epsilon$ -greedy)    {Depends on joint action a}
7:     See joint reward  $r$ , local observation  $o_i$ 
8:      $h'_i \leftarrow h_i a_i o_i$ 
9:      $Q_i(h_i, a_i) \leftarrow Q_i(h_i, a_i) + \alpha [r + \gamma \max_{a'_i} Q_i(h'_i, a'_i) - Q_i(h_i, a_i)]$ 
10:     $h_i \leftarrow h'_i$ 
11:   end for
12: end for
13: return  $Q_i$ 
```

Important hidden information

- Agents don't exist by themselves!
- Assumes other agents are acting according to some (fixed) policies
- Then learns as if in a POMDP where other agents are part of the environment:

$$Q_i(h_i, a_i) = \sum_{\mathbf{a} \in \mathbb{A}} \hat{P}(\mathbf{a}, \mathbf{h} | h_i, a_i) \left[r + \gamma \sum_{o_i} \hat{P}(o_i | \mathbf{h}, \mathbf{a}) \max_{a'_i} Q_i(h'_i, a'_i) \right]$$

- This is where non-stationarity comes from!
 - Other learning agents change their policies over time

\hat{P} s are empirical probabilities from data during training

IQL properties

- IQL may not converge (Tan ICML 93)
- Convergence properties of Q-learning in Dec-POMDPs is an open question!
- Usually performs poorly (often used as a baseline)
- Note even with optimal Q-values, agents may not select the optimal action without coordination when multiple actions are optimal (like equilibrium selection)

$$Q_1(h_1, a_1^1) = Q_1(h_1, a_1^2)$$

$$Q_2(h_2, a_2^1) = Q_2(h_2, a_2^2)$$

$$Q(h_1, h_2, a_1^1, a_2^2) = Q(h_1, h_2, a_1^2, a_2^1) < Q(h_1, h_2, a_1^2, a_2^2) = Q(h_1, h_2, a_1^1, a_2^1)$$

Extension to the deep case - IDRQN

[Tampuu et al. – Plos one 17](#)

- Just DRQN applied to the multi-agent case
- Still needs other agents to act

Algorithm 2 Independent DRQN (IDRQN) for agent i (finite-horizon*)

```
1: set  $\alpha$ ,  $\epsilon$ , and  $C$  (learning rate, exploration, and target update frequency)
2: Initialize network parameters  $\theta$  and  $\theta^-$  for  $Q_i$ 
3:  $\mathcal{D}_i \leftarrow \emptyset$ 
4:  $e \leftarrow 1$ 
5: for all episodes do
6:    $h_i \leftarrow \emptyset$                                 {initial history is empty}
7:   for  $t = 1$  to  $\mathcal{H}$  do
8:     Choose  $a_i$  at  $h_i$  from  $Q_i^\theta(h_i, \cdot)$  with exploration (e.g.,  $\epsilon$ -greedy)
9:     See joint reward  $r$ , local observation  $o_i$                                 {Depends on joint action a}
10:    append  $a_i, o_i, r$  to  $\mathcal{D}_i^e$ 
11:     $h_i \leftarrow h_i a_i o_i$                                      {update RNN state of the network}
12:   end for
13:   sample an episode from  $\mathcal{D}$                                 Based on other agents
14:   for  $t = 1$  to  $\mathcal{H}$  do
15:      $h_i \leftarrow \emptyset$ 
16:      $a_i, o_i, r \leftarrow \mathcal{D}_i^e(t)$ 
17:      $h'_i \leftarrow h_i a_i o_i$ 
18:      $y = r + \gamma \max_{a'_i} Q_i^{\theta^-}(h'_i, a'_i)$ 
19:     Perform gradient descent on parameters  $\theta$  with learning rate  $\alpha$  and loss:  $(y - Q_i^\theta(h_i, a_i))^2$ 

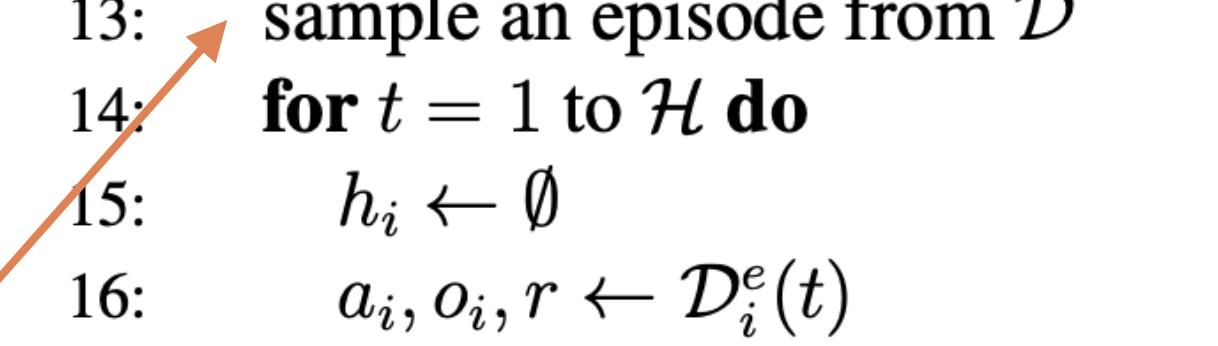
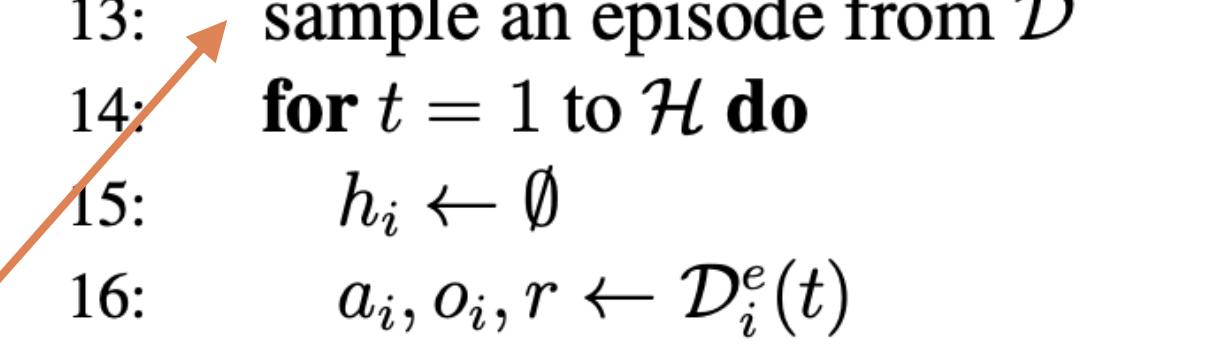
20:      $h_i \leftarrow h'_i$ 
21:   end for
22:   if  $e \bmod C = 0$  then
23:      $\theta^- \leftarrow \theta$ 
24:   end if
25:    $e \leftarrow e + 1$ 
26: end for
27: return  $Q_i$ 
```

Extension to the deep case - IDRQN

[Tampuu et al. – Plos one 17](#)

- Just DRQN applied to the multi-agent case
- Still needs other agents to act
- Independent buffers cause poor performance (non-stationarity)

Algorithm 2 Independent DRQN (IDRQN) for agent i (finite-horizon*)

```
1: set  $\alpha$ ,  $\epsilon$ , and  $C$  (learning rate, exploration, and target update frequency)
2: Initialize network parameters  $\theta$  and  $\theta^-$  for  $Q_i$ 
3:  $\mathcal{D}_i \leftarrow \emptyset$ 
4:  $e \leftarrow 1$ 
5: for all episodes do
6:    $h_i \leftarrow \emptyset$ 
7:   for  $t = 1$  to  $\mathcal{H}$  do
8:     Choose  $a_i$  at  $h_i$  from  $Q_i^\theta(h_i, \cdot)$  with exploration (e.g.,  $\epsilon$ -greedy)
9:     See joint reward  $r$ , local observation  $o_i$ 
10:    append  $a_i, o_i, r$  to  $\mathcal{D}_i^e$ 
11:     $h_i \leftarrow h_i a_i o_i$  
12:   end for
13:   sample an episode from  $\mathcal{D}$  
14:   for  $t = 1$  to  $\mathcal{H}$  do
15:      $h_i \leftarrow \emptyset$ 
16:      $a_i, o_i, r \leftarrow \mathcal{D}_i^e(t)$ 
17:      $h'_i \leftarrow h_i a_i o_i$ 
18:      $y = r + \gamma \max_{a'_i} Q_i^{\theta^-}(h'_i, a'_i)$ 
19:     Perform gradient descent on parameters  $\theta$  with learning rate  $\alpha$  and loss:  $(y - Q_i^\theta(h_i, a_i))^2$ 
20:      $h_i \leftarrow h'_i$ 
21:   end for
22:   if  $e \bmod C = 0$  then
23:      $\theta^- \leftarrow \theta$ 
24:   end if
25:    $e \leftarrow e + 1$ 
26: end for
27: return  $Q_i$ 
```

{episode index}

{initial history is empty}

{Depends on joint action a}

{update RNN state of the network}

Based on other agents

$(y - Q_i^\theta(h_i, a_i))^2$

Decentralized MARL (Dec-HDRQN)

[Omidshafiei, Pazis, Amato, How and Vian - ICML 17](#)

- Traditional Q-learning: estimate Q-value with (x can be state, observation or history)

$$Q(x, a) \leftarrow Q(x, a) + \alpha \delta \quad \text{For learning rate } \alpha$$
$$\delta = Q(x, a) - (r + \gamma \max_{a'} Q(x', a'))$$

- Hysteresis (Matignon et al., IROS 07): two learning rates α and β (with $\beta < \alpha$)

$$Q(x, a) \leftarrow Q(x, a) + \beta \delta \quad \text{if } \delta \leq 0 \quad \text{Helps with coordination}$$
$$Q(x, a) + \alpha \delta \quad \text{otherwise}$$

- Still use DRQN (Hausknecht and Stone, arXiv 15) if partially observable



Local history

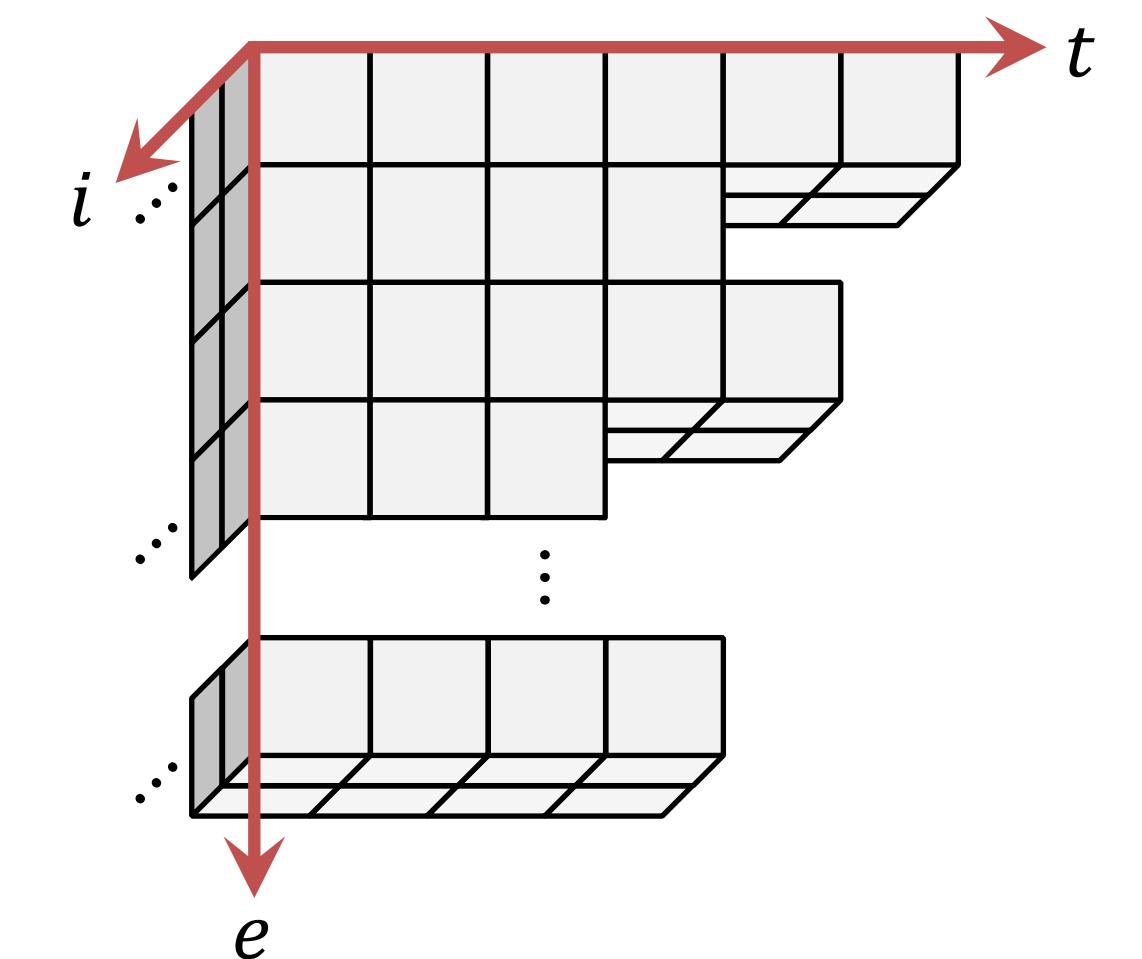
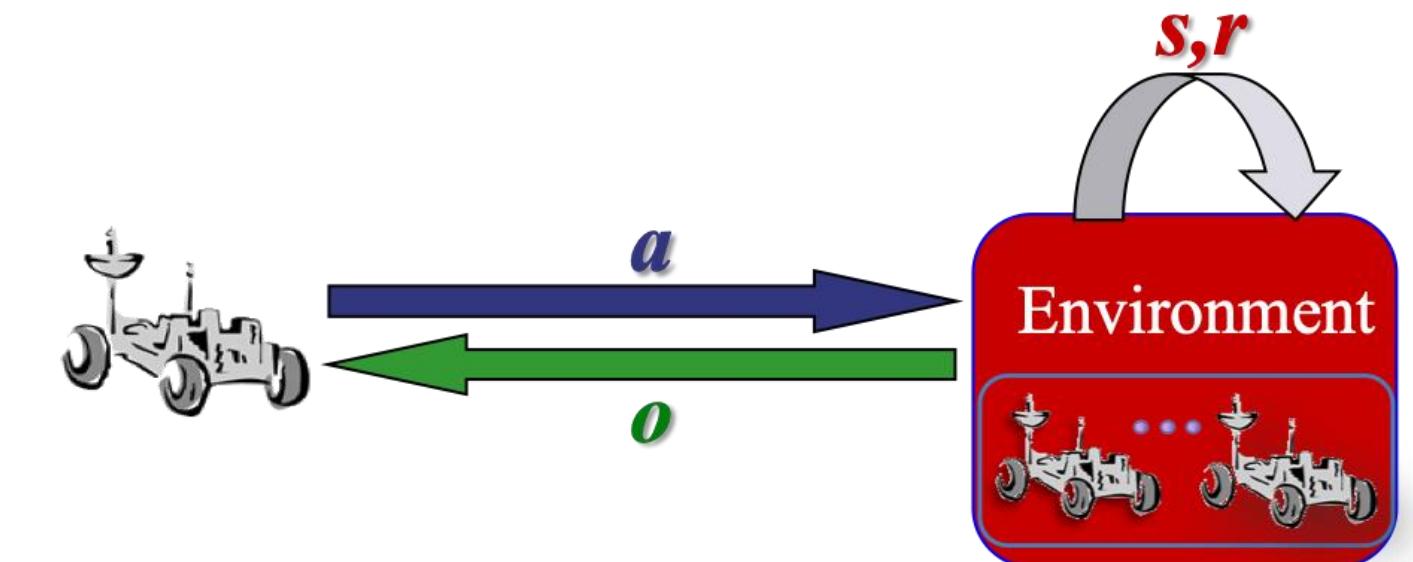
$$Q(h_i, a_i) \leftarrow Q(h_i, a_i) + \alpha \delta$$

$$\delta = Q(h_i, a_i) - (r + \gamma \max_{a'_i} Q(h_i, a'_i))$$

Decentralized Hysteretic DQN (Dec-HDRQN)

[Omidshafiei, Pazis, Amato, How and Vian - ICML 17](#)

- Dec-HDRQN algorithm overview
 - Use idea from previous slide to help with cooperation, scalability and partial observability
 - Each agent learns **concurrently** (not independently)
 - Use decentralized Concurrent Experience Replay Trajectories (CERTs) (synchronized buffers) to stabilize learning
 - Current decentralized methods (e.g., IPPO) also use some form of concurrent learning



Other deep decentralized methods

- Several other extensions of tabular and single agent methods
- Deep lenient Q-learning (Palmer et al. AAMAS 18)
 - Only for the fully observable case
 - Add leniency values to the replay buffer $(s_t, a_t, r_t, s_{t+1}, l(s_t, a_t))$ for $l(s_t, a_t) = 1 - e^{-K*T(\phi(s_t), a_t)}$
- Likelihood Q-learning (Lyu et al. AAMAS 20)
 - Uses distributional RL to estimate when other agents are exploring and use that info to adjust learning rate

Decentralized Policy Gradient Methods

Decentralized REINFORCE, IAC, IPPO

Decentralized REINFORCE

[Peshkin et al. – UAI 00](#)

- Extends single agent REINFORCE (Williams 92)
- Simple but has convergence guarantees!
 - joint gradient can be decomposed into decentralized gradients
 - I.e., this algorithm converges to the same values as a centralized algorithm (over decentralized policies)
 - Assumes concurrent learning

Algorithm 3 Decentralized REINFORCE for agent i (finite-horizon)

Require: Individual actor models $\pi_i(a_i|h_i)$, parameterized by ψ_i

```
1: set  $\alpha$  (learning rate)
2: for all episodes do
3:    $h_{i,0} \leftarrow \emptyset$                                 {Empty initial history}
4:    $ep \leftarrow \emptyset$                                 {Empty episode}
5:   for  $t = 0$  to  $\mathcal{H} - 1$  do
6:     Choose  $a_{i,t}$  at  $h_{i,t}$  from  $\pi_i(a_i|h_{i,t})$           Based on other agents
7:     See joint reward  $r_t$ , local observation  $o_{i,t}$           {Depends on joint action a}
8:     append  $a_{i,t}, o_{i,t}, r_t$  to  $ep$ 
9:      $h_{i,t+1} \leftarrow h_{i,t} a_{i,t} o_{i,t}$                   {Append new action and obs to previous history}
10:    end for
11:    for  $t = 0$  to  $\mathcal{H} - 1$  do
12:      Compute return at  $t$  from  $ep$ :  $G_{i,t} \leftarrow \sum_{k=t}^{\mathcal{H}-1} \gamma^{k-t} r_k$           Monte Carlo returns
13:      Update parameters:  $\psi_i \leftarrow \psi_i + \alpha \gamma^t G_{i,t} \nabla \log \pi_i(a_i|h_{i,t})$ 
14:    end for
15:  end for
```

Independent actor critic (IAC)

Foerster et al. – AAAI 18

- Extends Decentralized REINFORCE to the Actor Critic case

Algorithm 4 Independent Actor-Critic (IAC) (finite-horizon)

Require: Individual actor models $\pi_i(a_i|h_i)$, parameterized by ψ_i

Require: Individual critic models $\hat{V}_i(h)$, parameterized by θ_i

```
1: for all episodes do
2:    $h_{i,0} \leftarrow \emptyset$                                 {Empty initial history}
3:   for  $t = 0$  to  $\mathcal{H} - 1$  do
4:     Choose  $a_{i,t}$  at  $h_{i,t}$  from  $\pi_i(a_i|h_{i,t})$ 
5:     See joint reward  $r_t$ , local observation  $o_{i,t}$           {Depends on joint action a}
6:      $h_{i,t+1} \leftarrow h_{i,t} a_{i,t} o_{i,t}$                   {Append new action and obs to previous history}
7:     Compute value TD error:  $\delta_{i,t} \leftarrow r_t + \gamma \hat{V}_i(h_{i,t+1}) - \hat{V}_i(h_{i,t})$     ← On-policy error
8:     Compute actor gradient estimate:  $\gamma^t \delta_{i,t} \nabla \log \pi_i(a_{i,t}|h_{i,t})$ 
9:     Update actor parameters  $\psi_i$  using gradient estimate (e.g.,  $\psi_i \leftarrow \psi_i + \alpha \gamma^t \delta_{i,t} \nabla \log \pi_i(a_{i,t}|h_{i,t})$ )
10:    Compute critic gradient estimate:  $\delta_{i,t} \nabla \hat{V}_i(h_{i,t})$ 
11:    Update critic parameters  $\theta_i$  using gradient estimate (e.g.,  $\theta_i \leftarrow \theta_i + \beta \gamma \delta_{i,t} \nabla \hat{V}_i(h_{i,t})$ )    ← Update both models
12:   end for
13: end for
```

Policy and value model

{Empty initial history}

{Depends on joint action a}

{Append new action and obs to previous history}

On-policy error

$\psi_i \leftarrow \psi_i + \alpha \gamma^t \delta_{i,t} \nabla \log \pi_i(a_{i,t}|h_{i,t})$

Update both models

Other decentralized PG methods

- Can extend any single-agent PG method to the multi-agent case
- Independent PPO (IPPO) (de Witt et al. 20)
 - A version of IAC with PPO as the base RL method
 - Yu et al. (22) version uses parameter sharing (not DTE)
 - More about IPPO and MAPPO in the CTDE discussion
- Not a very active area

Other topics

- **Parameter sharing**
 - Agents share the same copy of policy and/or value networks
 - I consider this a form of CTDE (since it assumes centralized info)
 - Decentralized methods can easily use parameter sharing to potentially improve performance
- **Relationship with CTDE**
 - Centralized PG equal to decentralized PG so maybe not that different?
- **Other forms of decentralization**
 - Communication during execution using ‘networked’ agents, e.g., (Zhang et al. 18)

Centralized Training for Decentralized Execution (CTDE) MARL

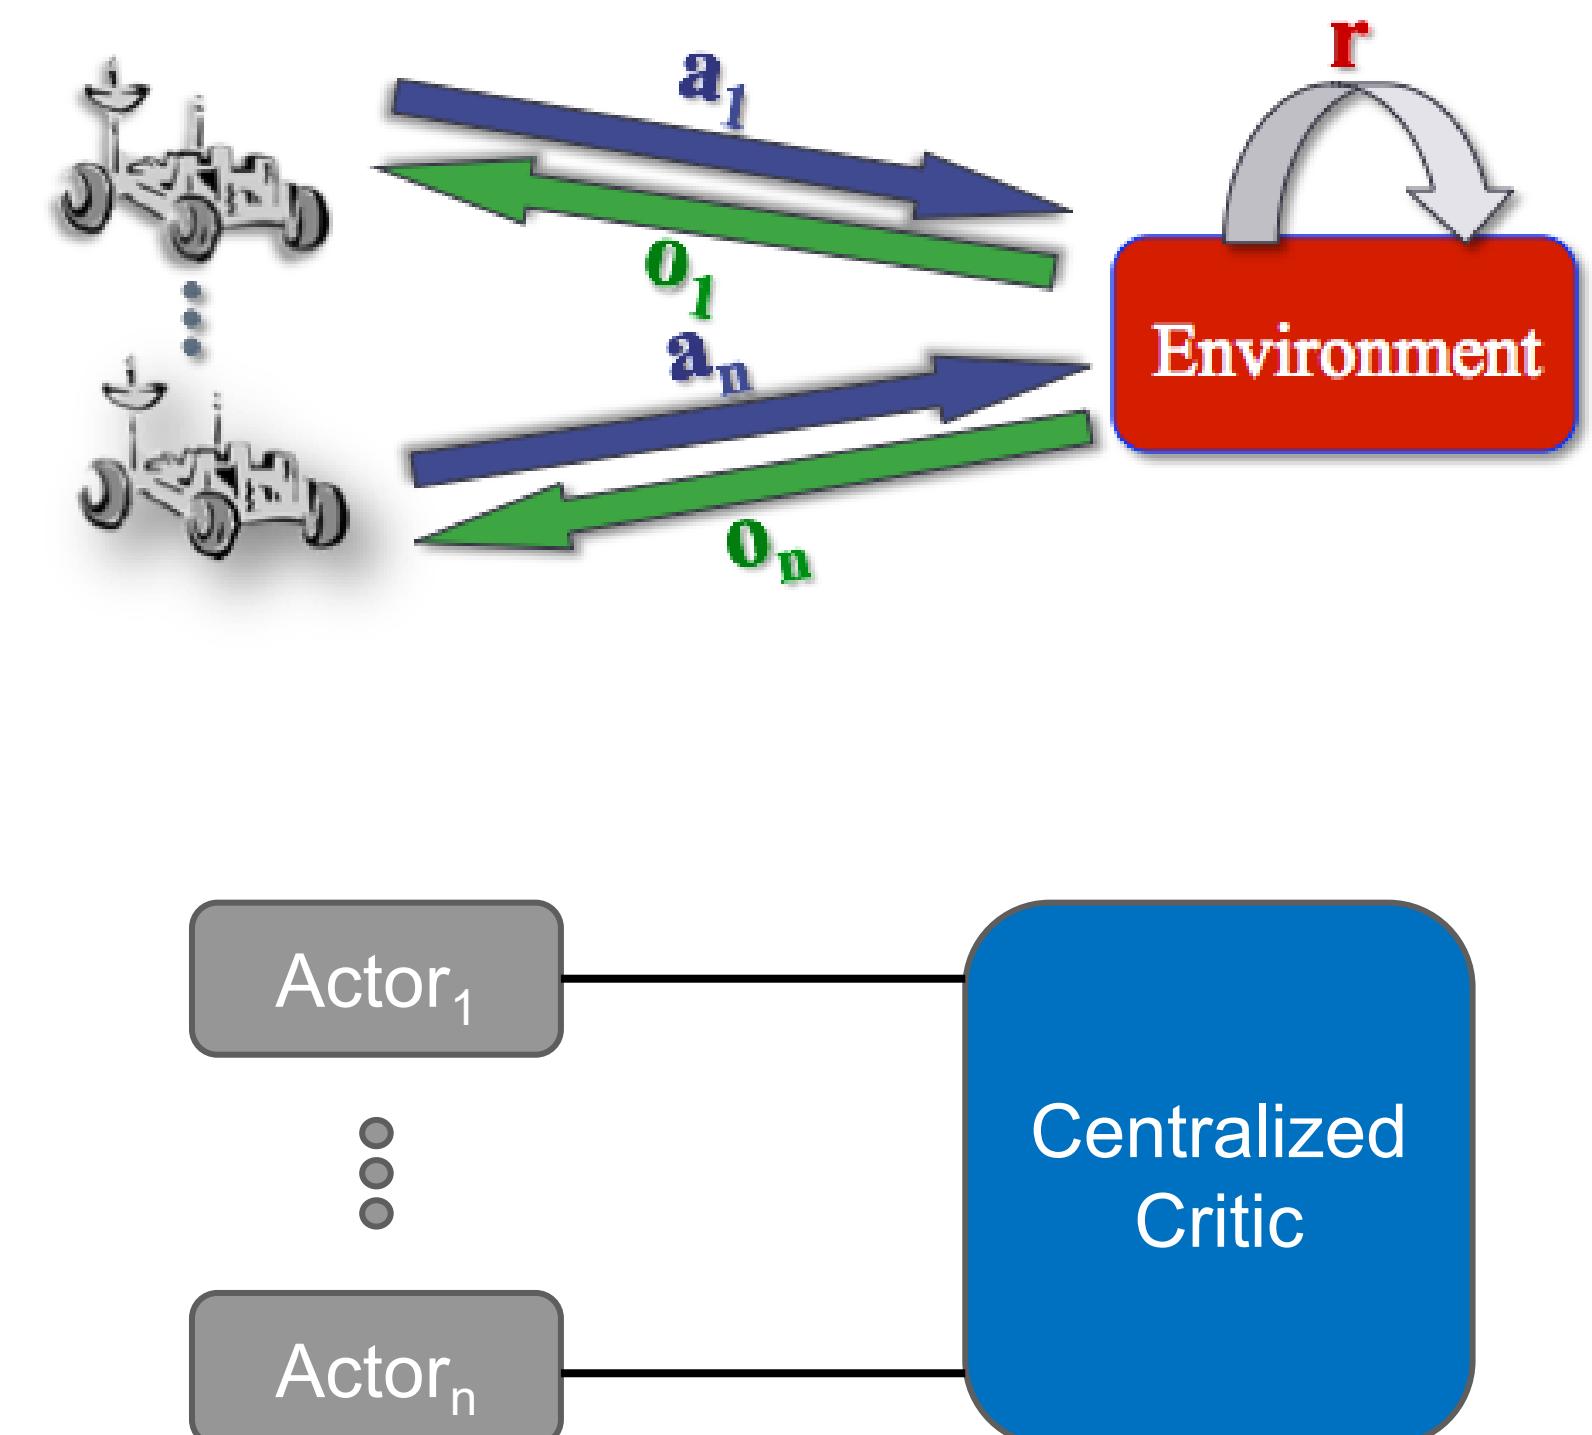
Models and methods

Centralized training for decentralized execution (CTDE)

Assumptions

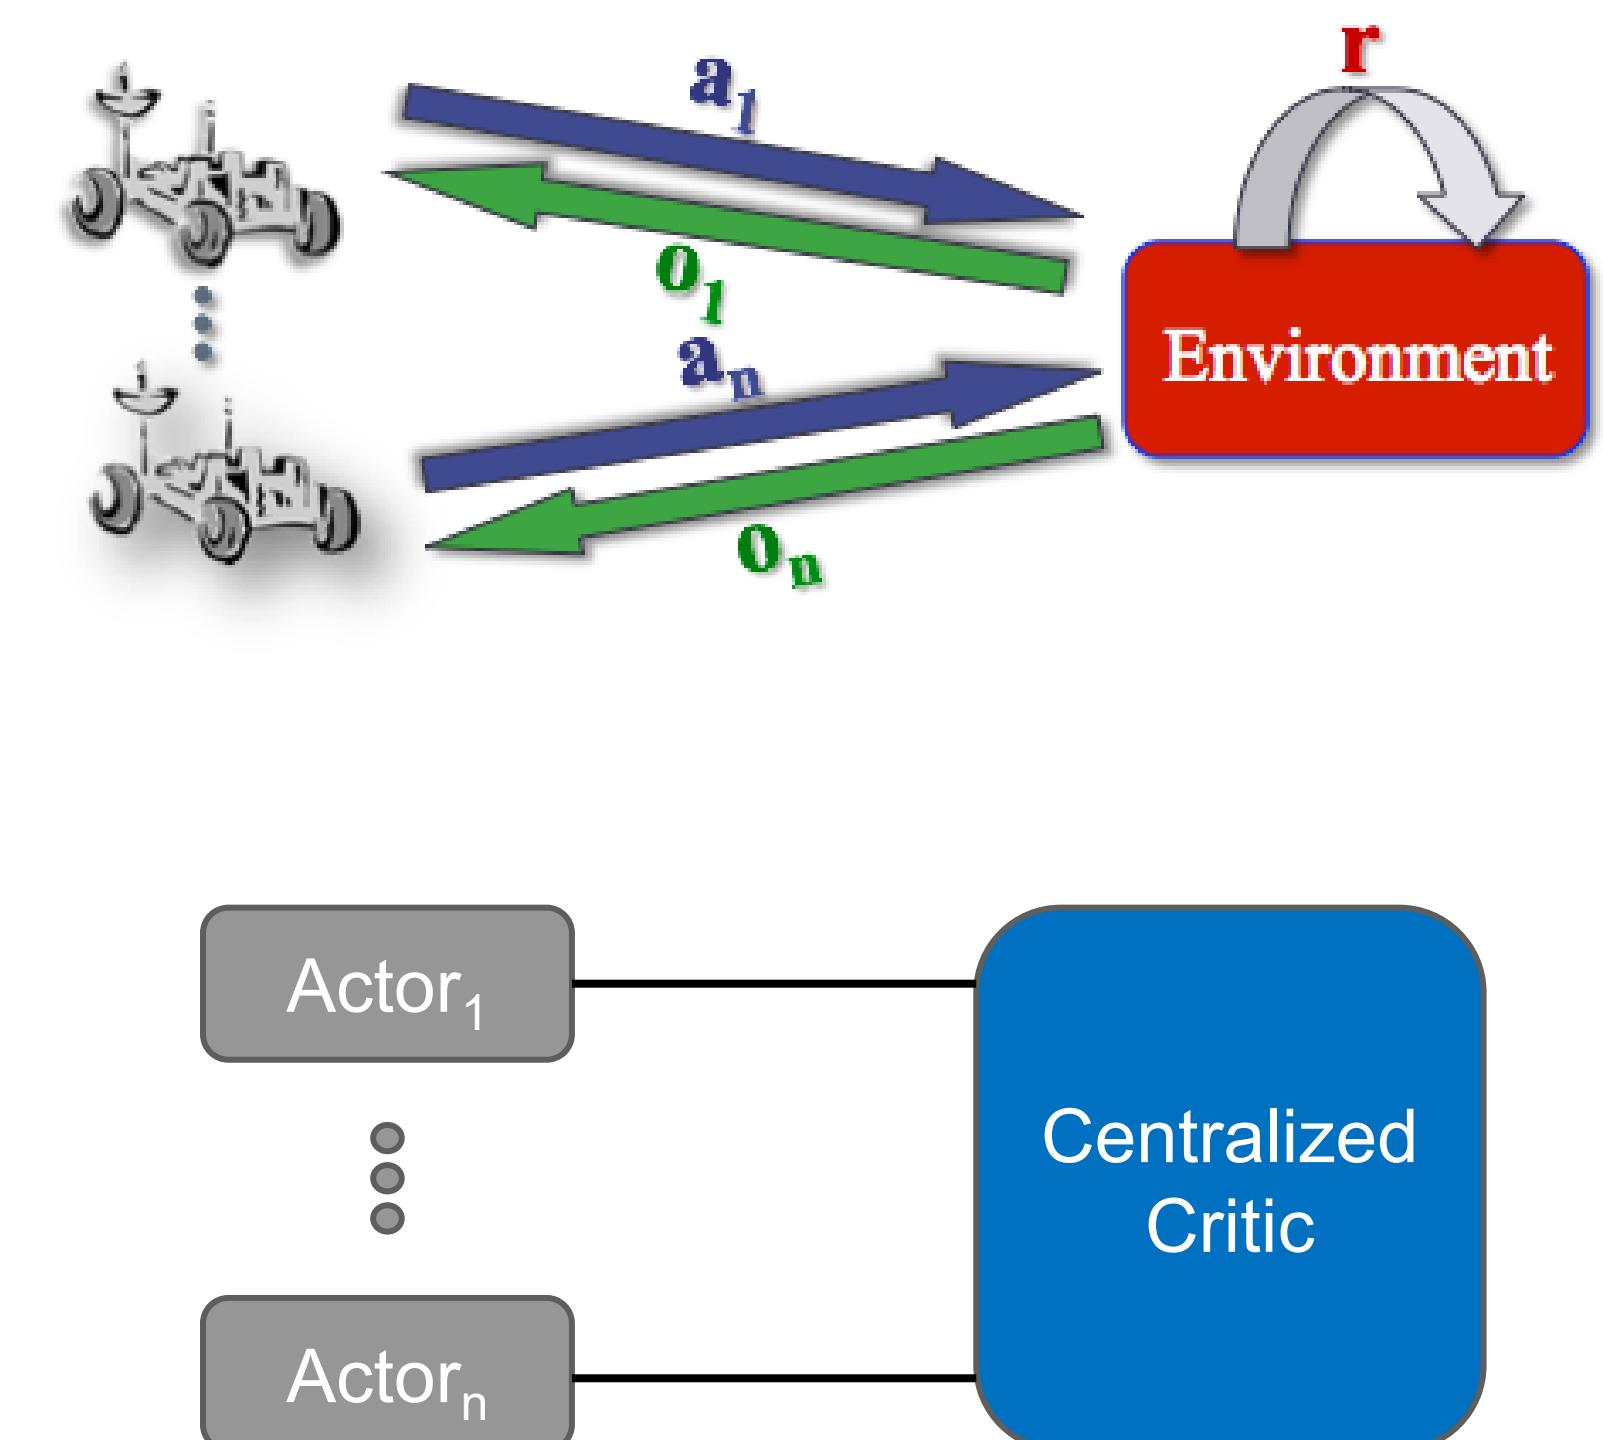
- each agent, i , observes its current observation, o_i , and takes action a_i at the resulting history, h_i , **like DTE**
- the (centralized) algorithm/controller observes joint information o and a and the joint reward r (and possibly other information such as the underlying state s) **like CTE**

By far the most common type of (cooperative) MARL



Centralized training for decentralized execution (CTDE)

- Train offline for online execution
- Can use centralized info offline
- Still need to execute in a decentralized manner
- CTDE has become the dominant form of (cooperative) MARL
- Many methods: MADDPG, NeurIPS-17; COMA AAAI-18; QMIX, ICML-18; QPLEX, ICML-21; MAPPO, NeurIPS DB-22

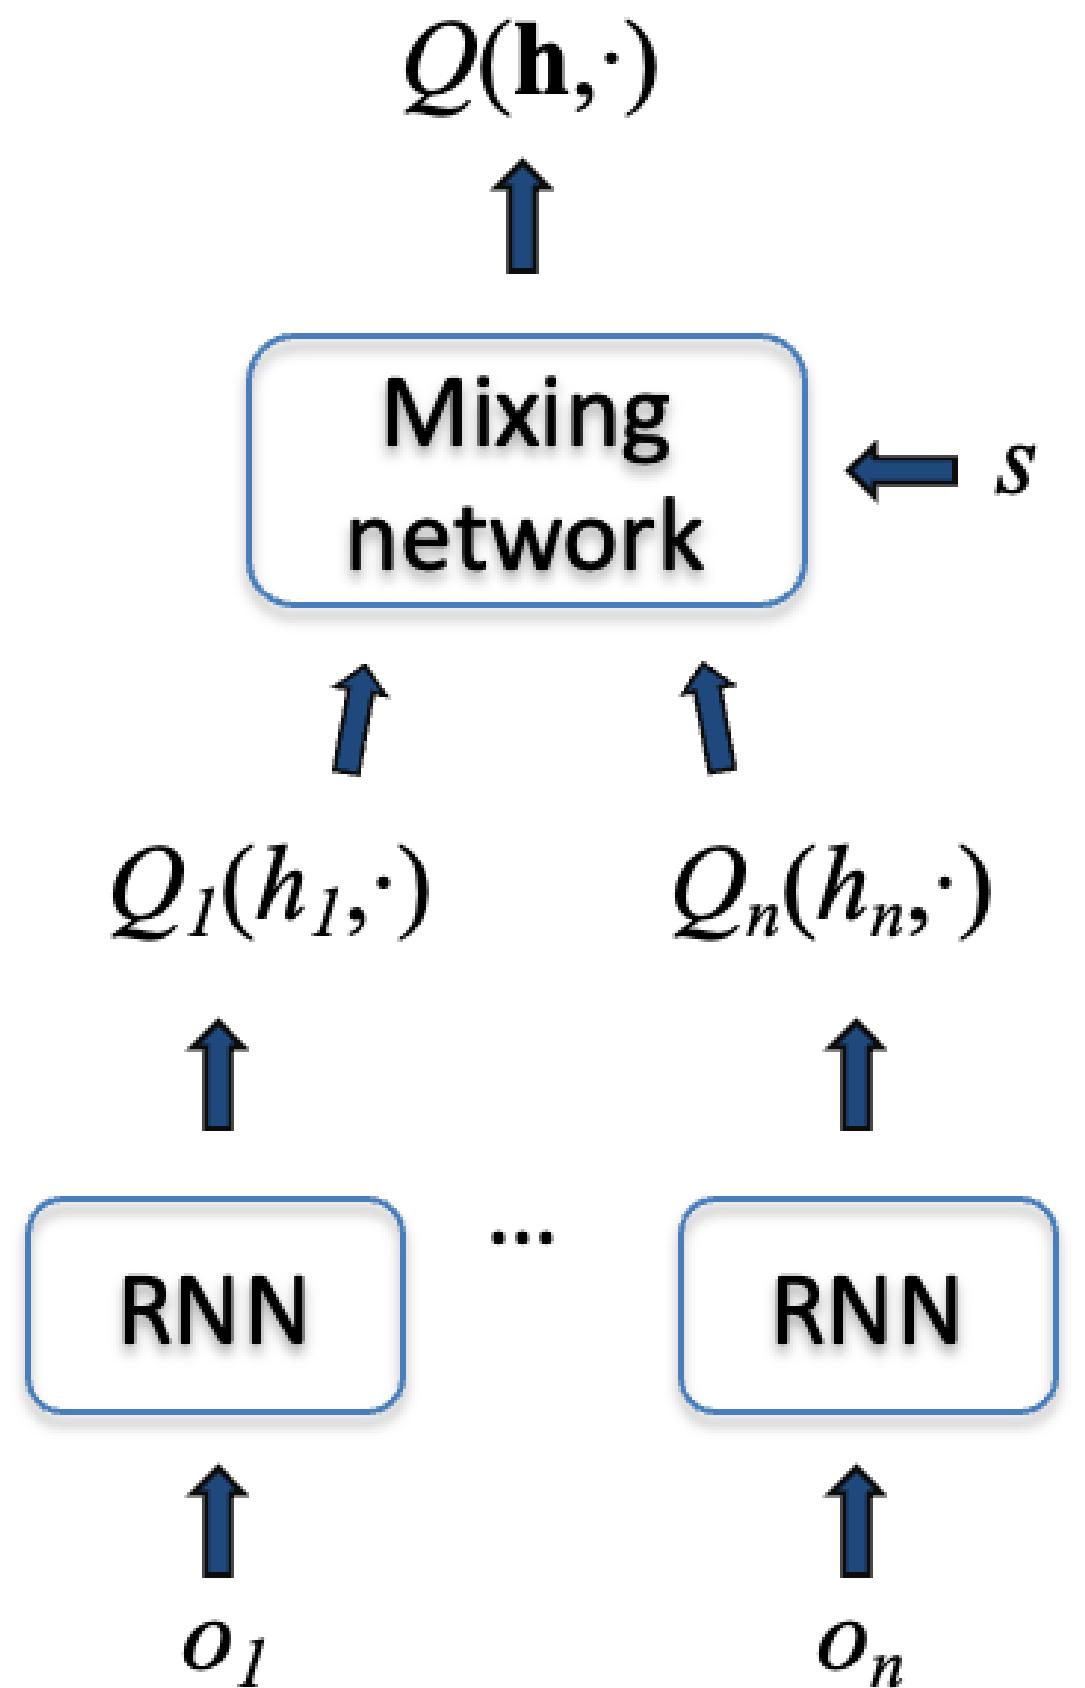


CTDE Action-Value Methods

Value function factorization: VDN, QMIX, and QPLEX

Value function factorization methods

- Basic idea:
 - Learn individual Q-values per agent as well as a form of joint Q-function
 - During training, learn individual Q-values from joint one
 - During execution, each agent uses individual Q-values to select actions



Value decomposition networks (VDN)

[Sunehag et al. – arXiv 17](#)

- The first deep value function factorization/decomposition method

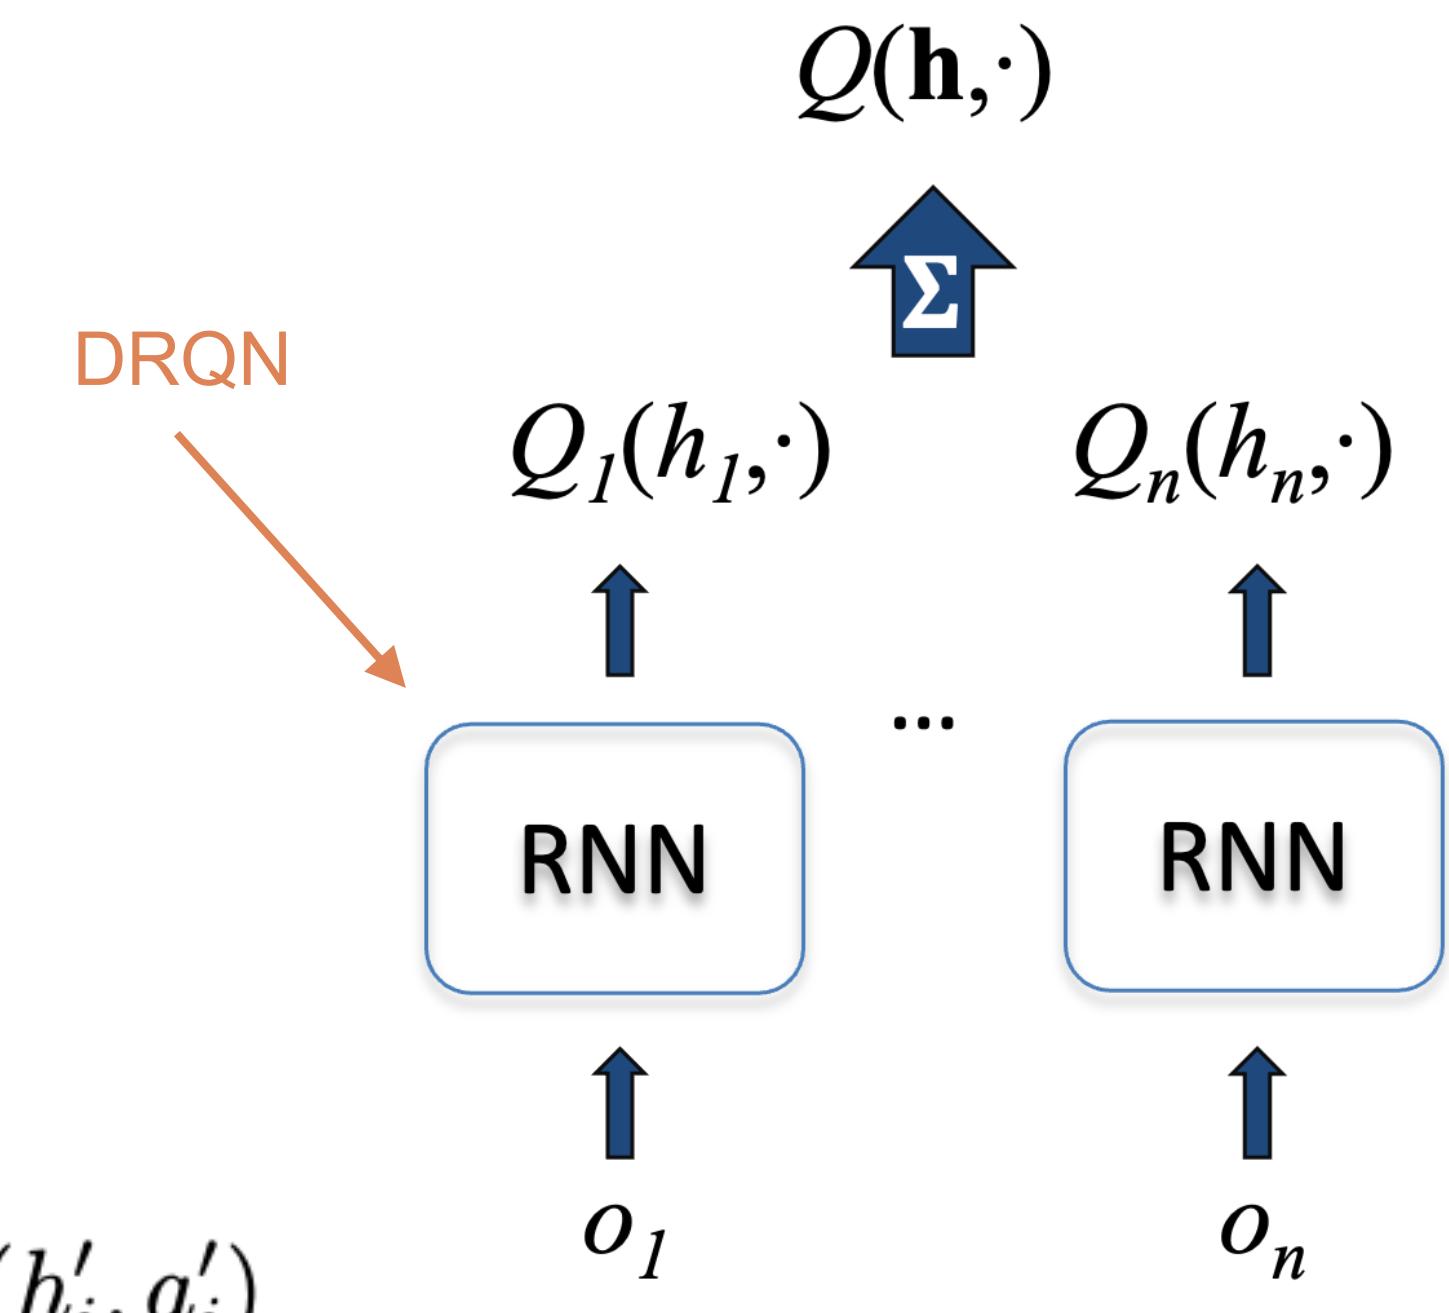
- Represents joint Q-value as a sum of individual Q-values:

$$Q(\mathbf{h}, \mathbf{a}) \approx \sum_{i \in \mathbb{I}}^n Q_i(h_i, a_i)$$

- Trains solely based on (joint) RL loss

$$\mathcal{L}(\theta) = \mathbb{E}_{\langle \mathbf{h}, \mathbf{a}, r, \mathbf{o} \rangle \sim \mathcal{D}} \left[(y - \sum_i^n Q_i^\theta(h_i, a_i))^2 \right], \text{ where } y = r + \gamma \sum_i^n \max_{a'_i} Q_i^{\theta^-}(h'_i, a'_i)$$

- Simple, scalable, but limited joint Q-value representation



QMIX

[Rashid et al. – ICML 18](#)

- Extends VDN to represent monotonic functions

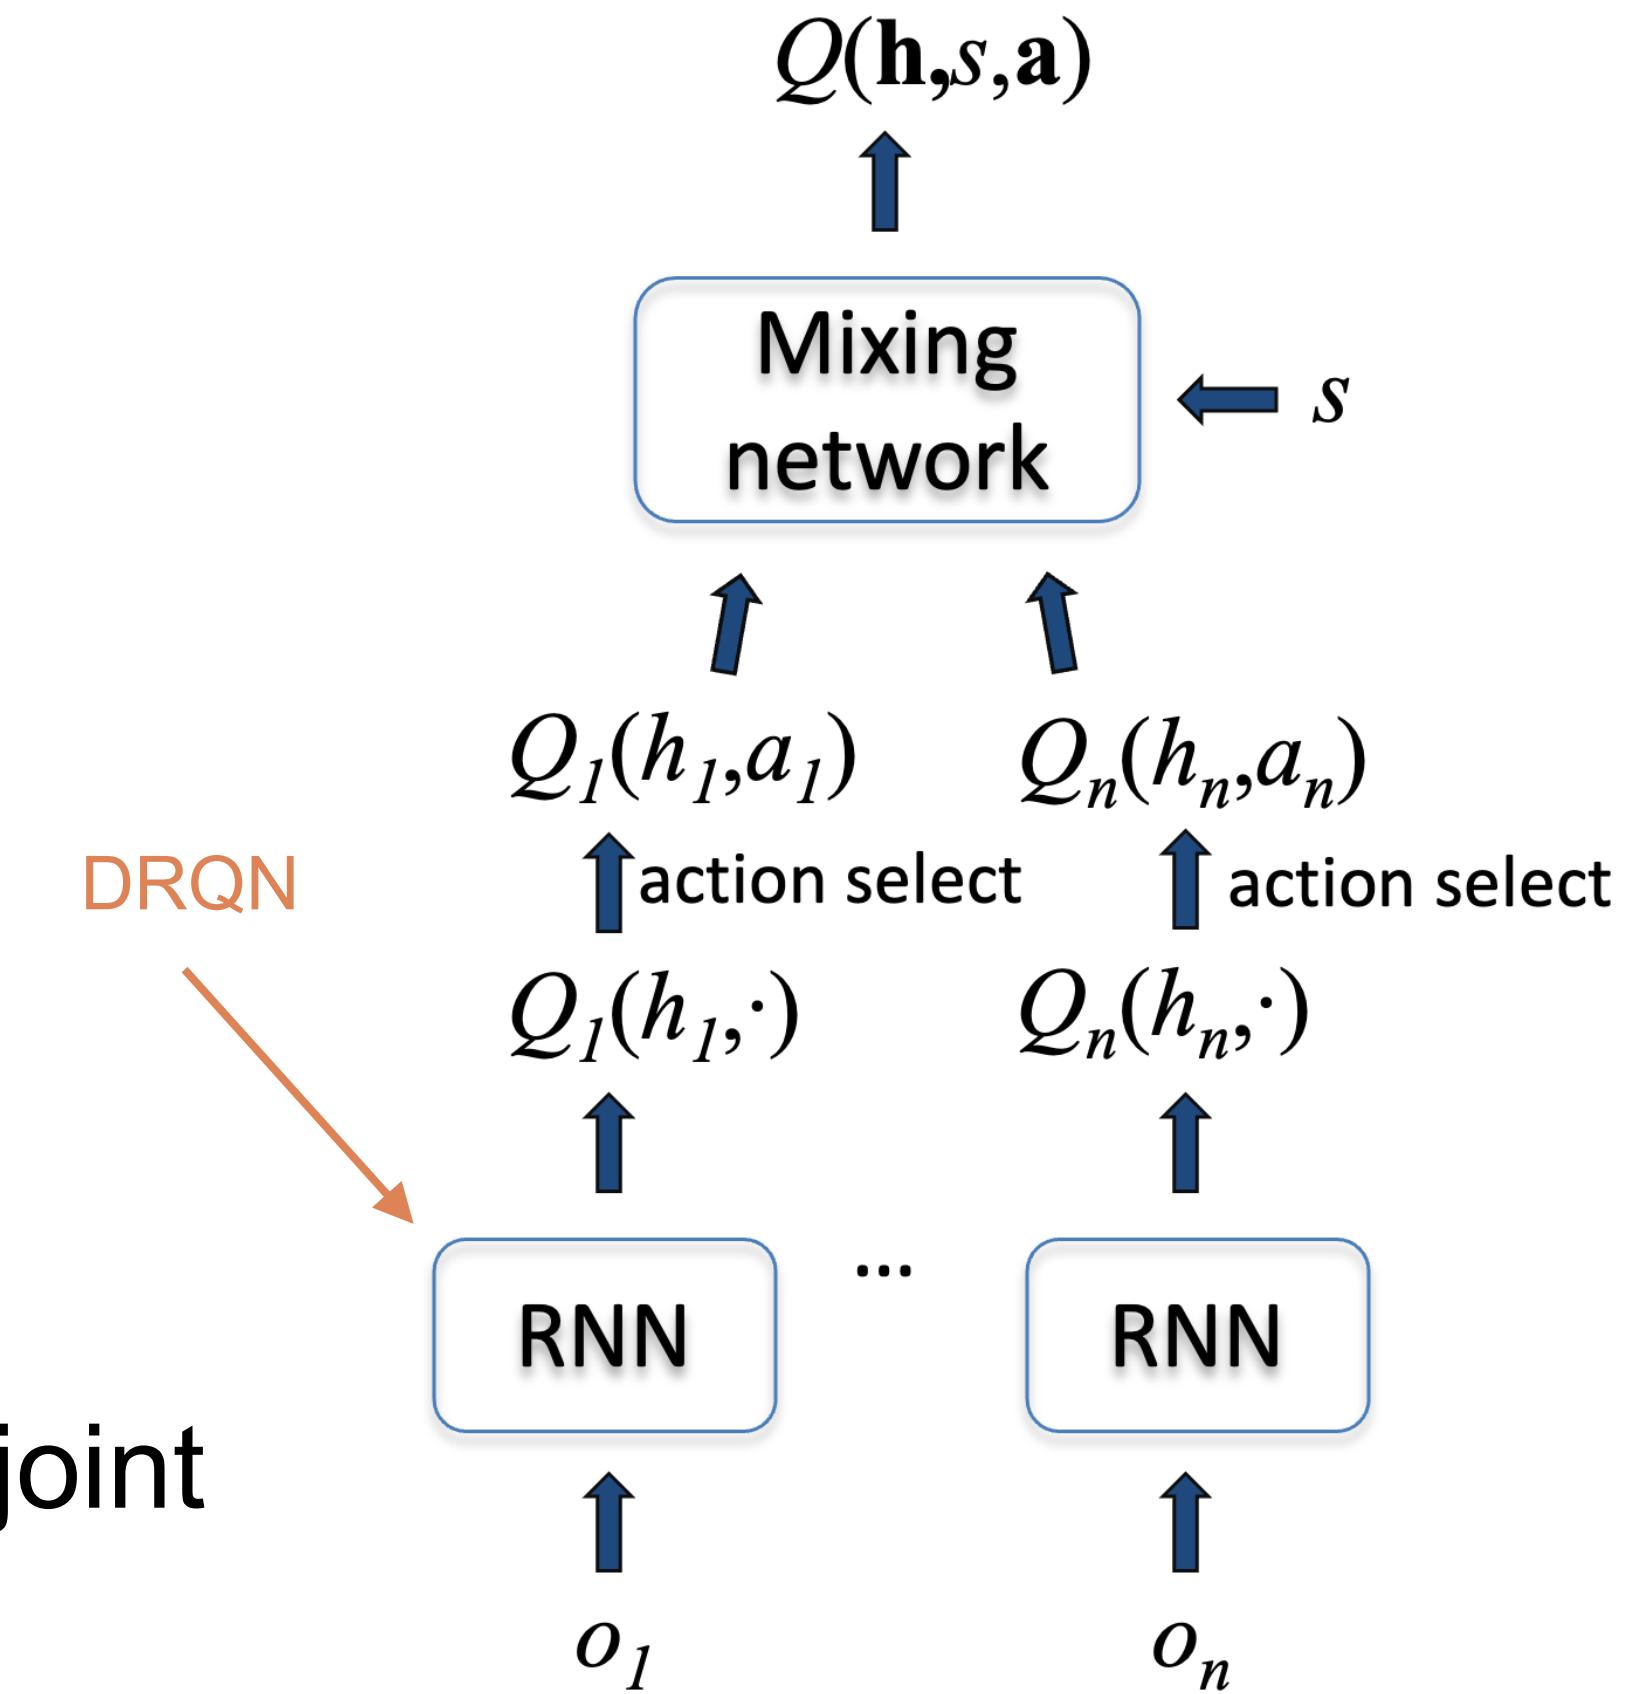
$$Q(\mathbf{h}, \mathbf{a}) \approx f_{mono}(Q_1(h_1, a_1), \dots, Q_n(h_n, a_n))$$

- (implemented with positive weights in mixer)
- Also, use state as input to mixer (with hypernetwork)
- Still argmax over indiv. Q-functions and train based on the joint loss

$$\mathcal{L}(\theta) = \mathbb{E}_{\langle \mathbf{h}, s, \mathbf{a}, r, \mathbf{o}, s' \rangle \sim \mathcal{D}} \left[(y - \mathbf{Q}^\theta(\mathbf{h}, s, \mathbf{a}))^2 \right], \text{ where } y = r + \gamma \mathbf{Q}^{\theta^-}(\mathbf{h}', s', \tilde{\mathbf{a}}'),$$

and $\tilde{\mathbf{a}}' = \langle \underset{a'_1}{\operatorname{argmax}} Q_1(h'_1, a'_1), \dots, \underset{a'_n}{\operatorname{argmax}} Q_n(h'_n, a'_n) \rangle$

- Can't represent all Q-functions but still a state-of-the-art method



Individual Global-Max (IGM)

[Son et al.– ICML 19](#) (QTRAN)

Definition: Individual-Global-Max

For a joint action-value function $\mathbf{Q}(\mathbf{h}, \mathbf{a})$ where $\mathbf{h} = \langle h_1, \dots, h_n \rangle$ is a joint action-observation history, if there exist individual functions $[Q_i]$ such that:

$$\underset{\mathbf{a}}{\operatorname{argmax}} \mathbf{Q}(\mathbf{h}, \mathbf{a}) = \begin{pmatrix} \operatorname{argmax}_{a_1} Q_1(h_1, a_1) \\ \vdots \\ \operatorname{argmax}_{a_n} Q_n(h_n, a_n) \end{pmatrix}$$

Then $[Q_i]$ satisfy IGM for \mathbf{Q} at \mathbf{h}

- This is the main principle of value factorization/decomposition methods: the argmax of the joint value function is the same as the argmax of the individual Q-functions
- VDN and QPLEX satisfy this (as do QTRAN, QPLEX, etc.)

QPLEX

[Wang et al.– ICLR 21](#)

Extends IGM to the advantage case

Definition: Advantage-based IGM

For joint and individual advantages:

$$\mathbf{A}(\mathbf{h}, \mathbf{a}) = \mathbf{Q}(\mathbf{h}, \mathbf{a}) - \mathbf{V}(\mathbf{h}) \text{ where } \mathbf{V}(\mathbf{h}) = \max_{\mathbf{a}} \mathbf{Q}(\mathbf{h}, \mathbf{a}) \text{ and } A_i(h_i, a_i) = Q_i(h_i, a_i) - V_i(h_i) \text{ where } V_i(h_i) = \max_{a_i} Q_i(h_i, a_i)$$

For a joint action-value function $\mathbf{Q}(\mathbf{h}, \mathbf{a})$ where $\mathbf{h} = \langle h_1, \dots, h_n \rangle$ is a joint action-observation history, if there exist individual functions $[Q_i]$ such that:

$$\operatorname{argmax}_{\mathbf{a}} \mathbf{A}(\mathbf{h}, \mathbf{a}) = \begin{pmatrix} \operatorname{argmax}_{a_1} A_1(h_1, a_1) \\ \vdots \\ \operatorname{argmax}_{a_n} A_n(h_n, a_n) \end{pmatrix}$$

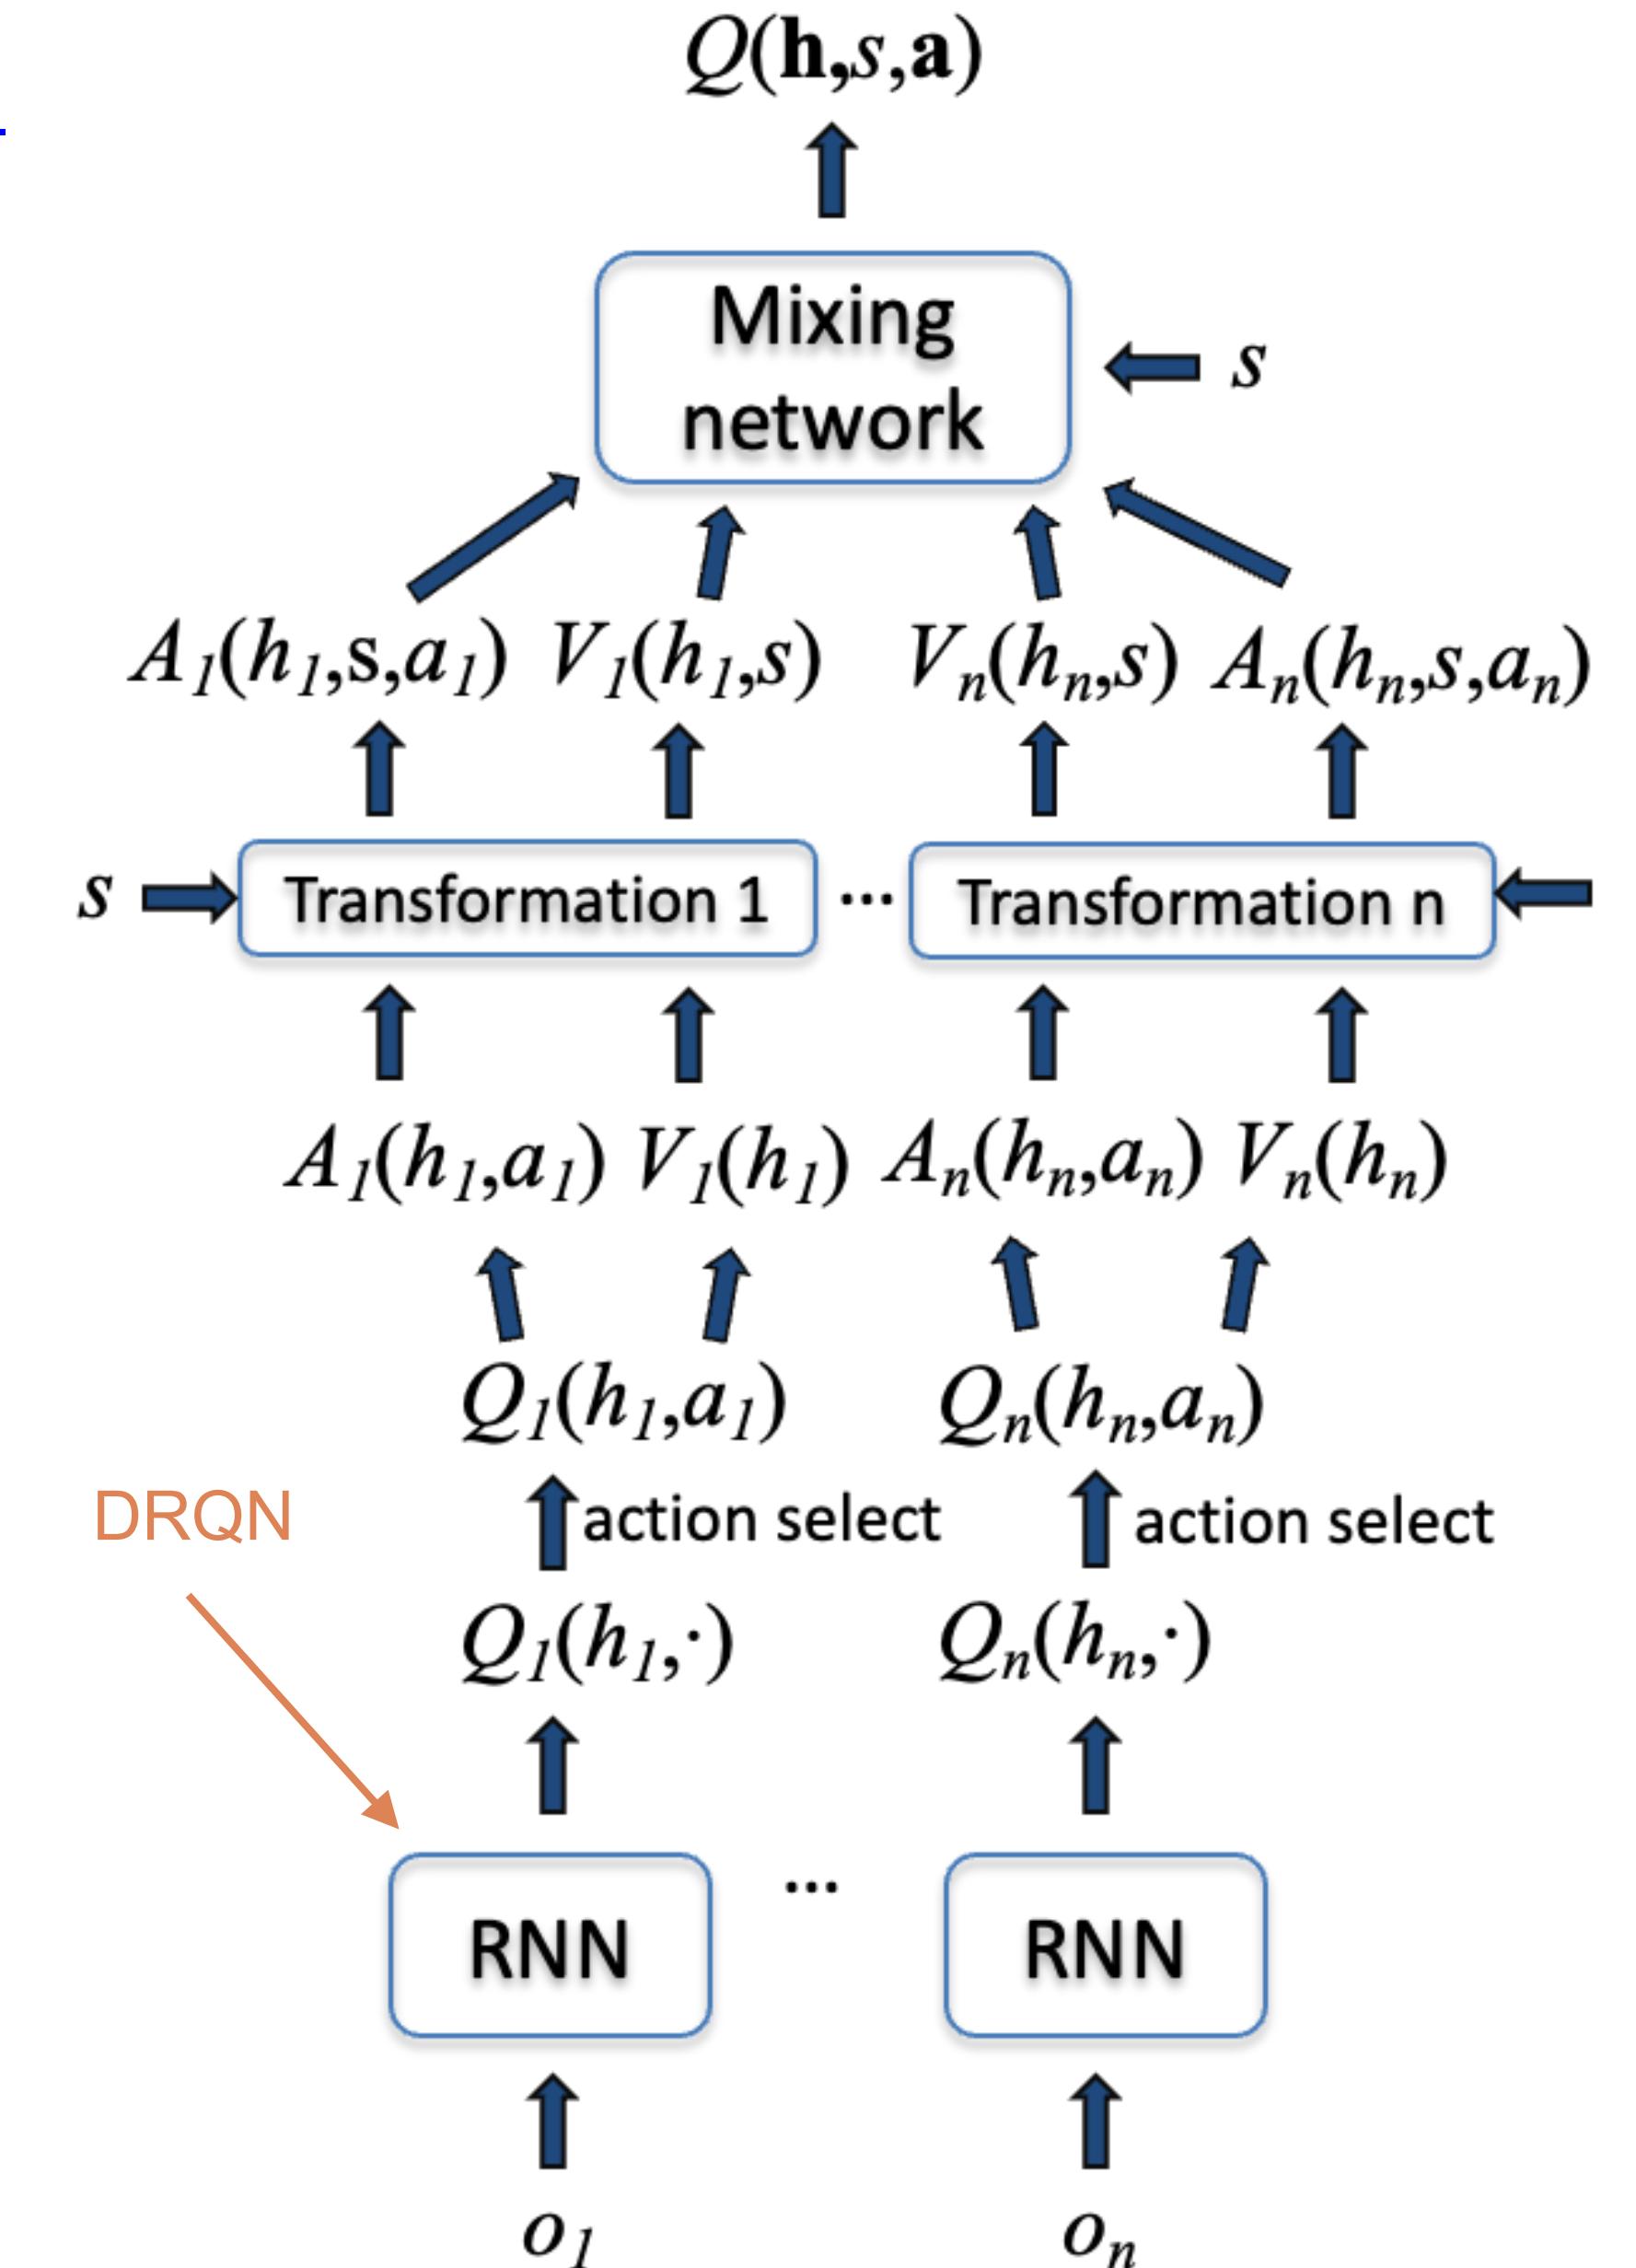
Then $[Q_i]$ satisfy IGM for \mathbf{Q} at \mathbf{h}

- This is subtle but important! Non-standard advantage makes them 0 for optimal action and negative otherwise! Used as a constraint to represent the full IGM function class

QPLEX architecture

[Wang et al.– ICLR 21](#)

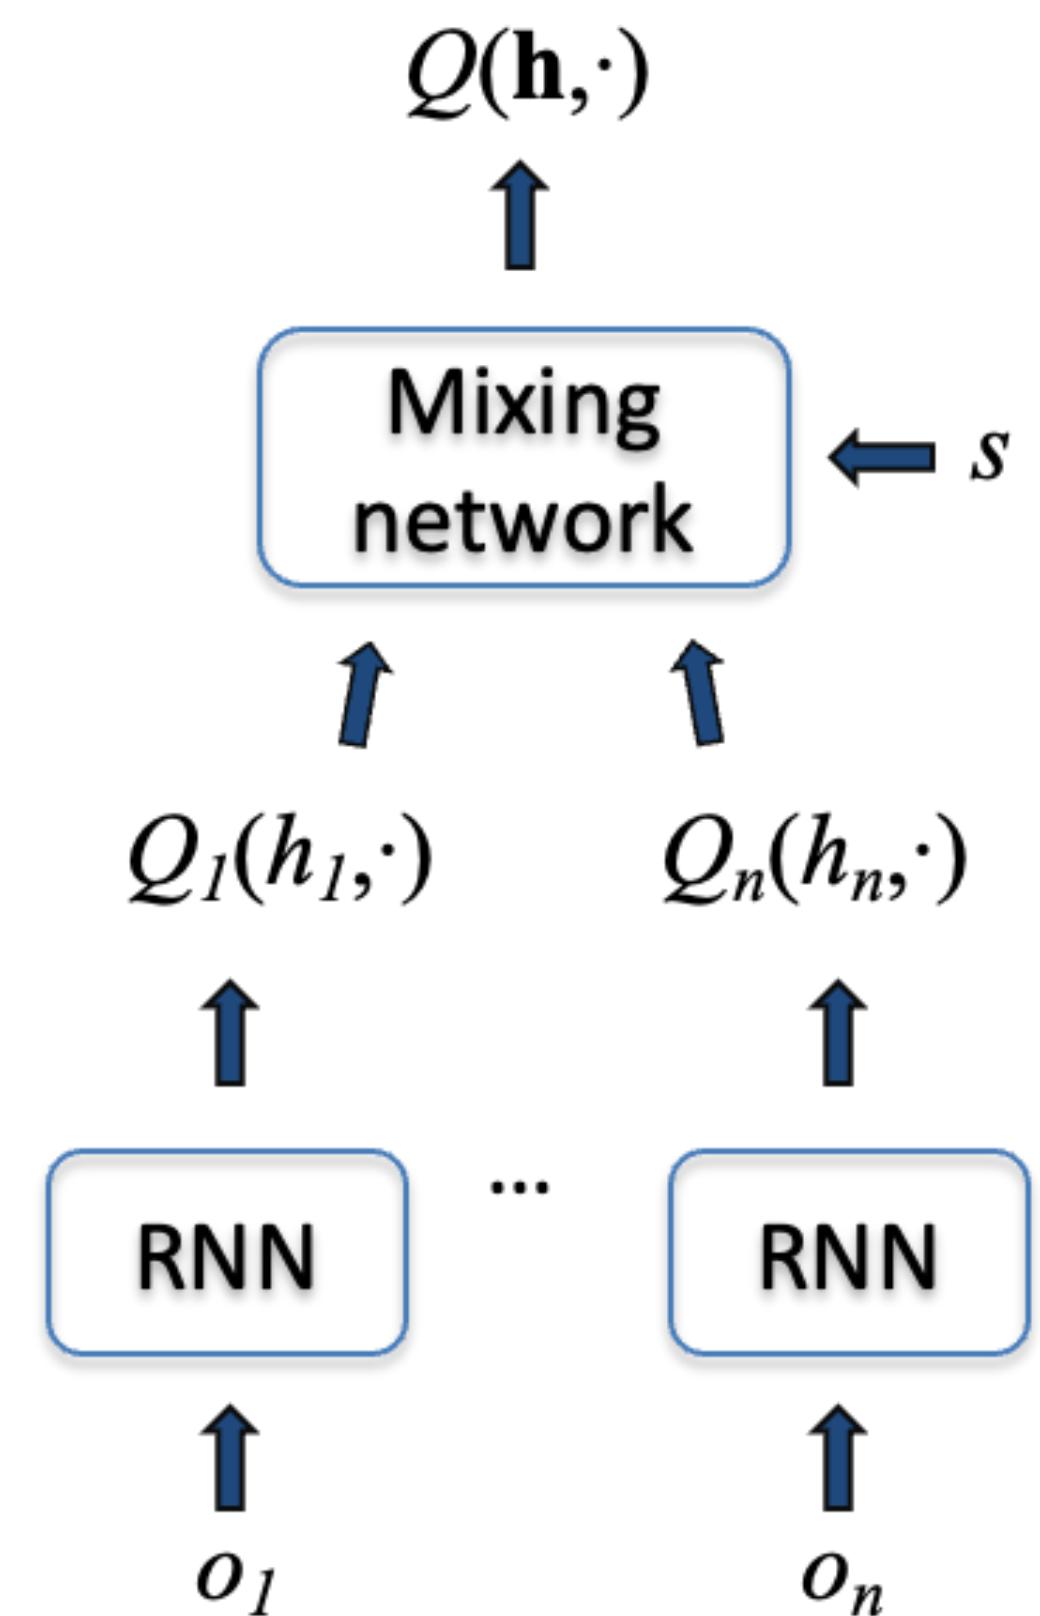
- Architecture is a bit complicated but it performs well
- Can sometimes outperform QMIX and is a state-of-the-art method
- Other recent value factorization/decomposition methods but not clear they outperform QMIX and QPLEX



State in value function factorization

Marchesini et al.,--AAMAS 25

- Is it cheating/wrong to use state during training?
- **QMIX**: Sound since state information gets marginalized out
- **QPLEX**:
 - Sound since similar to QMIX
 - Less general with state (can't represent all IGM functions)
- **Weighted QMIX**: Probably not sound as uses separate state-conditioned weights



Note: The paper also introduces a new algorithm DualMIX which I don't discuss here

State in value function factorization

Marchesini et al.,--AAMAS 25

Why is the state helpful?

Benefit of state unclear in theory but may be helpful in practice

Tried the methods with state (s), a random (r) value, or a 0 value

Other information can outperform state info!

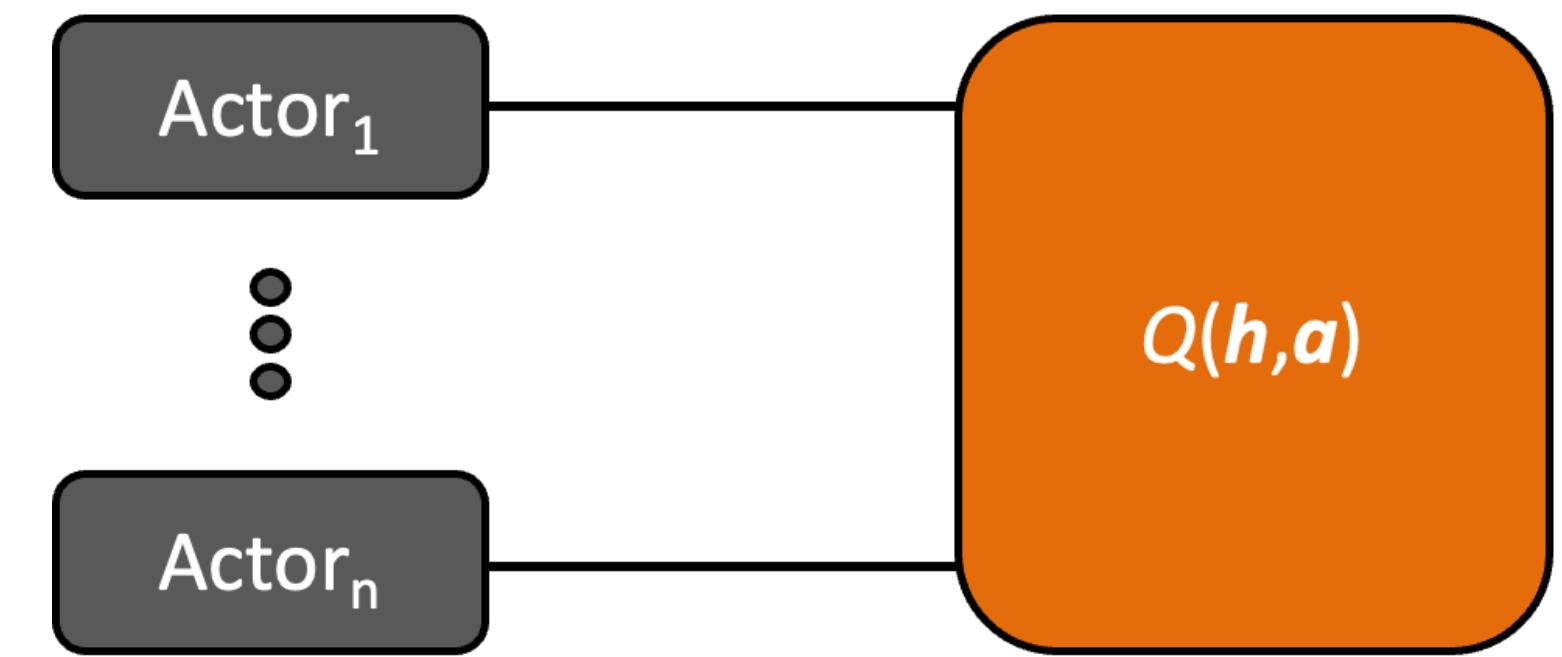
	(fine-tuned ↓)	5s10z
QMIX		
<i>s</i>	15.8 ± 0.4	
<i>r</i>	14.5 ± 1.4	
<i>c</i>	14.7 ± 0.1	
QPLEX		
<i>s</i>	16.2 ± 2.1	
<i>r</i>	18.0 ± 0.6	
<i>c</i>	18.3 ± 0.8	

CTDE Policy Gradient Methods

Centralized critics: MADDPG, COMA, and MAPPO

Actor critic with a centralized critic

- Have an actor for each agent
- Learn a 'centralized' Q-function
- Update each actor using this joint Q-value:



$$\nabla_{\psi_i} J = \mathbb{E}_{\langle \mathbf{h}, \mathbf{a} \rangle \sim \mathcal{D}} [\mathbf{Q}^\pi(\mathbf{h}, \mathbf{a}) \nabla_{\psi_i} \log \pi_i(a_i | h_i)]$$

- Update the joint Q-value using the joint info:

$$\mathcal{L}(\theta) = \mathbb{E}_{\langle \mathbf{h}, \mathbf{a}, r, \mathbf{h}' \rangle \sim \mathcal{D}} \left[(y - \hat{\mathbf{Q}}(\mathbf{h}, \mathbf{a}))^2 \right], \text{ where } y = r + \gamma \hat{\mathbf{Q}}(\mathbf{h}', \mathbf{a}')$$

A basic centralized critic approach

A policy network for each agent
 A joint value network
 Joint error calculation
 The gradient using
 $\mathcal{L}(\theta) = \mathbb{E}_{\langle h, a, r, h' \rangle \sim \mathcal{D}} [(y - \hat{Q}(h, a))^2]$, where $y = r + \gamma \hat{Q}(h', a')$
 Loop over agents
 Use joint Q to update agent policies

Algorithm 6 Independent Actor Centralized Critic (IACC) (finite-horizon)

```

1: Initialize individual actor models  $\pi_i(a_i|h_i)$ , parameterized by  $\psi_i$ 
2: Initialize centralized critic model  $\hat{Q}(h, a)$ , parameterized by  $\theta$ 
3: for all episodes do
4:    $h_{i,0} \leftarrow \emptyset$                                      {Empty initial history}
5:   Denote  $h_t$  as  $\langle h_{1,0}, \dots, h_{n,0} \rangle$            {Notation for joint variables}
6:   for all  $i$ , choose  $a_{i,0}$  at  $h_{i,0}$  from  $\pi_i(a_i|h_{i,0})$ 
7:   Store  $a_t$  as  $\langle a_{1,0}, \dots, a_{n,0} \rangle$ 
8:   for  $t = 0$  to  $\mathcal{H} - 1$  do
9:     Take joint action  $a_t$ , see joint reward  $r_t$ , and observations  $o_t$ 
10:    for all  $i$ ,  $h_{i,t+1} \leftarrow h_{i,t} a_{i,t} o_{i,t}$           {Append new action and obs to previous history}
11:    for all  $i$ , choose  $a_{i,t+1}$  at  $h_{i,t+1}$  from  $\pi_i(a_i|h_{i,t+1})$ 
12:    Store  $a_{t+1}$  as  $\langle a_{1,t+1}, \dots, a_{n,t+1} \rangle$ 
13:     $\delta_t \leftarrow r_t + \gamma \hat{Q}(h_{t+1}, a_{t+1}) - \hat{Q}(h_t, a_t)$       {Compute centralized TD error}
14:    Compute critic gradient estimate:  $\delta_t \nabla_\theta \hat{Q}(h_t, a_t)$ 
15:    Update critic parameters  $\theta$  using gradient estimate (e.g.,  $\theta \leftarrow \theta + \beta \delta_t \nabla_\theta \hat{Q}(h_t, a_t)$  for learning rate  $\beta$ )
16:    for each agent  $i$  do
17:      Compute actor gradient estimate:  $\gamma^t \hat{Q}(h_t, a_t) \nabla_{\psi_i} \log \pi_i(a_{i,t}|h_{i,t})$ 
18:      Update actor parameters  $\psi_i$  using gradient estimate (e.g.,  $\psi_i \leftarrow \psi_i + \alpha \gamma^t \hat{Q}(h_t, a_t) \nabla_{\psi_i} \log \pi_i(a_{i,t}|h_{i,t})$  for learning rate  $\alpha$ )
19:    end for
20:  end for
21: end for

```

MADDPG

[Lowe et al.—NeurIPS 17](#)

- Designed for competitive or cooperative problems
- Off-policy (so uses reply buffer like DQN)
- Continuous action, so uses a Deterministic PG (Silver et al., ICML-14)

$$\nabla_{\psi_i} J = \mathbb{E}_{x, \mathbf{a} \sim \mathcal{D}} [\nabla_{\psi_i} \mu_i(o_i) \nabla_{\mathbf{a}} \mathbf{Q}^{\pi}(x, \mathbf{a}) \mid_{a_i = \mu_i(o_i)}]$$

- Defined policies based on a single observation but should be:
$$\nabla_{\psi_i} J = \mathbb{E}_{x, \mathbf{a} \sim \mathcal{D}} [\nabla_{\psi_i} \mu_i(h_i) \nabla_{\mathbf{a}} \mathbf{Q}^{\pi}(\mathbf{h}, \mathbf{a}) \mid_{a_i = \mu_i(h_i)}]$$
- Learn centralized critic from the reply buffer and using target network θ^-
$$\mathcal{L}(\theta) = \mathbb{E}_{\langle \mathbf{h}, \mathbf{a}, r, \mathbf{h}' \rangle \sim \mathcal{D}} [(y - Q_{\theta}(\mathbf{h}, \mathbf{a}))^2], \text{ where } y = r + \gamma Q_{\theta^-}(\mathbf{h}', \mathbf{a}') \mid_{a_i = \mu^-(h_i)} \forall i \in \mathbb{I}$$
- MADDPG is no longer widely used but the centralized critic have been adopted

Note: For the cooperative CTDE case we assume a single shared critic among agents, do not consider learning policy models of the other agents, and do not consider ensembles of other agent policies to improve robustness.

Counterfactual Multi-Agent Policy Gradients (COMA)

[Foerster et al.–AAAI 18](#)

- Centralized critic along with a counterfactual baseline to potentially help with variance and credit assignment
- Calculate a per-agent advantage considering that difference between with the agent did and the expected Q-value from policy and fixing other agents:

$$A_i(\mathbf{h}, \mathbf{a}) = \mathbf{Q}(\mathbf{h}, \mathbf{a}) - \sum_{a'_i} \pi_i(a'_i | h_i) \mathbf{Q}(\mathbf{h}, a'_i, \mathbf{a}_{-i})$$

- Is implemented with agent ids to only require a single centralized critic network (rather than one per agent)
- On-policy so the critic is updated as usual: $\mathcal{L}(\theta) = \mathbb{E}_{\langle \mathbf{h}, \mathbf{a}, r, \mathbf{h}' \rangle \sim \mathcal{D}} [(y - \hat{\mathbf{Q}}(\mathbf{h}, \mathbf{a}))^2]$, where $y = r + \gamma \hat{\mathbf{Q}}(\mathbf{h}', \mathbf{a}')$
- Policy network update uses A_i instead of Q : $\gamma^t A_i(\mathbf{h}_t, \mathbf{a}_t) \nabla_{\psi_i} \log \pi_i(a_{i,t} | h_{i,t})$
- COMA is also not widely used but very influential

Note: COMA originally used state instead of history in the advantage and Q-values but this is incorrect as I'll discuss later.

MAPPO

[Yu et al. -- NeurIPS DB&B 22](#)

- MAPPO is a form of a centralized critic method
- Just use PPO as the base RL method
- Actor loss: $\mathcal{L}_{clip}^{MAPPO}(\psi_i) = \min \left(r_{\psi_i,i} \mathbf{A}, \text{clip}(r_{\psi_i,i}, 1 - \epsilon, 1 + \epsilon) \mathbf{A} \right)$
 - Uses joint advantage: $\mathbf{A}(\mathbf{h}, \mathbf{a}) = \mathbf{Q}(\mathbf{h}, \mathbf{a}) - \mathbf{V}(\mathbf{h})$
 - Use GAE but can be computed from \mathbf{V} as $\delta = r_t + \gamma \hat{\mathbf{V}}(\mathbf{h}_{t+1}) - \hat{\mathbf{V}}(\mathbf{h}_t)$
 - Uses joint value function and local policy ratio: $r_{\psi_i,i} = \frac{\pi_{\psi_i}(a_i|h_i)}{\pi_{\psi_i,old}(a_i|h_i)}$
- Critic loss: $\mathcal{L}^{MAPPO}(\theta) = \max \left[(\mathbf{V}(\mathbf{h}_t) - \hat{R}_t)^2, \left(\text{clip}(\mathbf{V}(\mathbf{h}), \mathbf{V}_{old}(\mathbf{h}) - \epsilon, \mathbf{V}_{old}(\mathbf{h}) + \epsilon) - \hat{R}_t \right)^2 \right]$
- Can use other centralized info in the critic (more later)
- Simple, but works well and some form of this often works best

IPPO

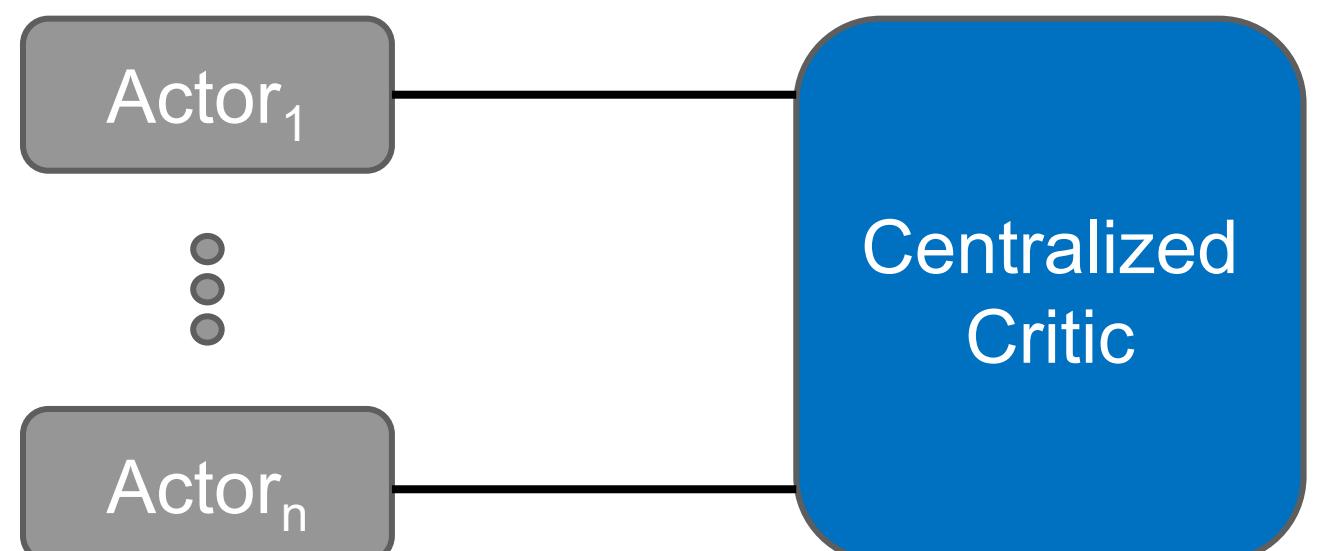
[de Witt et al. –arXiv 20](#)

- Actor loss: $\mathcal{L}_{clip}^{IPPO}(\psi_i) = \min \left(r_{\psi_i,i} A_i, \text{clip}(r_{\psi_i,i}, 1 - \epsilon, 1 + \epsilon) A_i \right)$
 - Uses local advantage: $\hat{A}_i = r_t + \gamma \hat{V}_i(h_{i,t+1}) - \hat{V}_i(h_{i,t})$
 - Can also use GAE or other methods (e.g., n-step)
 - Ratio same as before: $r_{\psi_i,i} = \frac{\pi_{\psi_i}(a_i|h_i)}{\pi_{\psi_i,old}(a_i|h_i)}$
 - The only difference is the use of A_i instead of A
- Critic loss (with clipping):
$$\mathcal{L}^{IPPO}(\theta) = \max \left[(V_i(h_{i,t})) - \hat{R}_t)^2, \left(\text{clip}(V_i(h_{i,t})), V_{i,old}(h_{i,t}) - \epsilon, V_{i,old}(h_{i,t}) + \epsilon \right) - \hat{R}_t \right]^2$$
- Often performs similarly to MAPPO but sometimes lower

Contrasting Centralized and Decentralized Critics in Multi-Agent Policy Gradient

[Lyu, Xiao, Daley and Amato – AAMAS21 Best Paper Nomination](#)

- Centralized critic widely used but misunderstood
- We show in theory:
 - Centralized Critic does not foster cooperation any better than Decentralized Critics
 - Both unbiased estimates of the decentralized policy
 - Centralized Critic exhibits more variance in policy gradient
- In practice:
 - Centralized Critic – less bias, more variance
 - Decentralized Critics – more bias, less variance

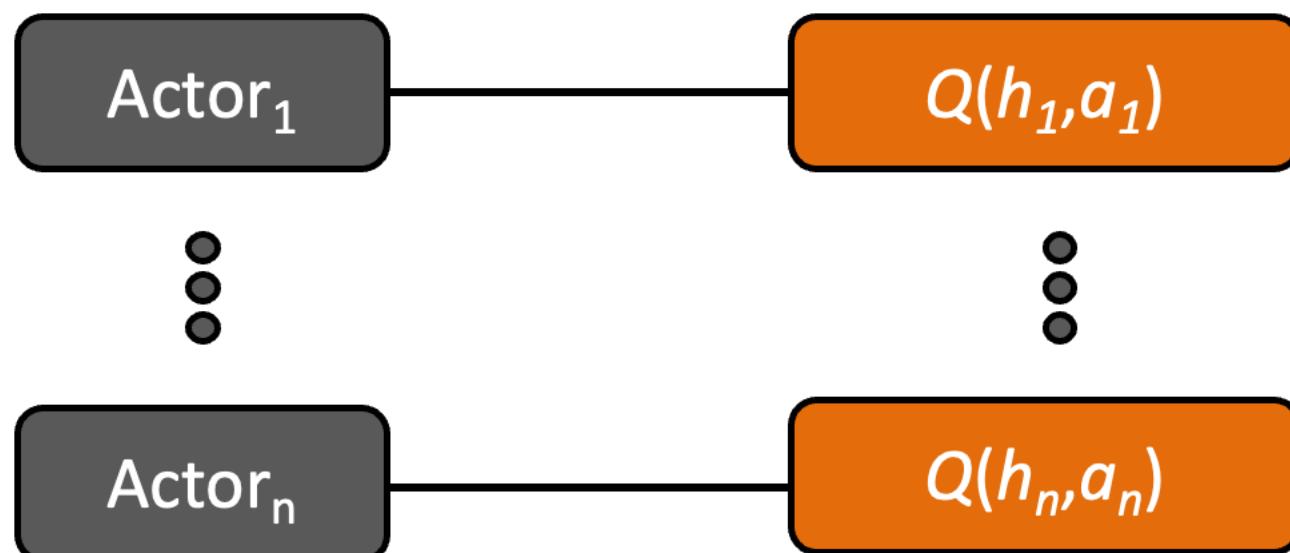


Multi-Agent Actor Critic

Decentralized and Centralized Critic

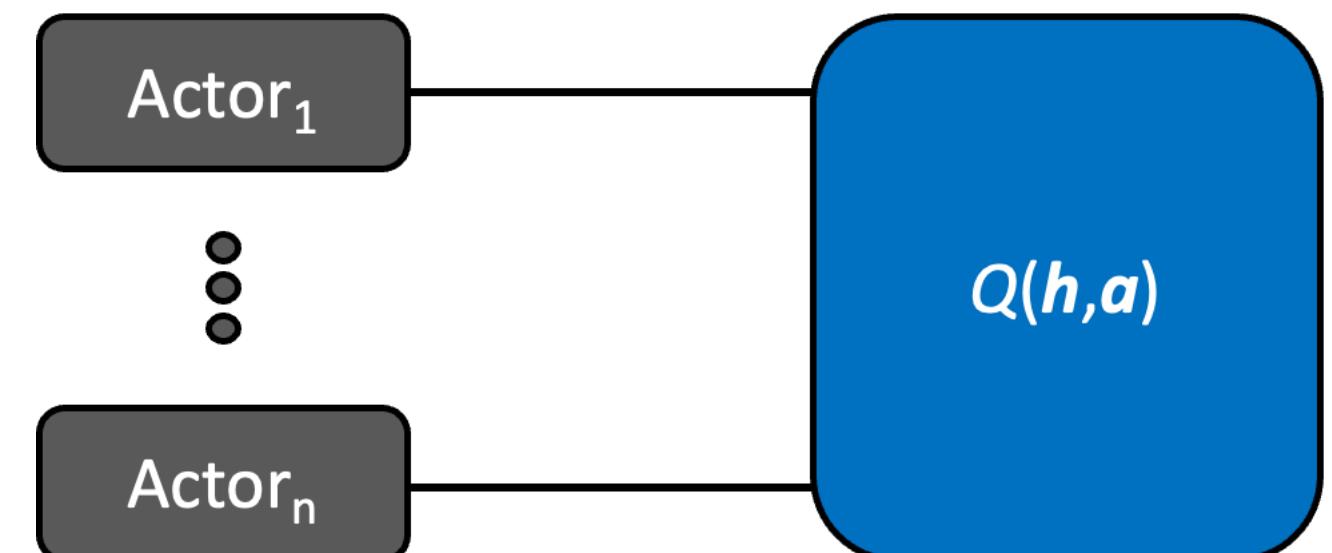
```
Initialize  $\theta, \phi$ 
for each training rollout  $e$  do
    Empty and fill buffer with experience data using actors  $\pi$ 
    for Each batch  $t$  do
        Unroll RNN using observations, actions and rewards
        for each agent  $i$  do
            Calculate TD targets  $y_t^i$ 
             $\phi_i = \phi_i - \alpha \nabla_{\phi^i} (y_t^i - Q^i(h_t^i, \mathbf{a}))^2$  // update critic weights
             $\theta^i = \theta^i + \alpha \nabla_{\theta^i} \log \pi^i(a | h_t^i) Q^i(h_t^i, a_t^i)$  // update actor weights
        end for
    end for
end for
```

Decentralized actor and critic: pretend the other agents are part of the environment (independent per agent)



```
Initialize  $\theta, \phi$ 
for each training rollout  $e$  do
    Empty and fill buffer with experience data using actors  $\pi$ 
    for Each batch  $t$  do
        Unroll RNN using observations, actions and rewards
        Calculate TD targets  $y_t$ 
         $\phi = \phi - \alpha \nabla_{\phi} (y_t - \mathbf{Q}(\mathbf{h}_t, \mathbf{a}_t))^2$  // update critic weights
        for each agent  $i$  do
             $\theta^i = \theta^i + \alpha \nabla_{\theta^i} \log \pi^i(a | h_t^i) \mathbf{Q}(\mathbf{h}_t, \mathbf{a}_t)$  // update actor weights
        end for
    end for
end for
```

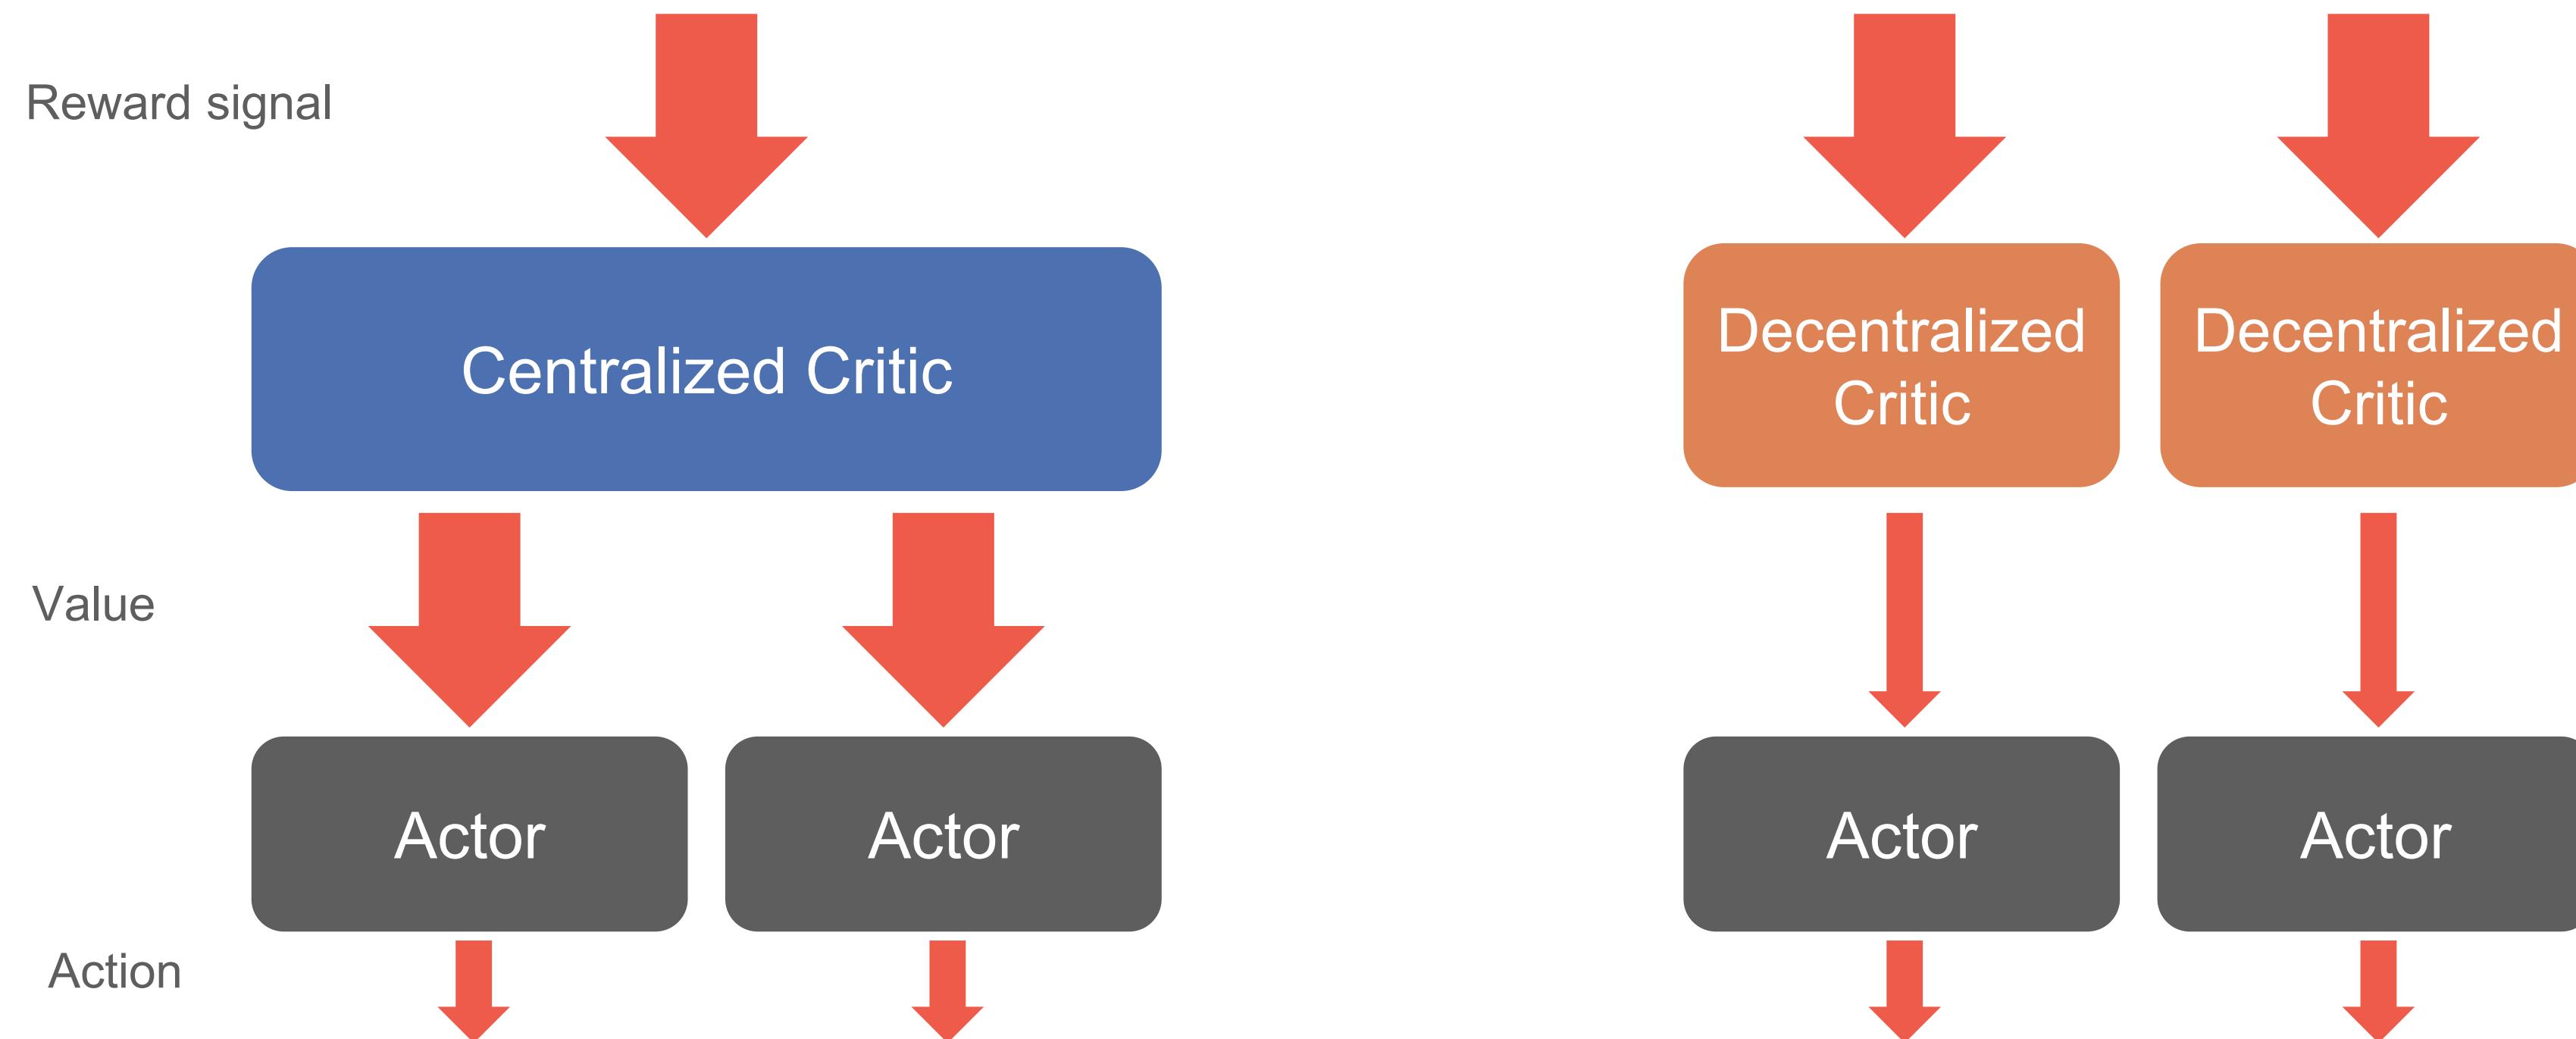
Decentralized actor and centralized critic: update critic based on centralized Q-value and then update each agent's actor



Learning Value Functions

Joint*
Local*

* the return/value/action in the joint/local action-history space

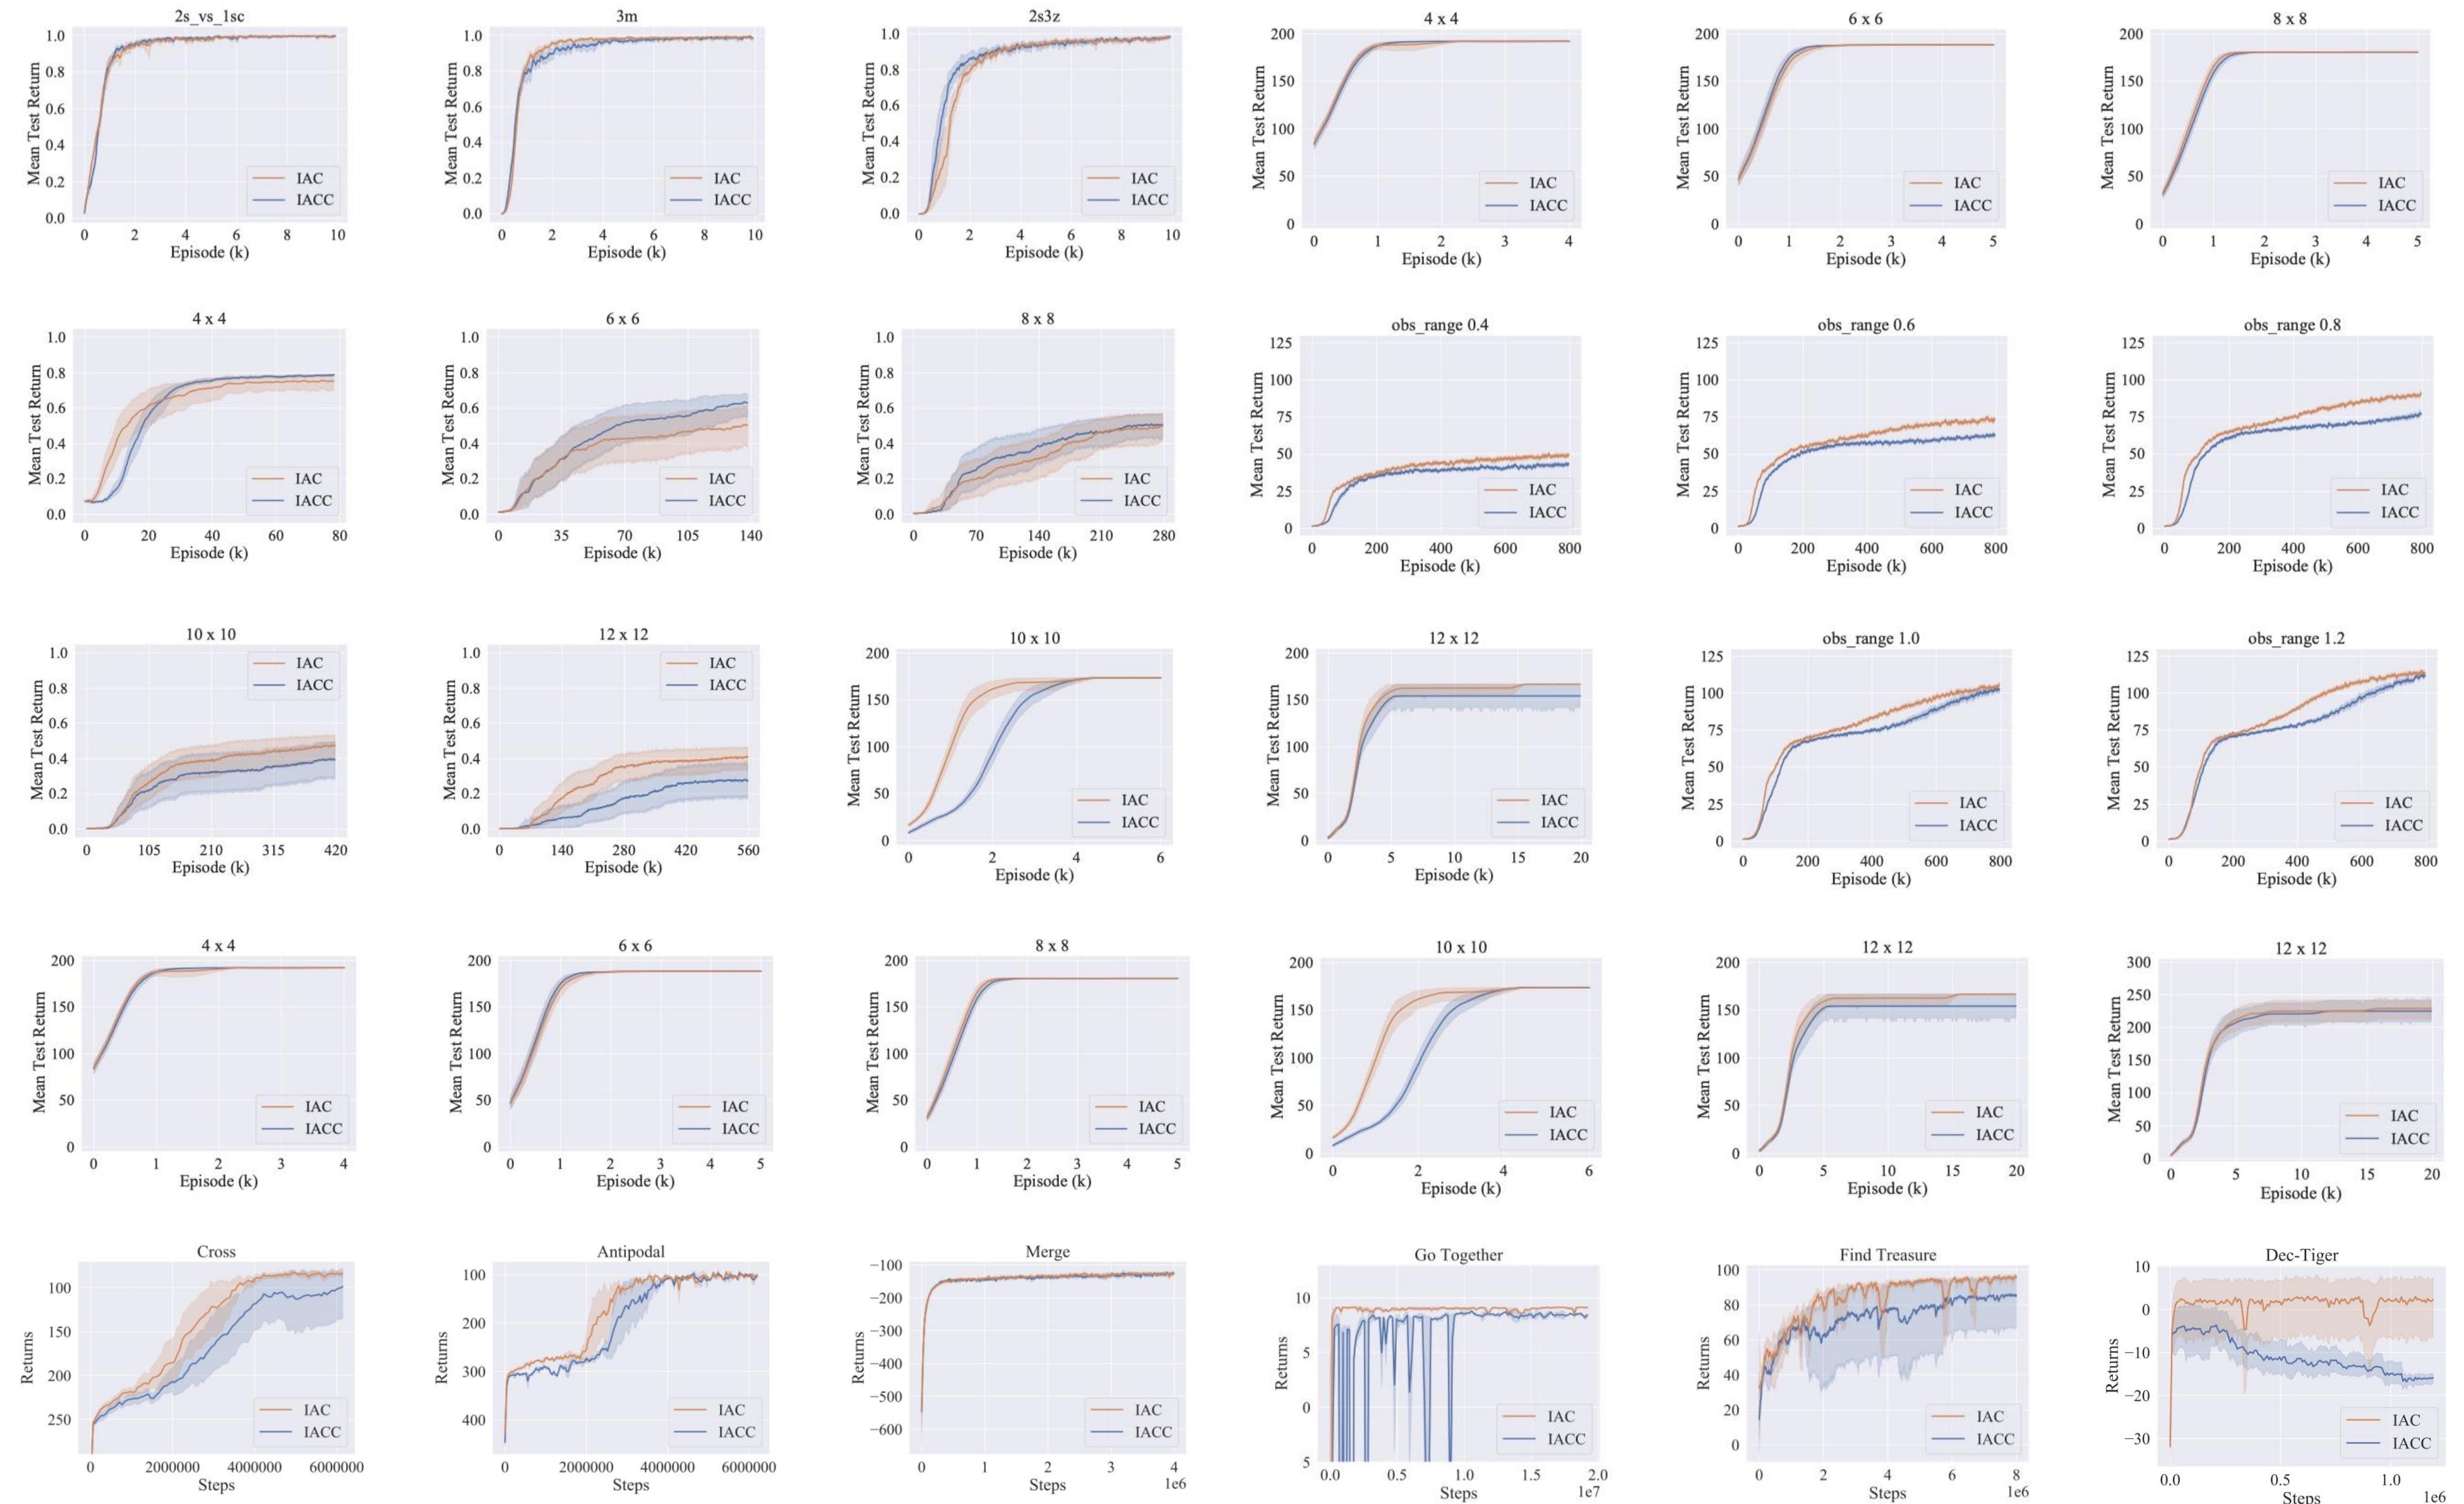


$$\begin{aligned}\nabla_{\theta_i} J_c(\theta_i) &= \mathbb{E}_{\mathbf{a}, \mathbf{h}} [\nabla \log \pi_i(a_i | h_i; \theta_i) Q^{\boldsymbol{\pi}}(\mathbf{h}, \mathbf{a}; \phi)] \\ &= \mathbb{E}_{a_i, h_i} \left[\nabla \log \pi_i(a_i | h_i; \theta_i) \mathbb{E}_{a_{-i}, h_{-i}} [Q^{\boldsymbol{\pi}}(h_i, h_{-i}, a_i, a_{-i})] \right]\end{aligned}$$

Both estimating and updating decentralized policies

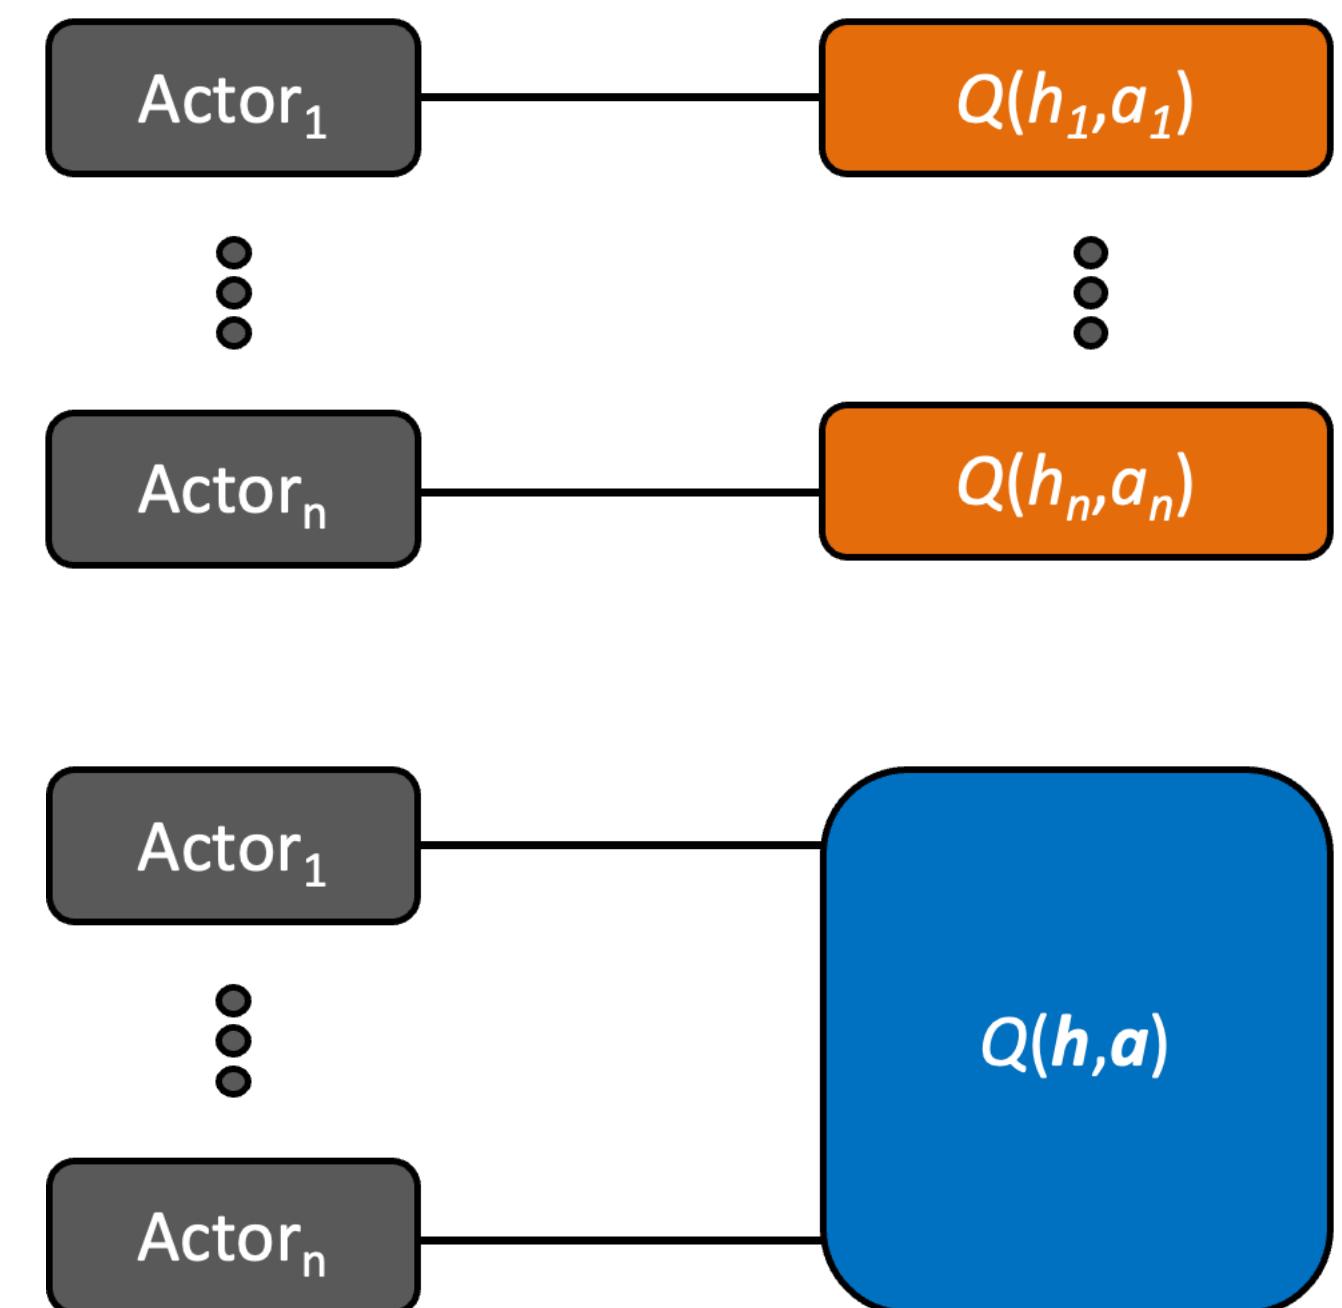
Centralized and Decentralized Critic Performance

on StarCraft Multi-Agent Challenge (SMAC), Box Pushing, Particle environments, Target Capture, etc.



Decentralized vs centralized critics

- Theoretically equivalent
 - But that assumes learned critics
- Decentralized critics can be harder to learn
 - When other agents change policies
 - Higher bias
- Centralized critics can be harder to learn
 - Large domains (action, obs, agents)
 - Higher variance to marginalize out other agents



State-based Centralized Critics

State information is often available offline in a simulator

Implemented by pioneering Centralized Critic methods

COMA (Foerster et al. 2018), MADDPG (Lowe et al. 2017)

Followed by later methods

SQDDPG (Wang et al. 2020), LIIR (Du et al. 2019), LICA (Zhou et al. 2020), VDAC-mix (Su, Adams, and Beling 2021), DOP (Wang et al. 2021) and MACKRL (Schroeder de Witt et al. 2019)

Obvious Advantages of State-based Centralized Critic

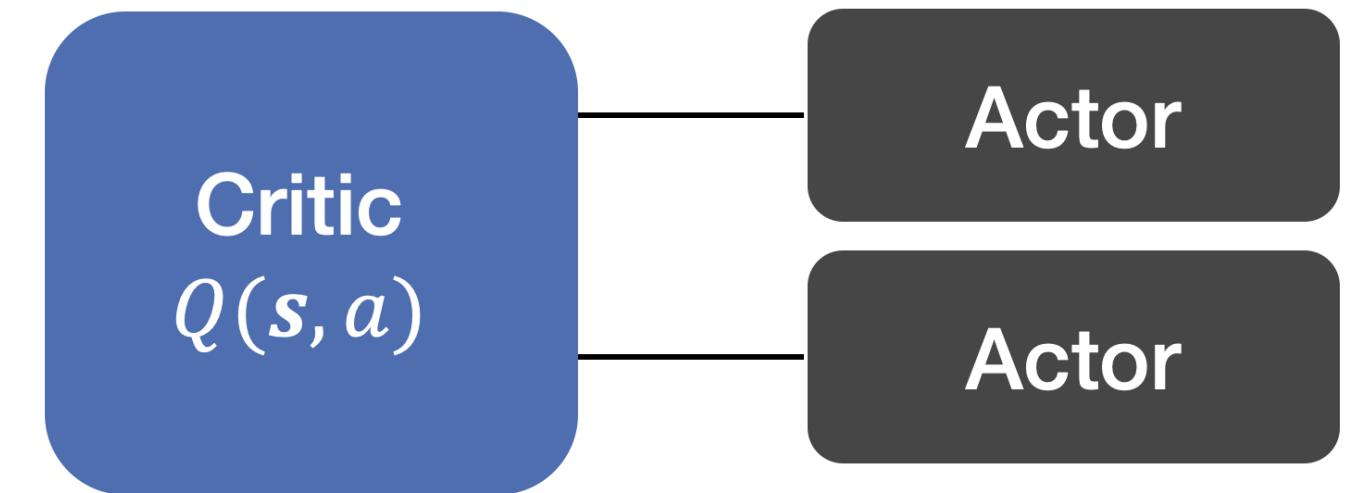
Compact, Fully Observable

Obvious Disadvantages of History-based Centralized Critic

Complexity from (potentially long) time horizon

Complexity from combining observations (and actions) from multiple agents

Partially Observable



A Deeper Understanding of State-Based Critics in Multi-Agent Reinforcement Learning

[Lyu, Baisero, Xiao and Amato – AAAI22](#)

State-based critics in MARL are popular but misunderstood

We show in theory:

State-based critics may be biased compared to **History-based Critics**

State-based critics may produce higher variance

We show empirically:

Both critics work well in different domains

Common benchmarks lack partial observability

The **state-history-based critic** is robust to various domains

Centralized critics

Centralized critic

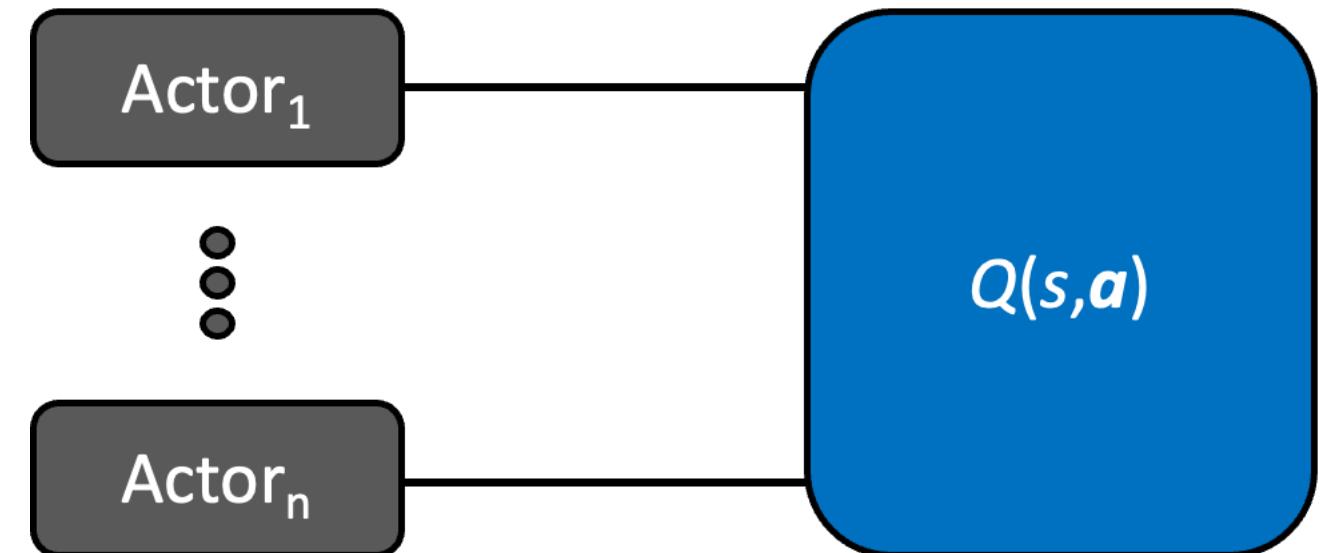
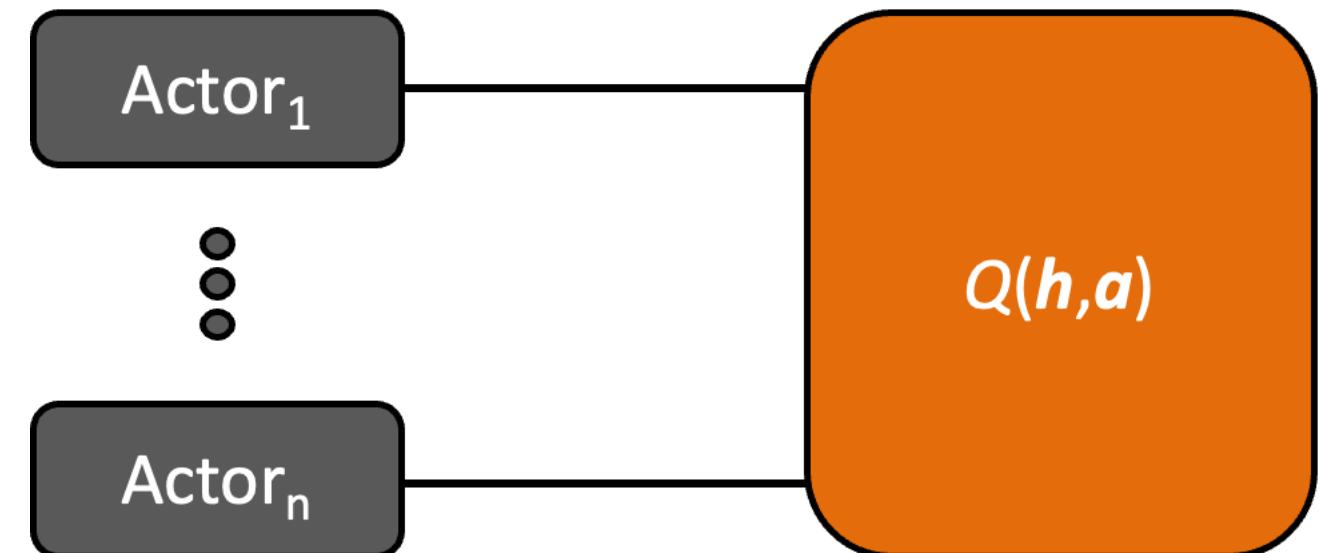
Conditions on history of all agents (joint history \mathbf{h})

$$\nabla_i J_{\mathbf{h}} = \mathbb{E}_{\mathbf{h} \sim \rho(\mathbf{h}), \mathbf{a} \sim \pi(\mathbf{h})} [Q^{\pi}(\mathbf{h}, \mathbf{a}) \nabla_{\theta_i} \log \pi_i(a_i; h_i)]$$

State-based centralized critic

Conditions on the world state s

$$\nabla_i J_s = \mathbb{E}_{\mathbf{h}, s \sim \rho(\mathbf{h}, s), \mathbf{a} \sim \pi(\mathbf{h})} [Q^{\pi}(s, \mathbf{a}) \nabla_{\theta_i} \log \pi_i(a_i; h_i)]$$

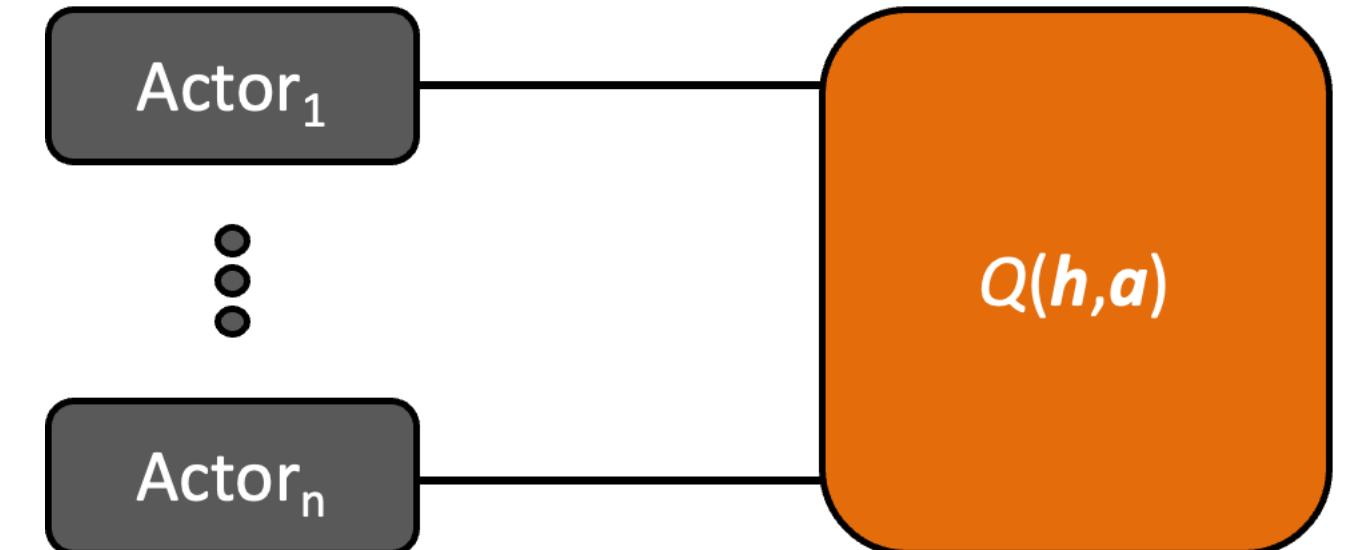


Centralized critics

Centralized critic

Conditions on history of all agents (joint history \mathbf{h})

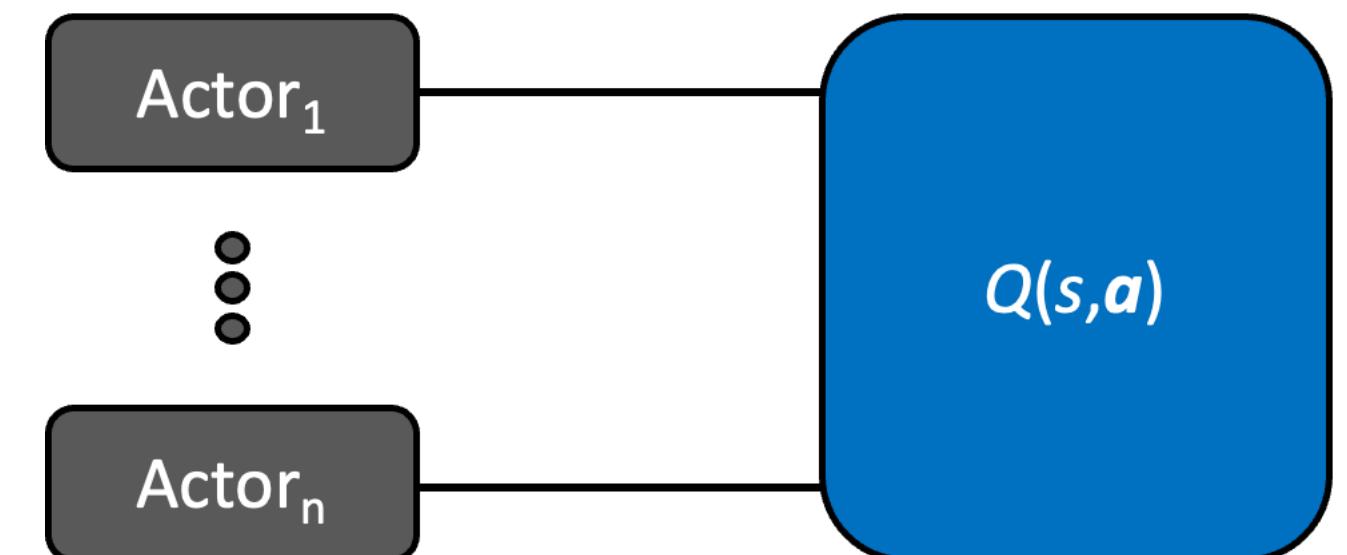
$$\nabla_i J_{\mathbf{h}} = \mathbb{E}_{\mathbf{h} \sim \rho(\mathbf{h}), \mathbf{a} \sim \pi(\mathbf{h})} [Q^{\pi}(\mathbf{h}, \mathbf{a}) \nabla_{\theta_i} \log \pi_i(a_i; h_i)]$$



State-based centralized critic

Conditions on the world state s

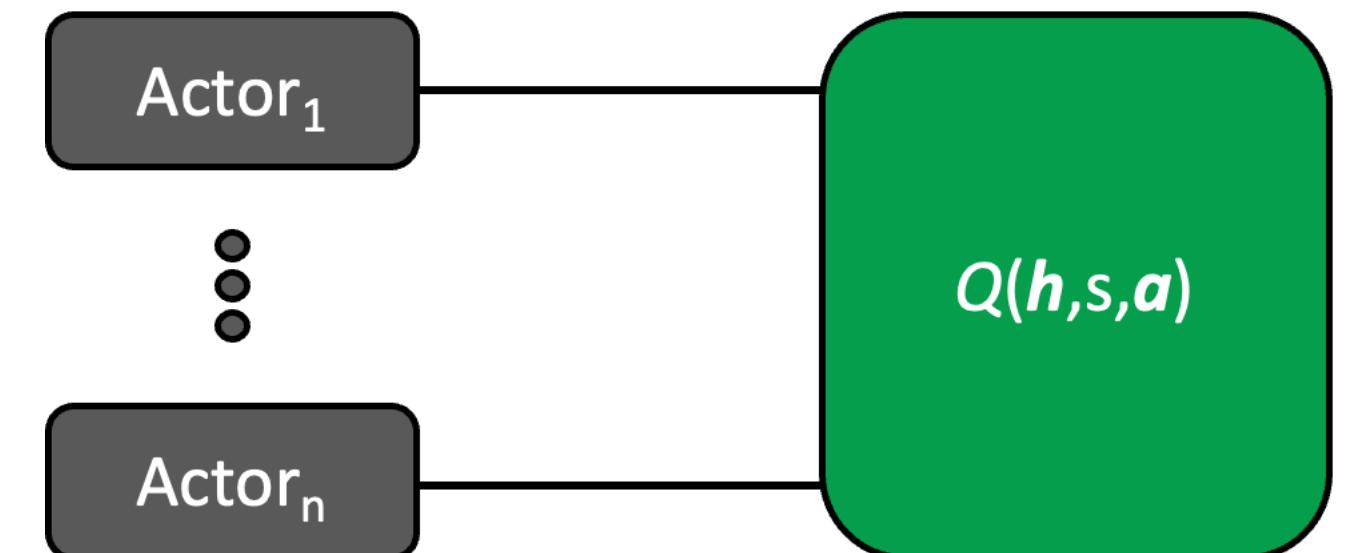
$$\nabla_i J_s = \mathbb{E}_{\mathbf{h}, s \sim \rho(\mathbf{h}, s), \mathbf{a} \sim \pi(\mathbf{h})} [Q^{\pi}(s, \mathbf{a}) \nabla_{\theta_i} \log \pi_i(a_i; h_i)]$$



History-state-based centralized critic

Conditions on the joint history \mathbf{h} and world state s

$$\nabla_i J_s = \mathbb{E}_{\mathbf{h}, s \sim \rho(\mathbf{h}, s), \mathbf{a} \sim \pi(\mathbf{h})} [Q^{\pi}(s, \mathbf{h}, \mathbf{a}) \nabla_{\theta_i} \log \pi_i(a_i; h_i)]$$



Experiments

Tested with advantage actor critic (A2C)

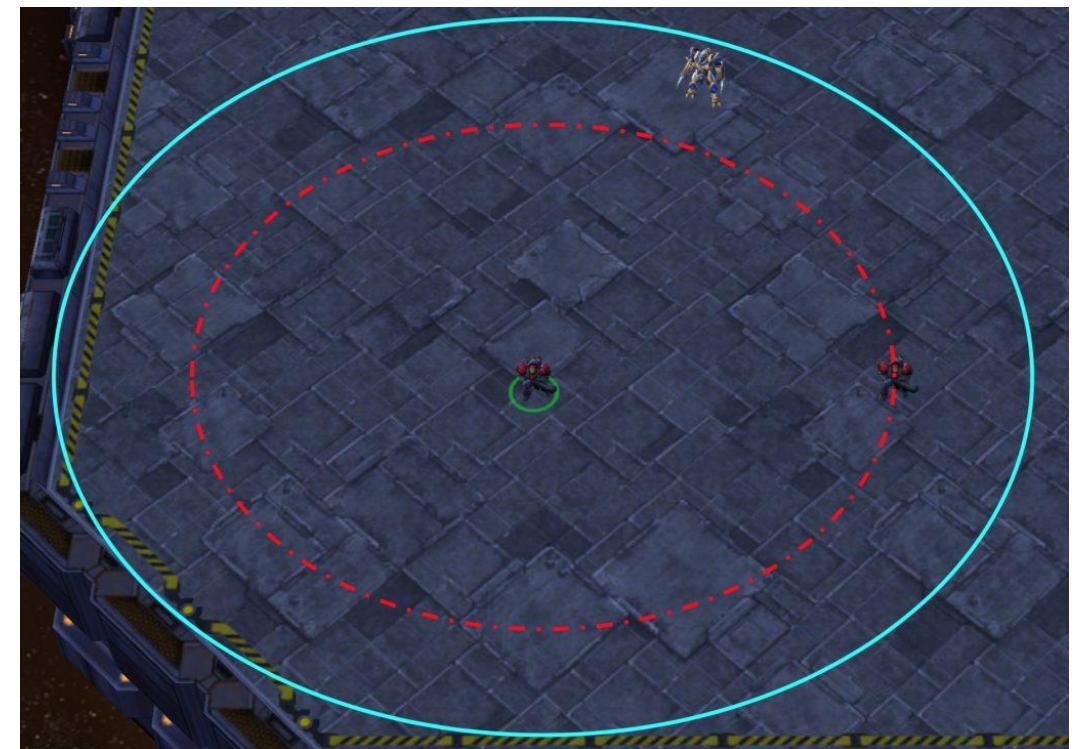
History critic

State critic

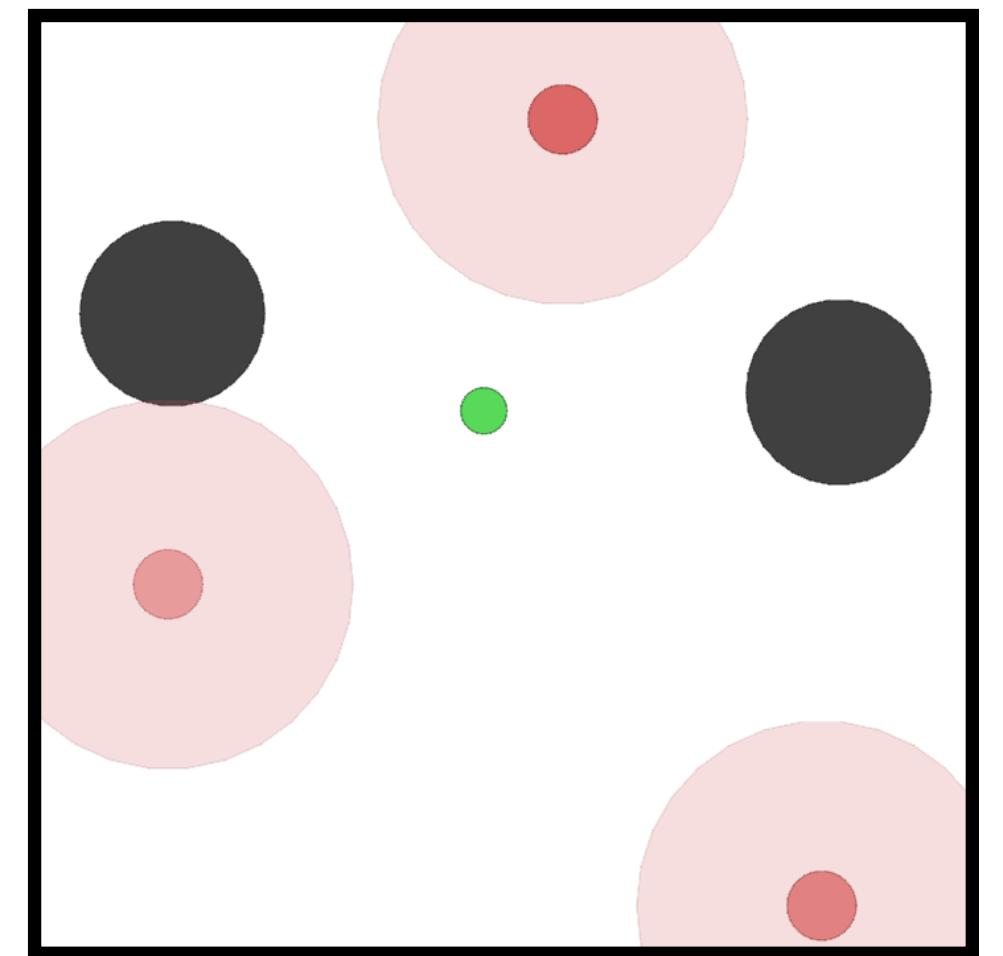
State-history critic

Used standard domains: small common domains, SMACv1 (Starcraft) and partially observable particle environments

Have additional experiments and base actor-critic methods in the paper

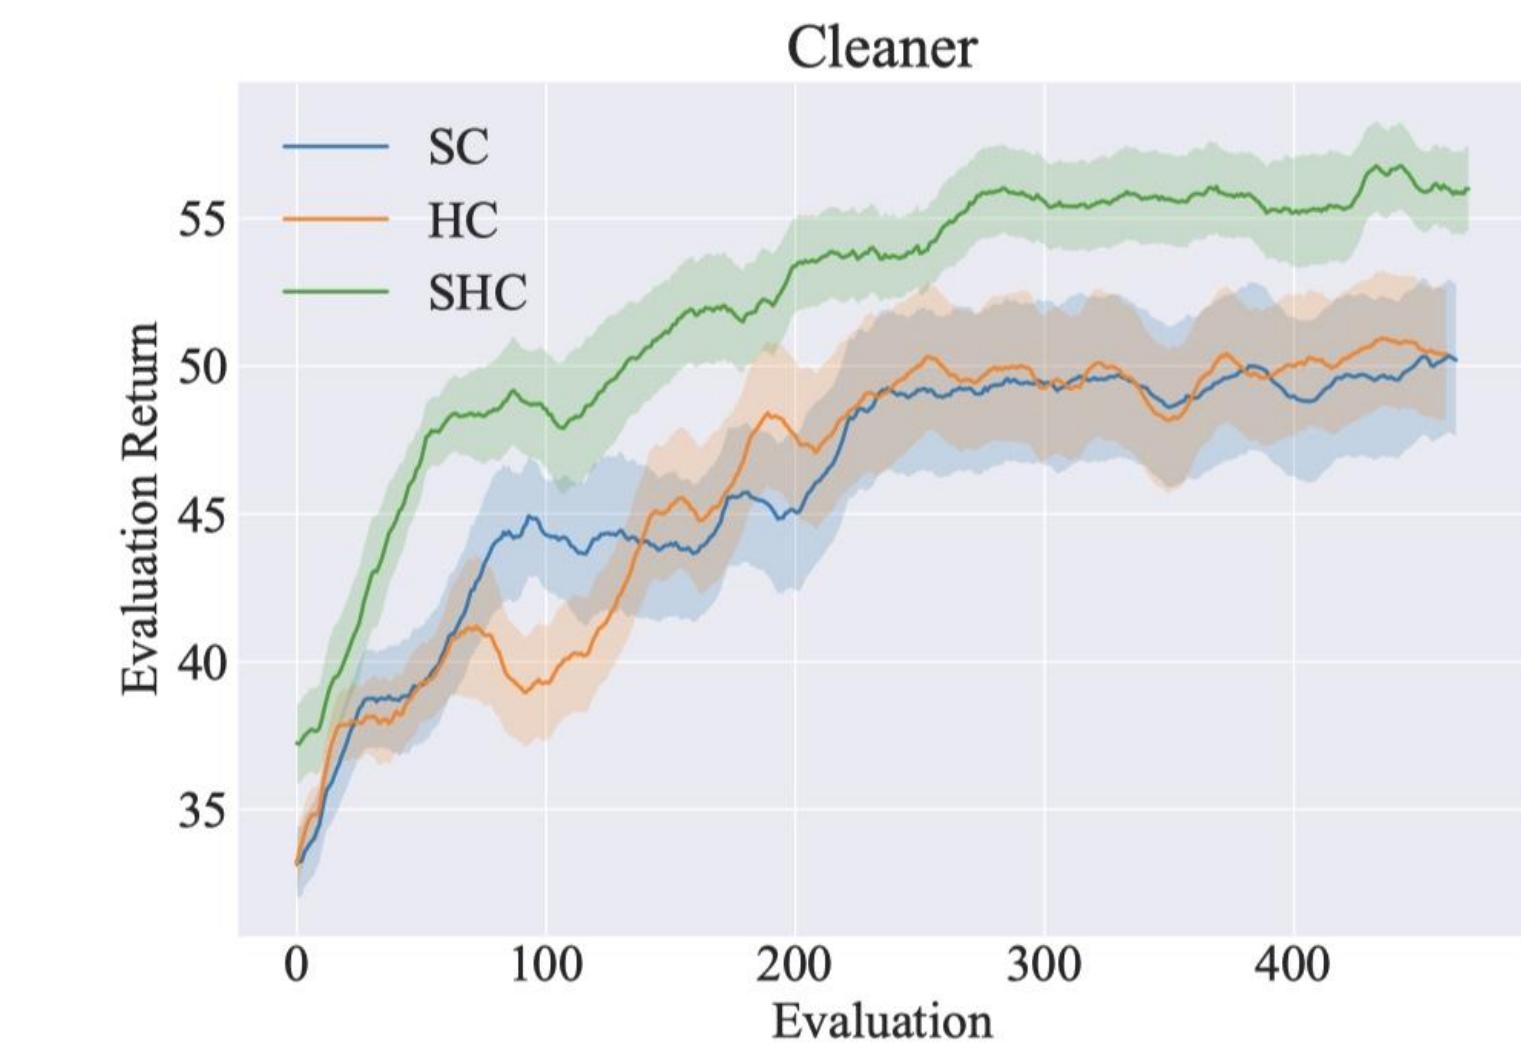
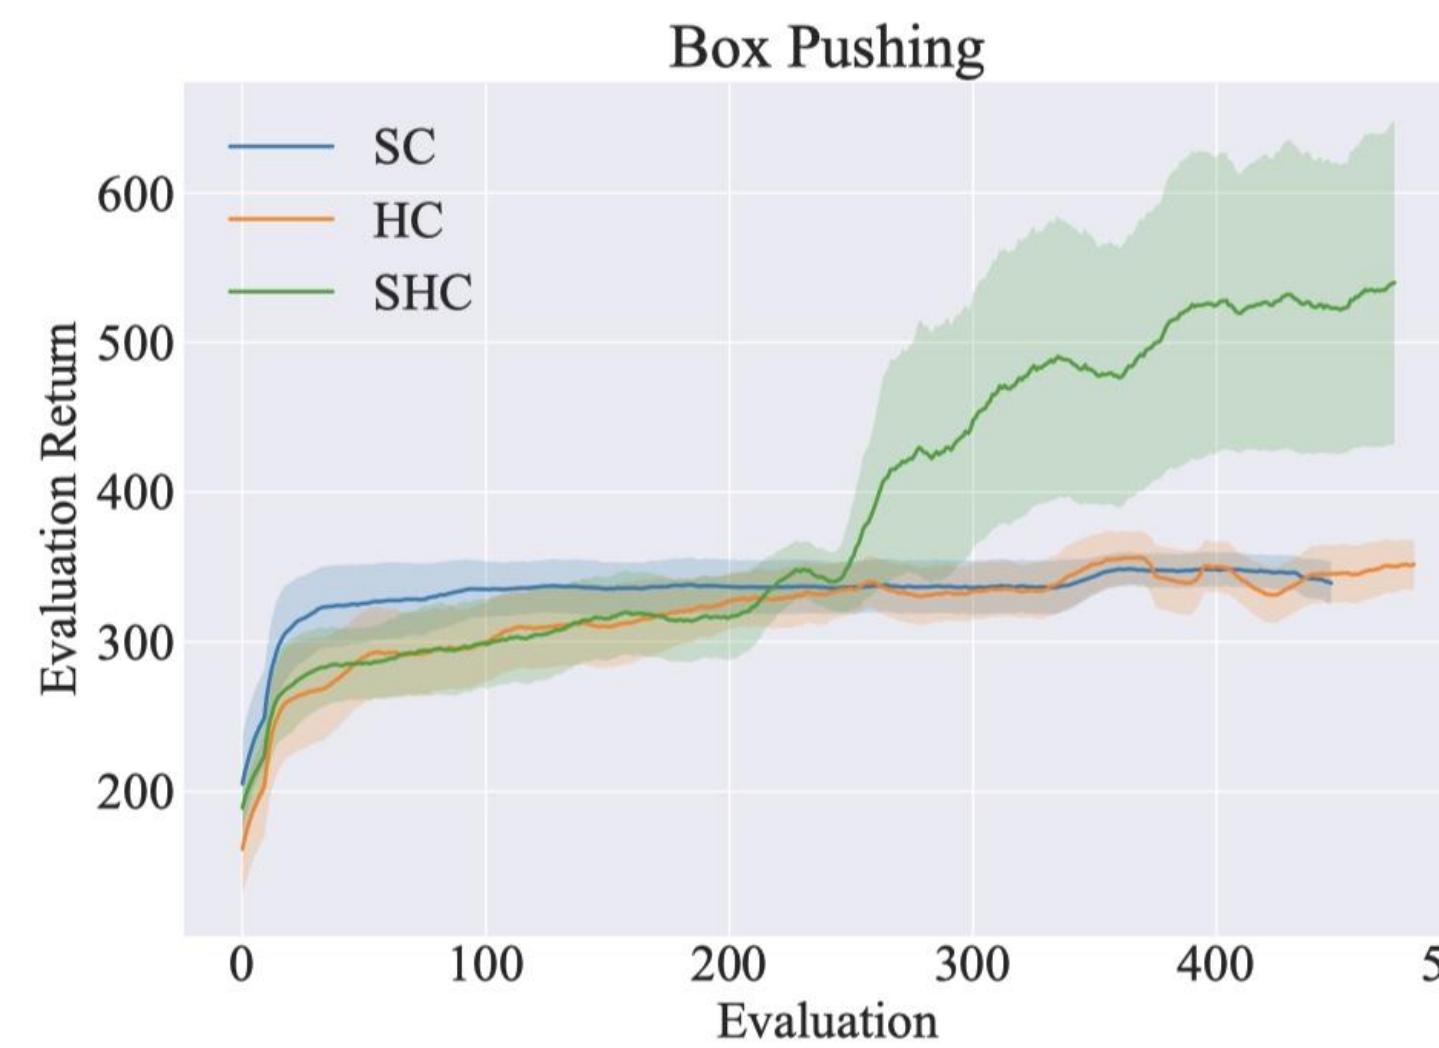
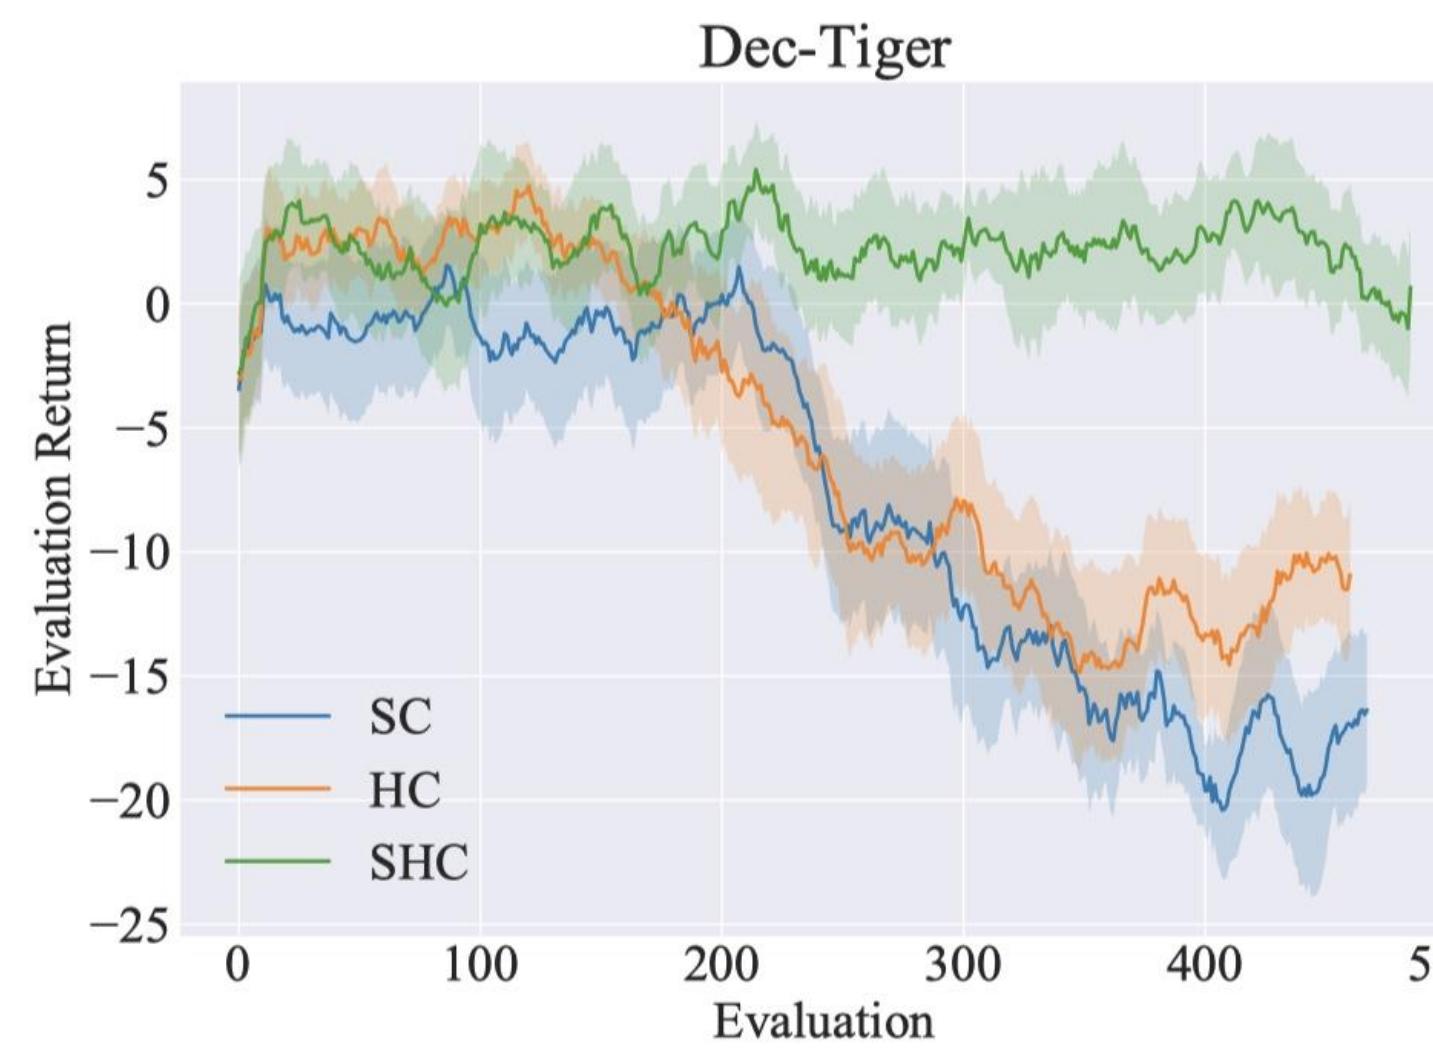
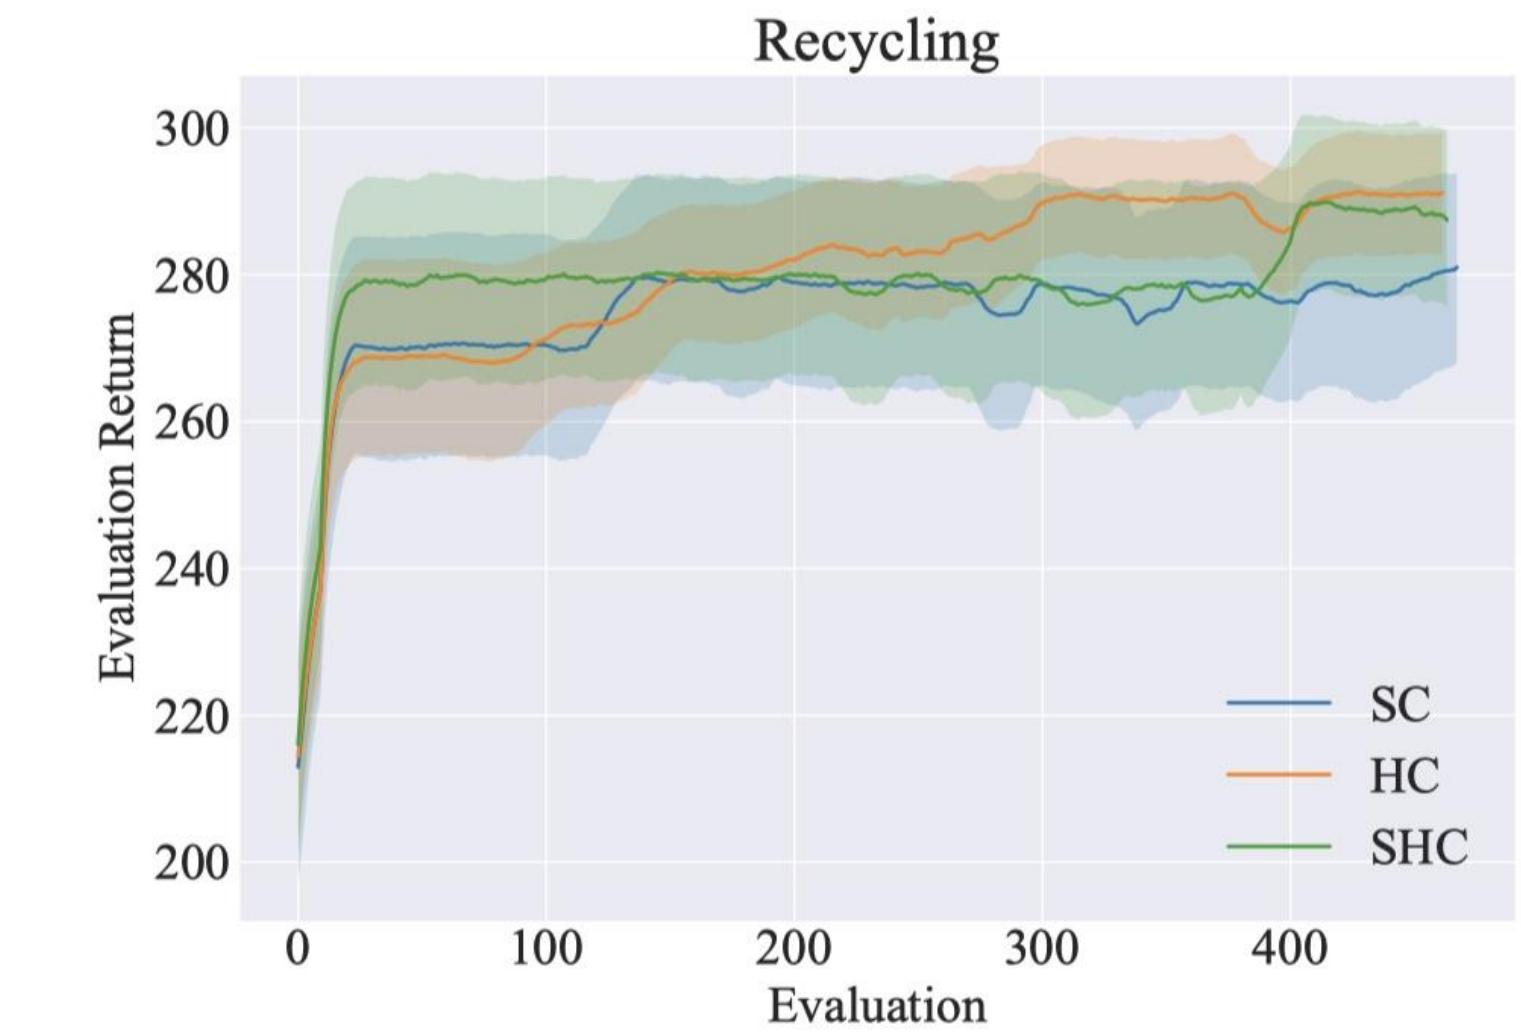
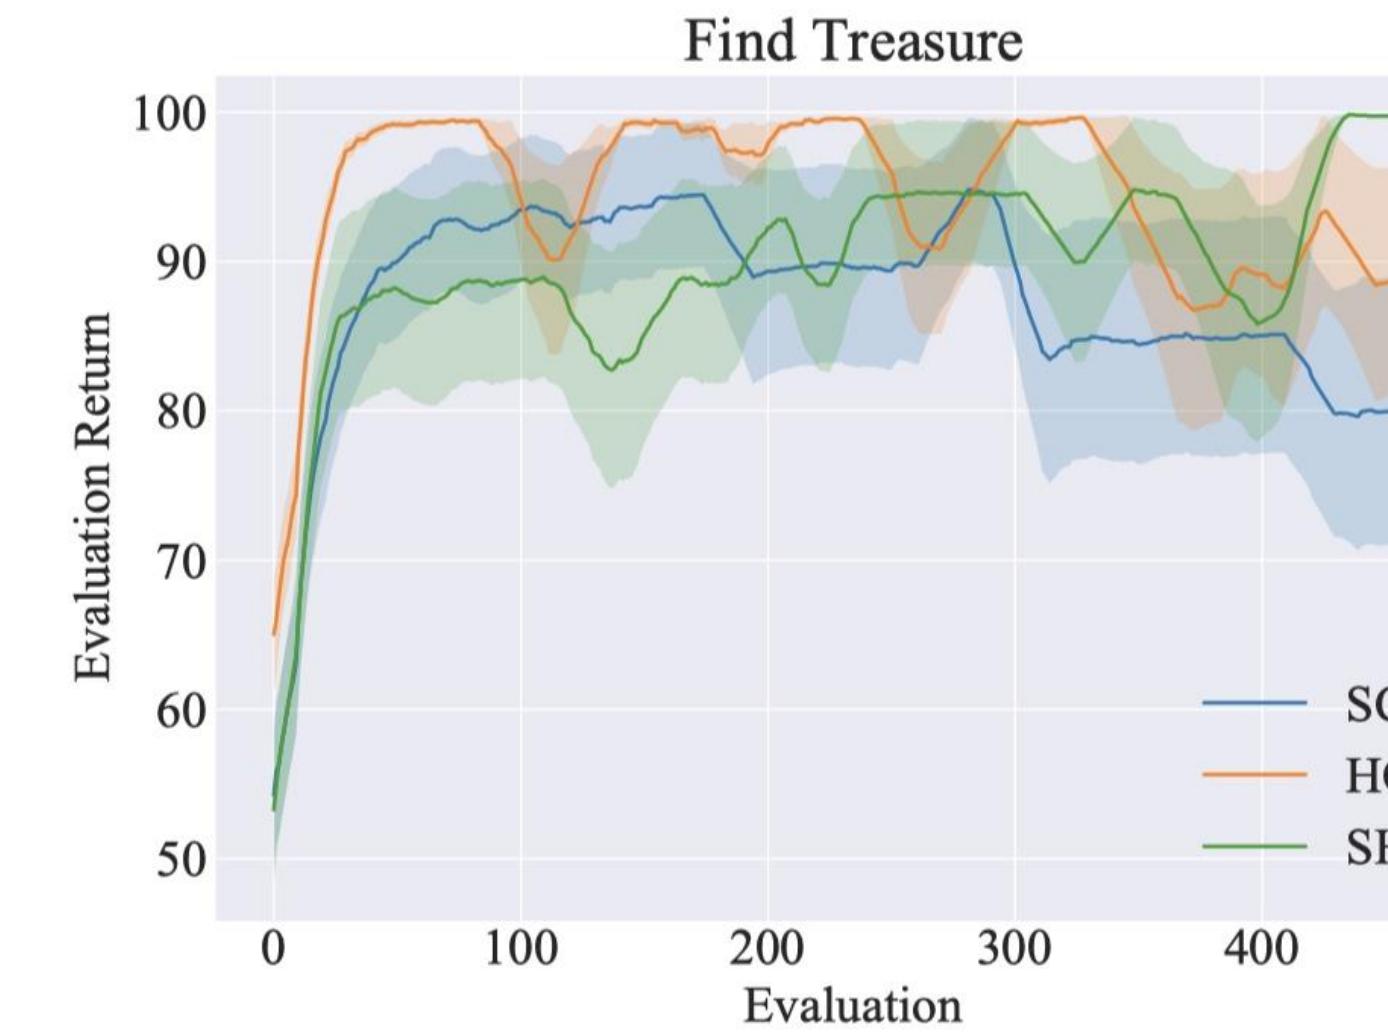
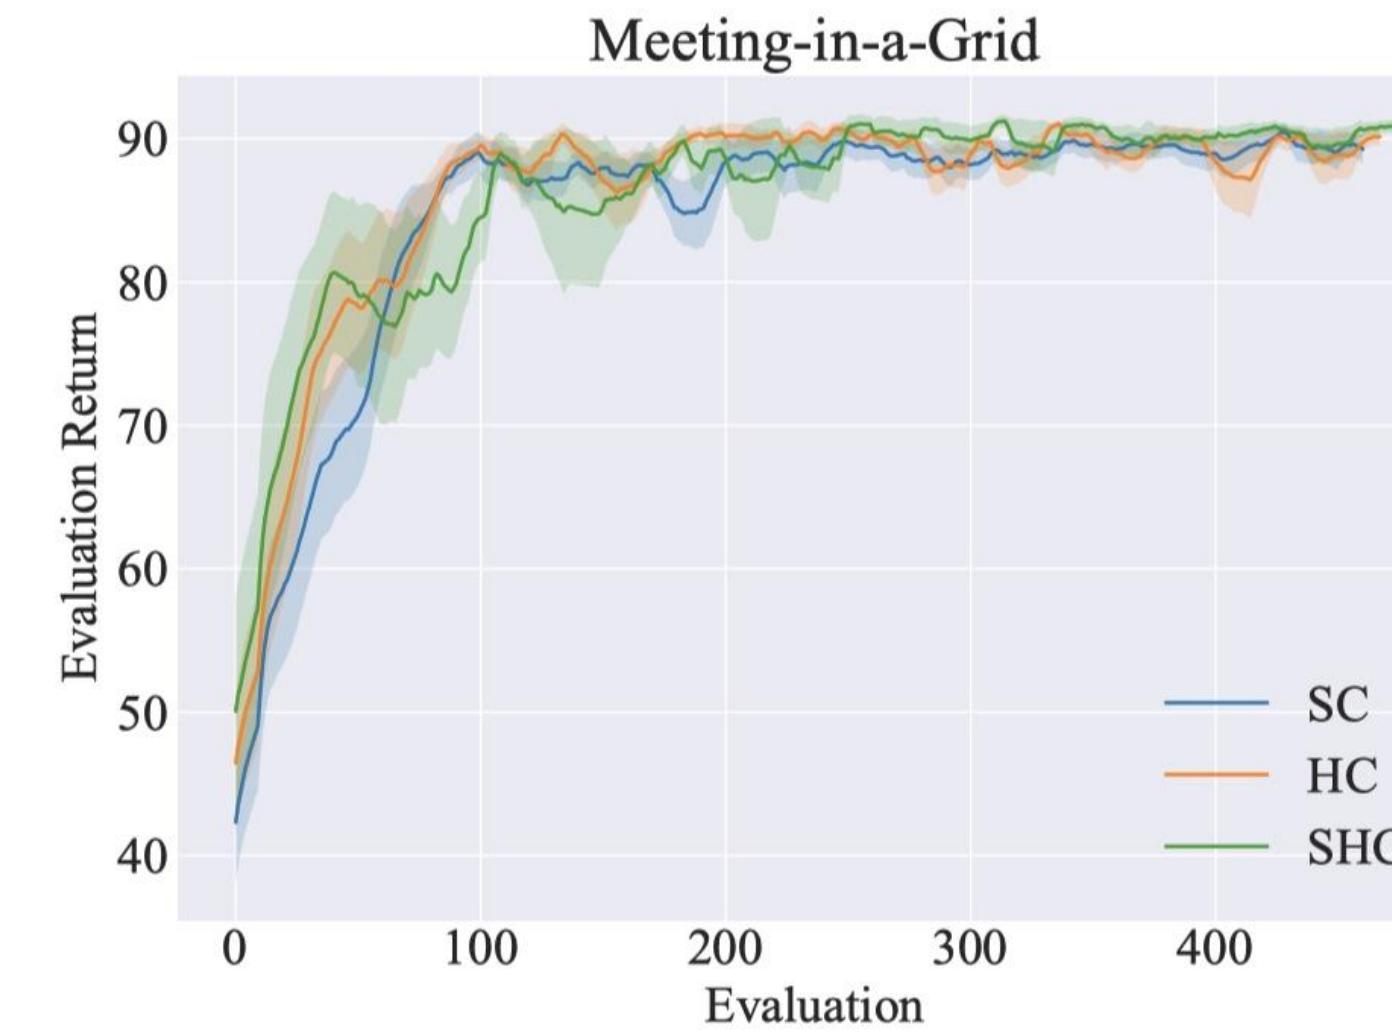


SMAC v1

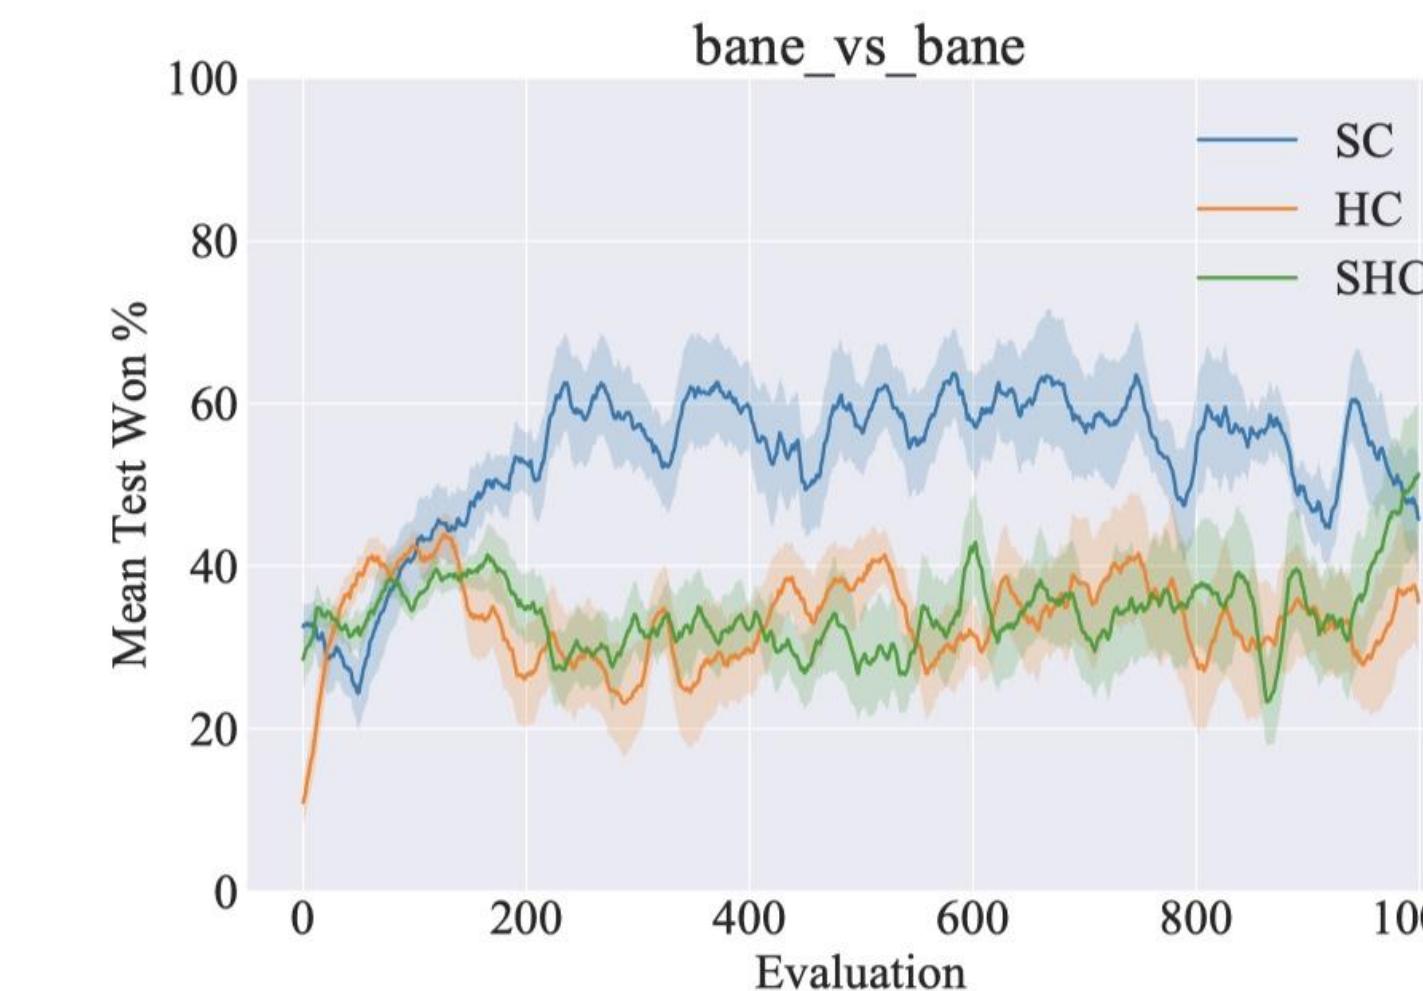
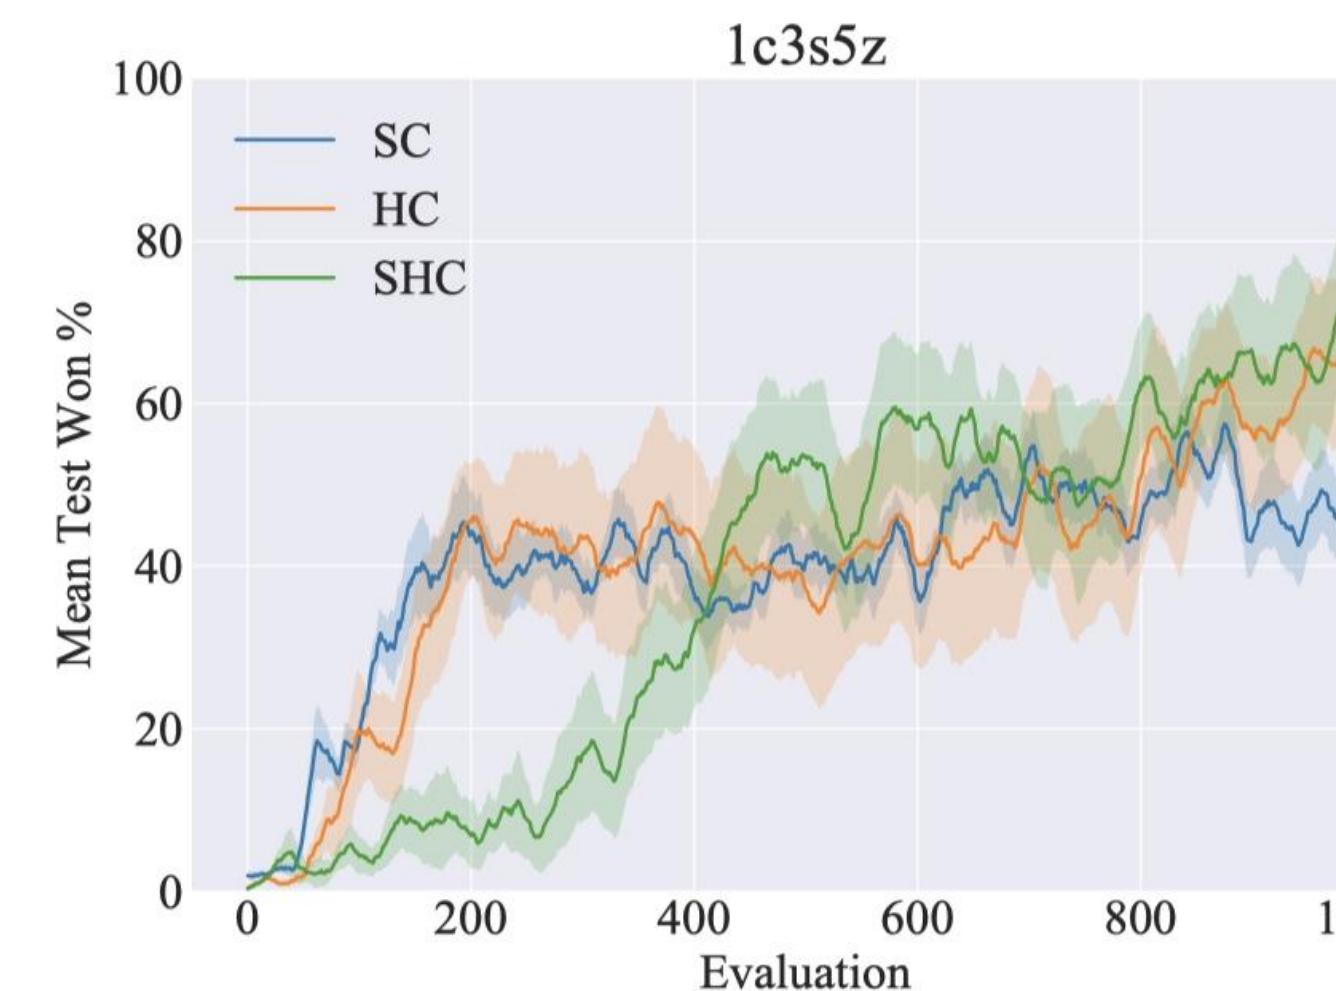
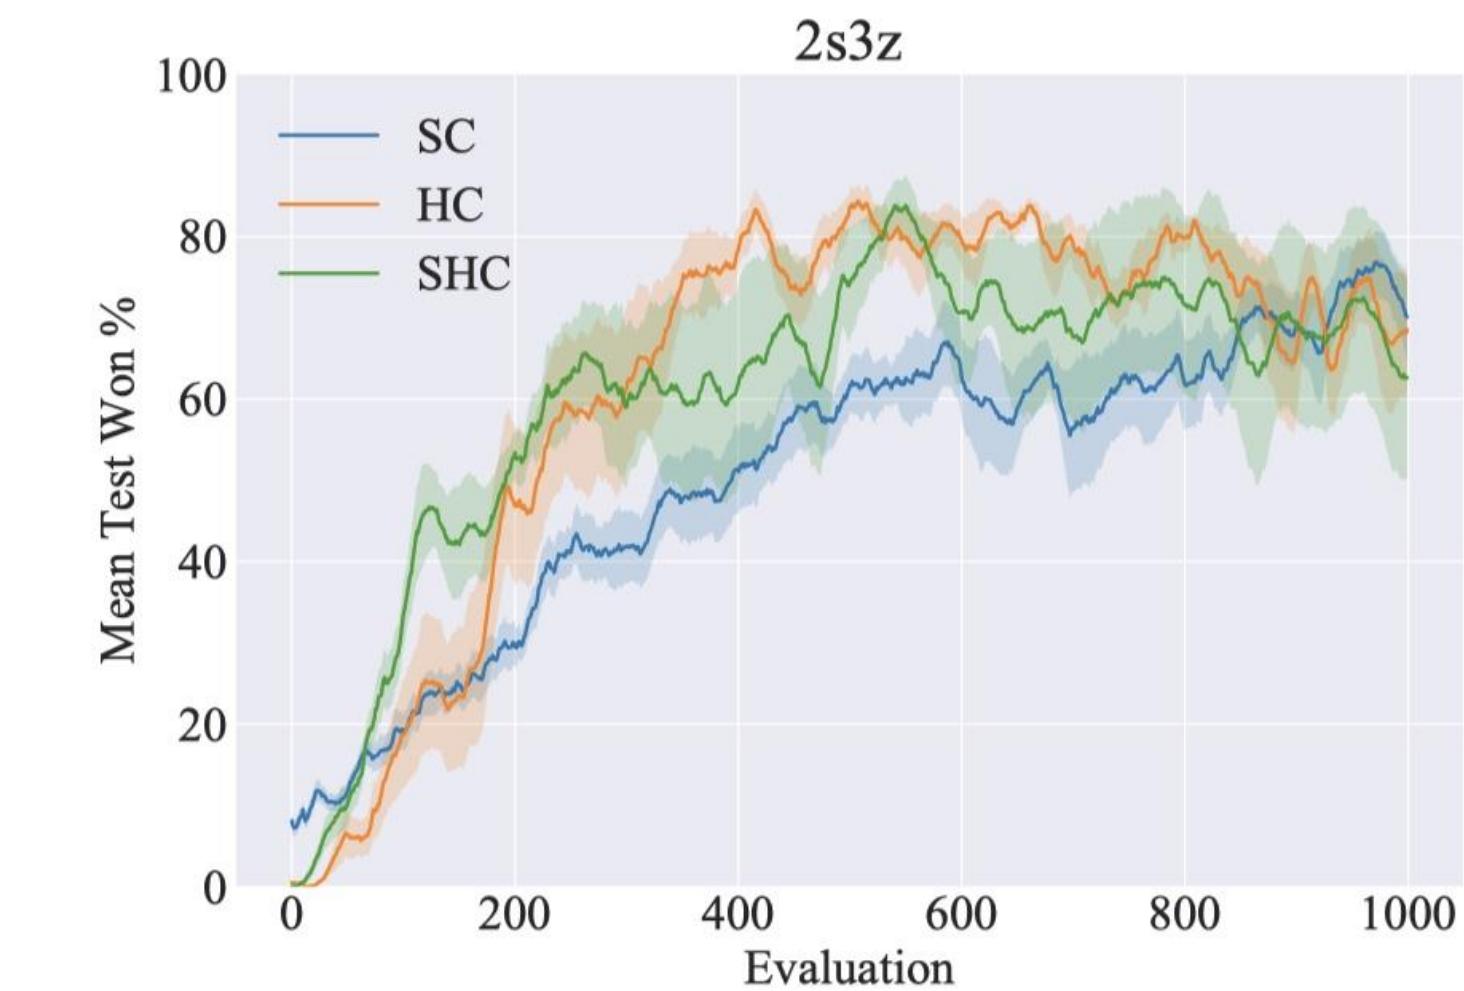
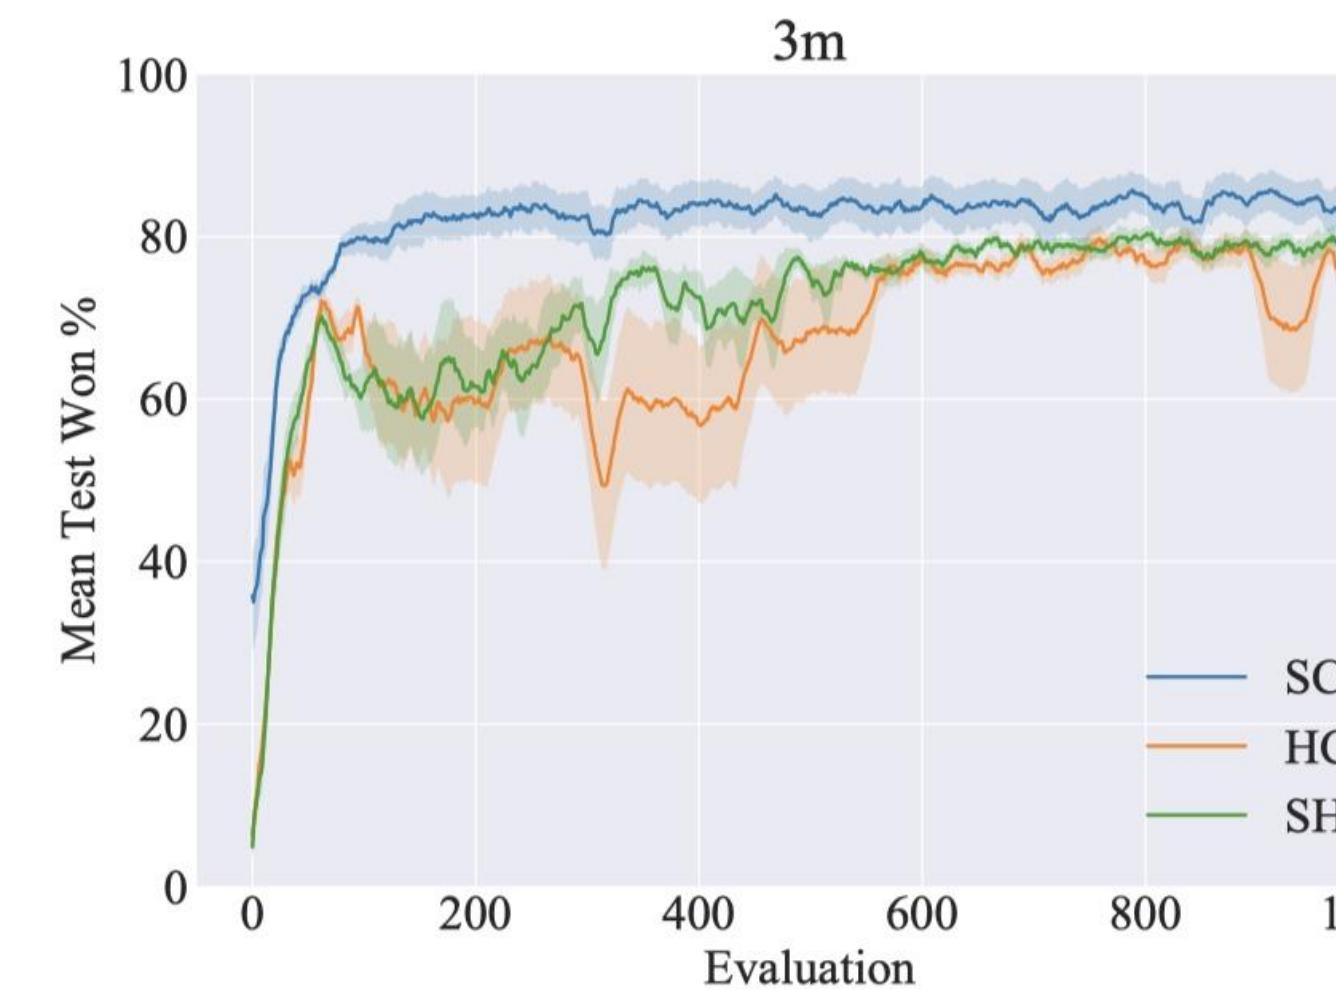
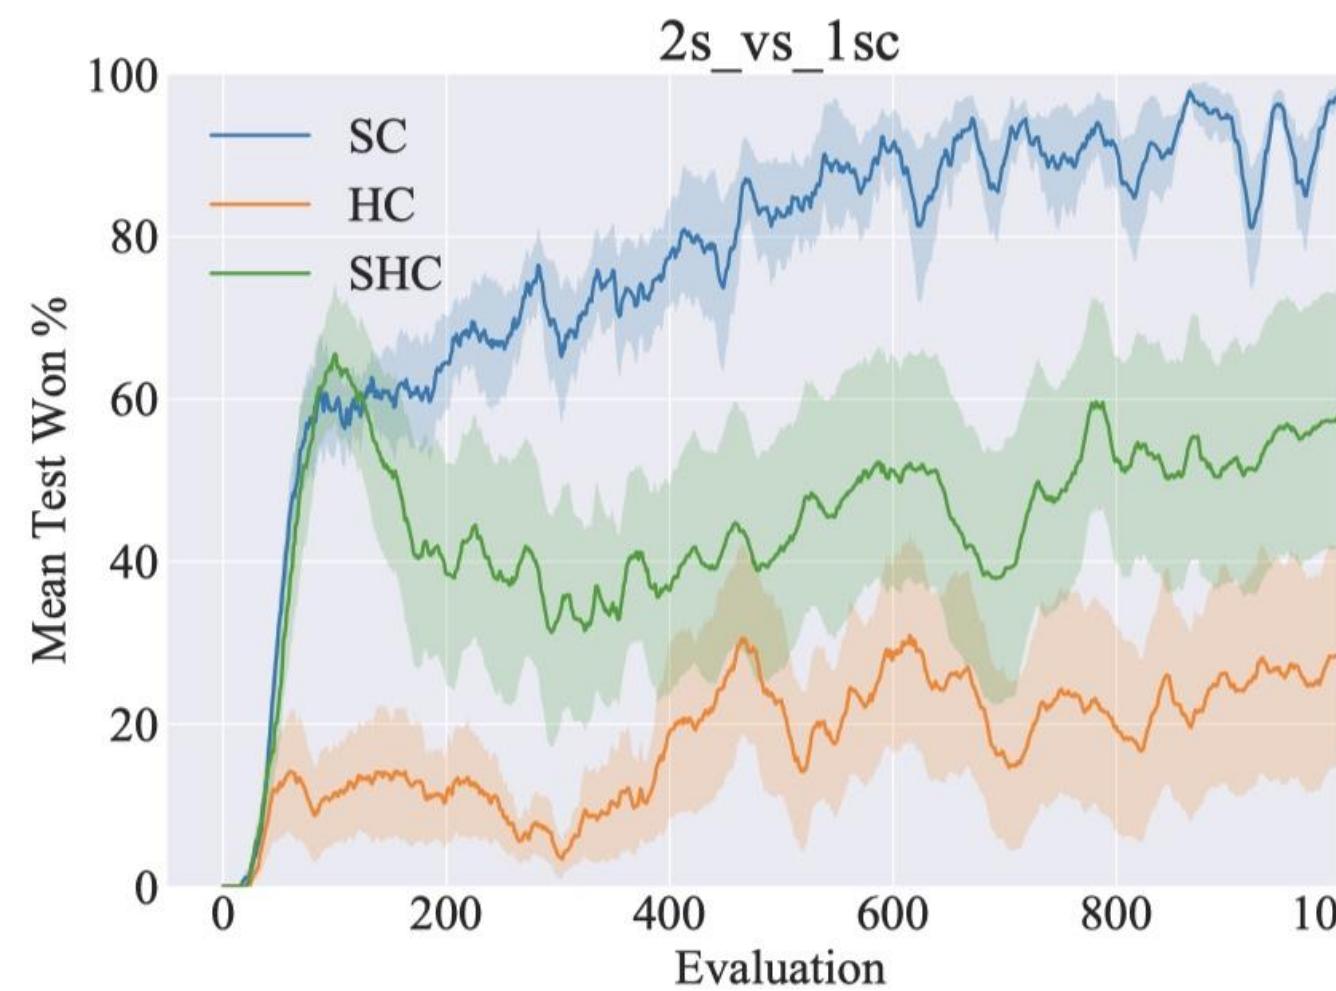


Partially observable particle envs

Common small environments

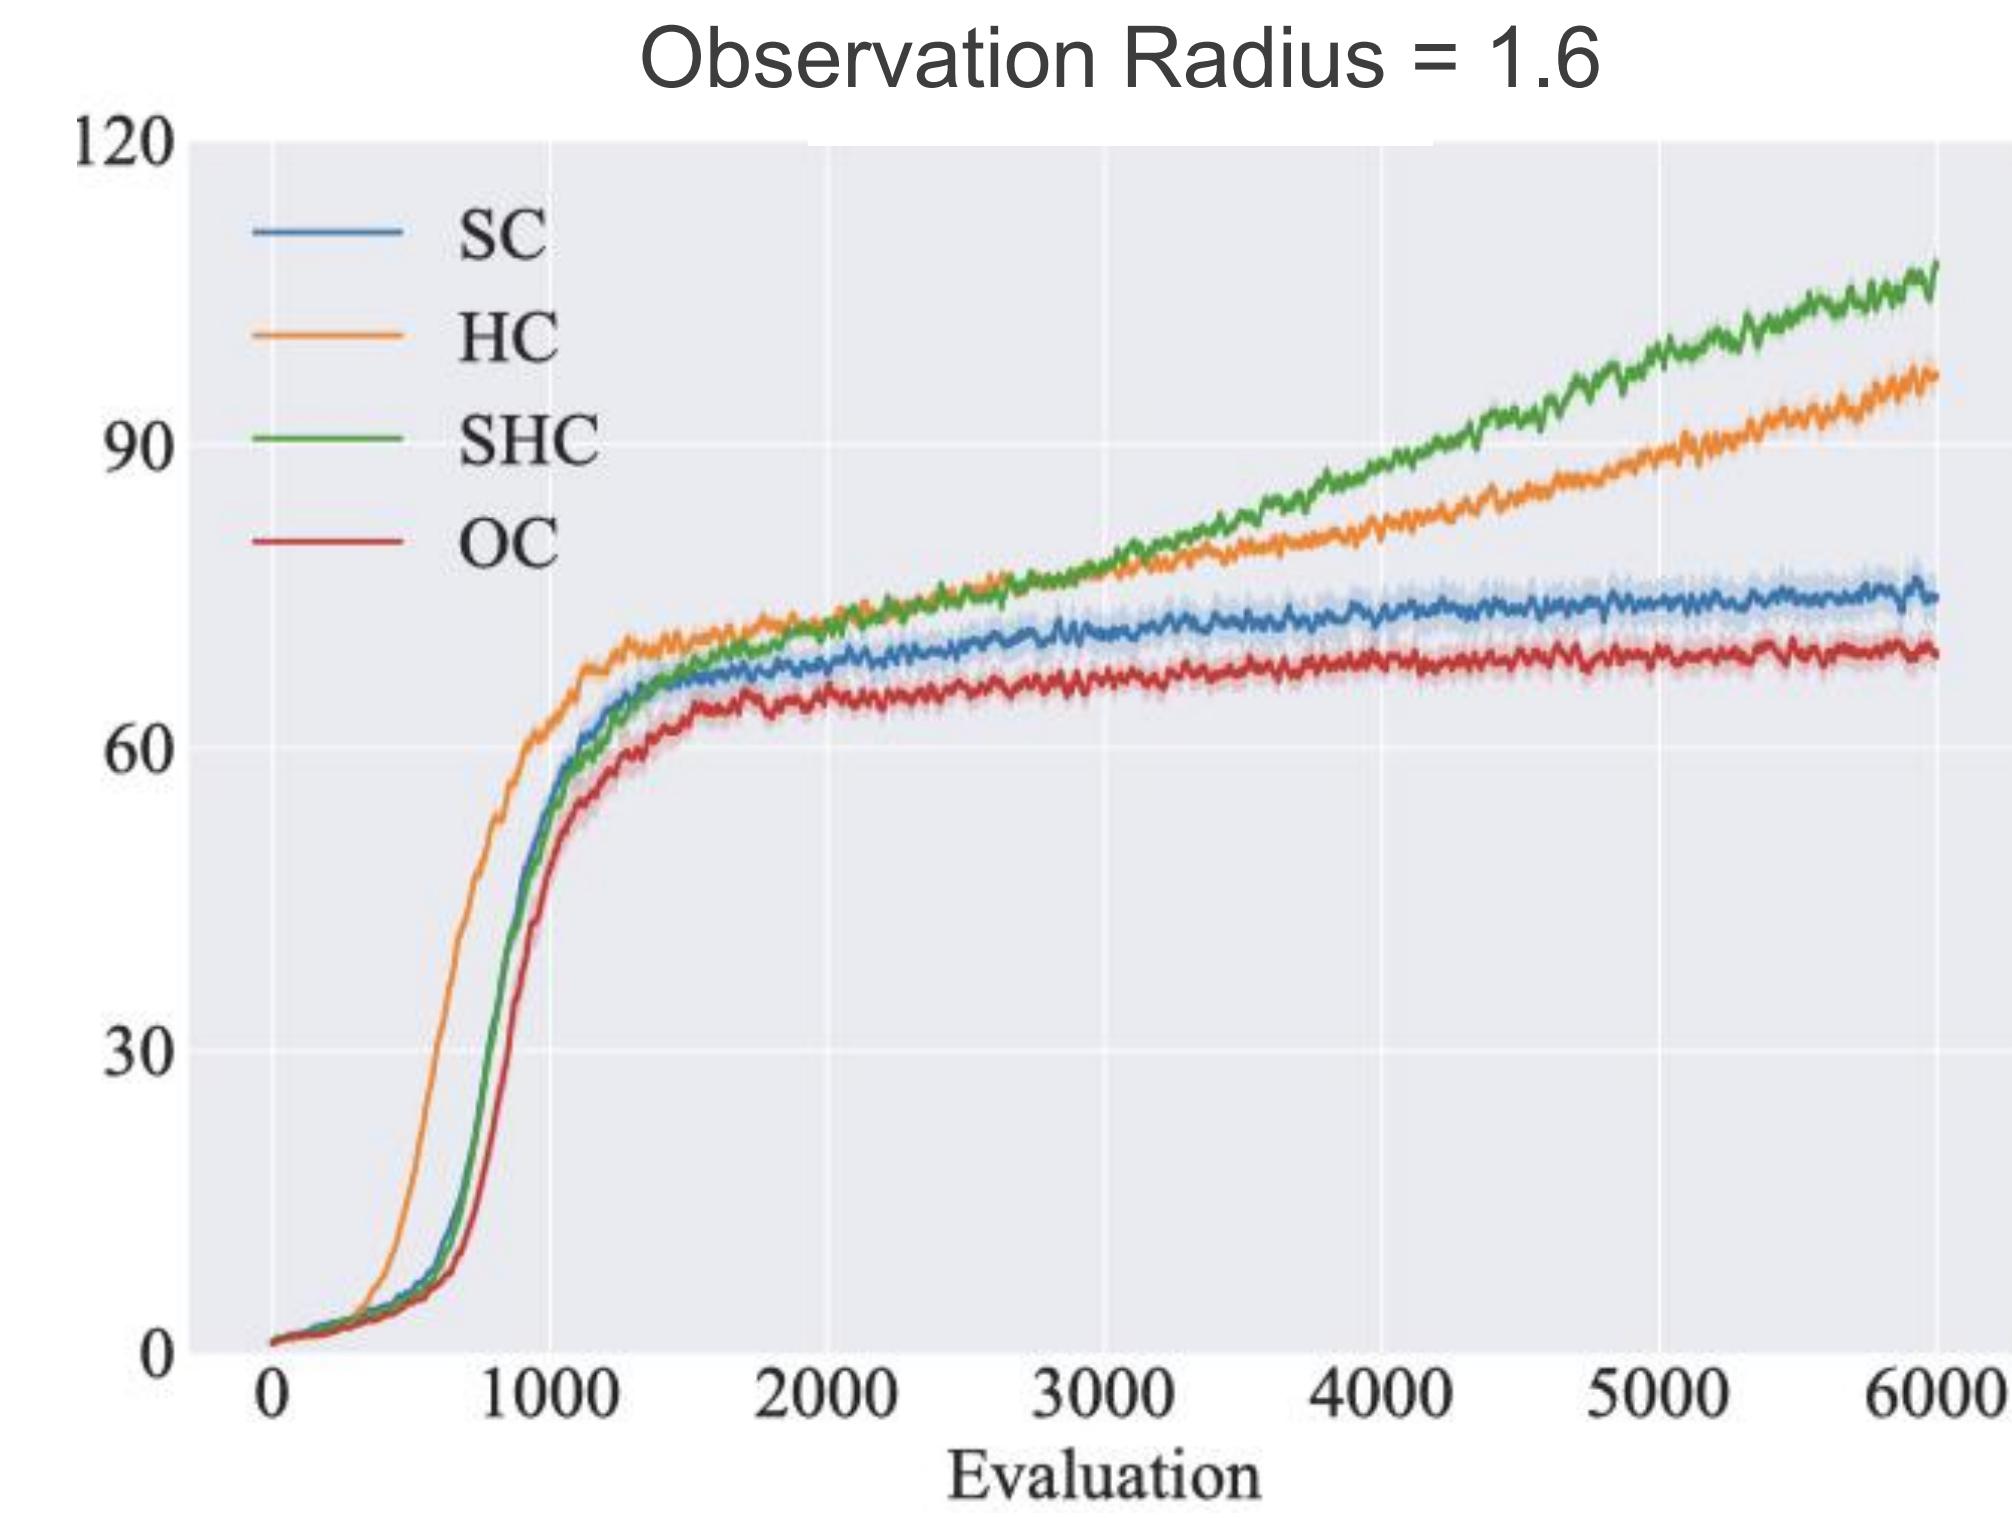
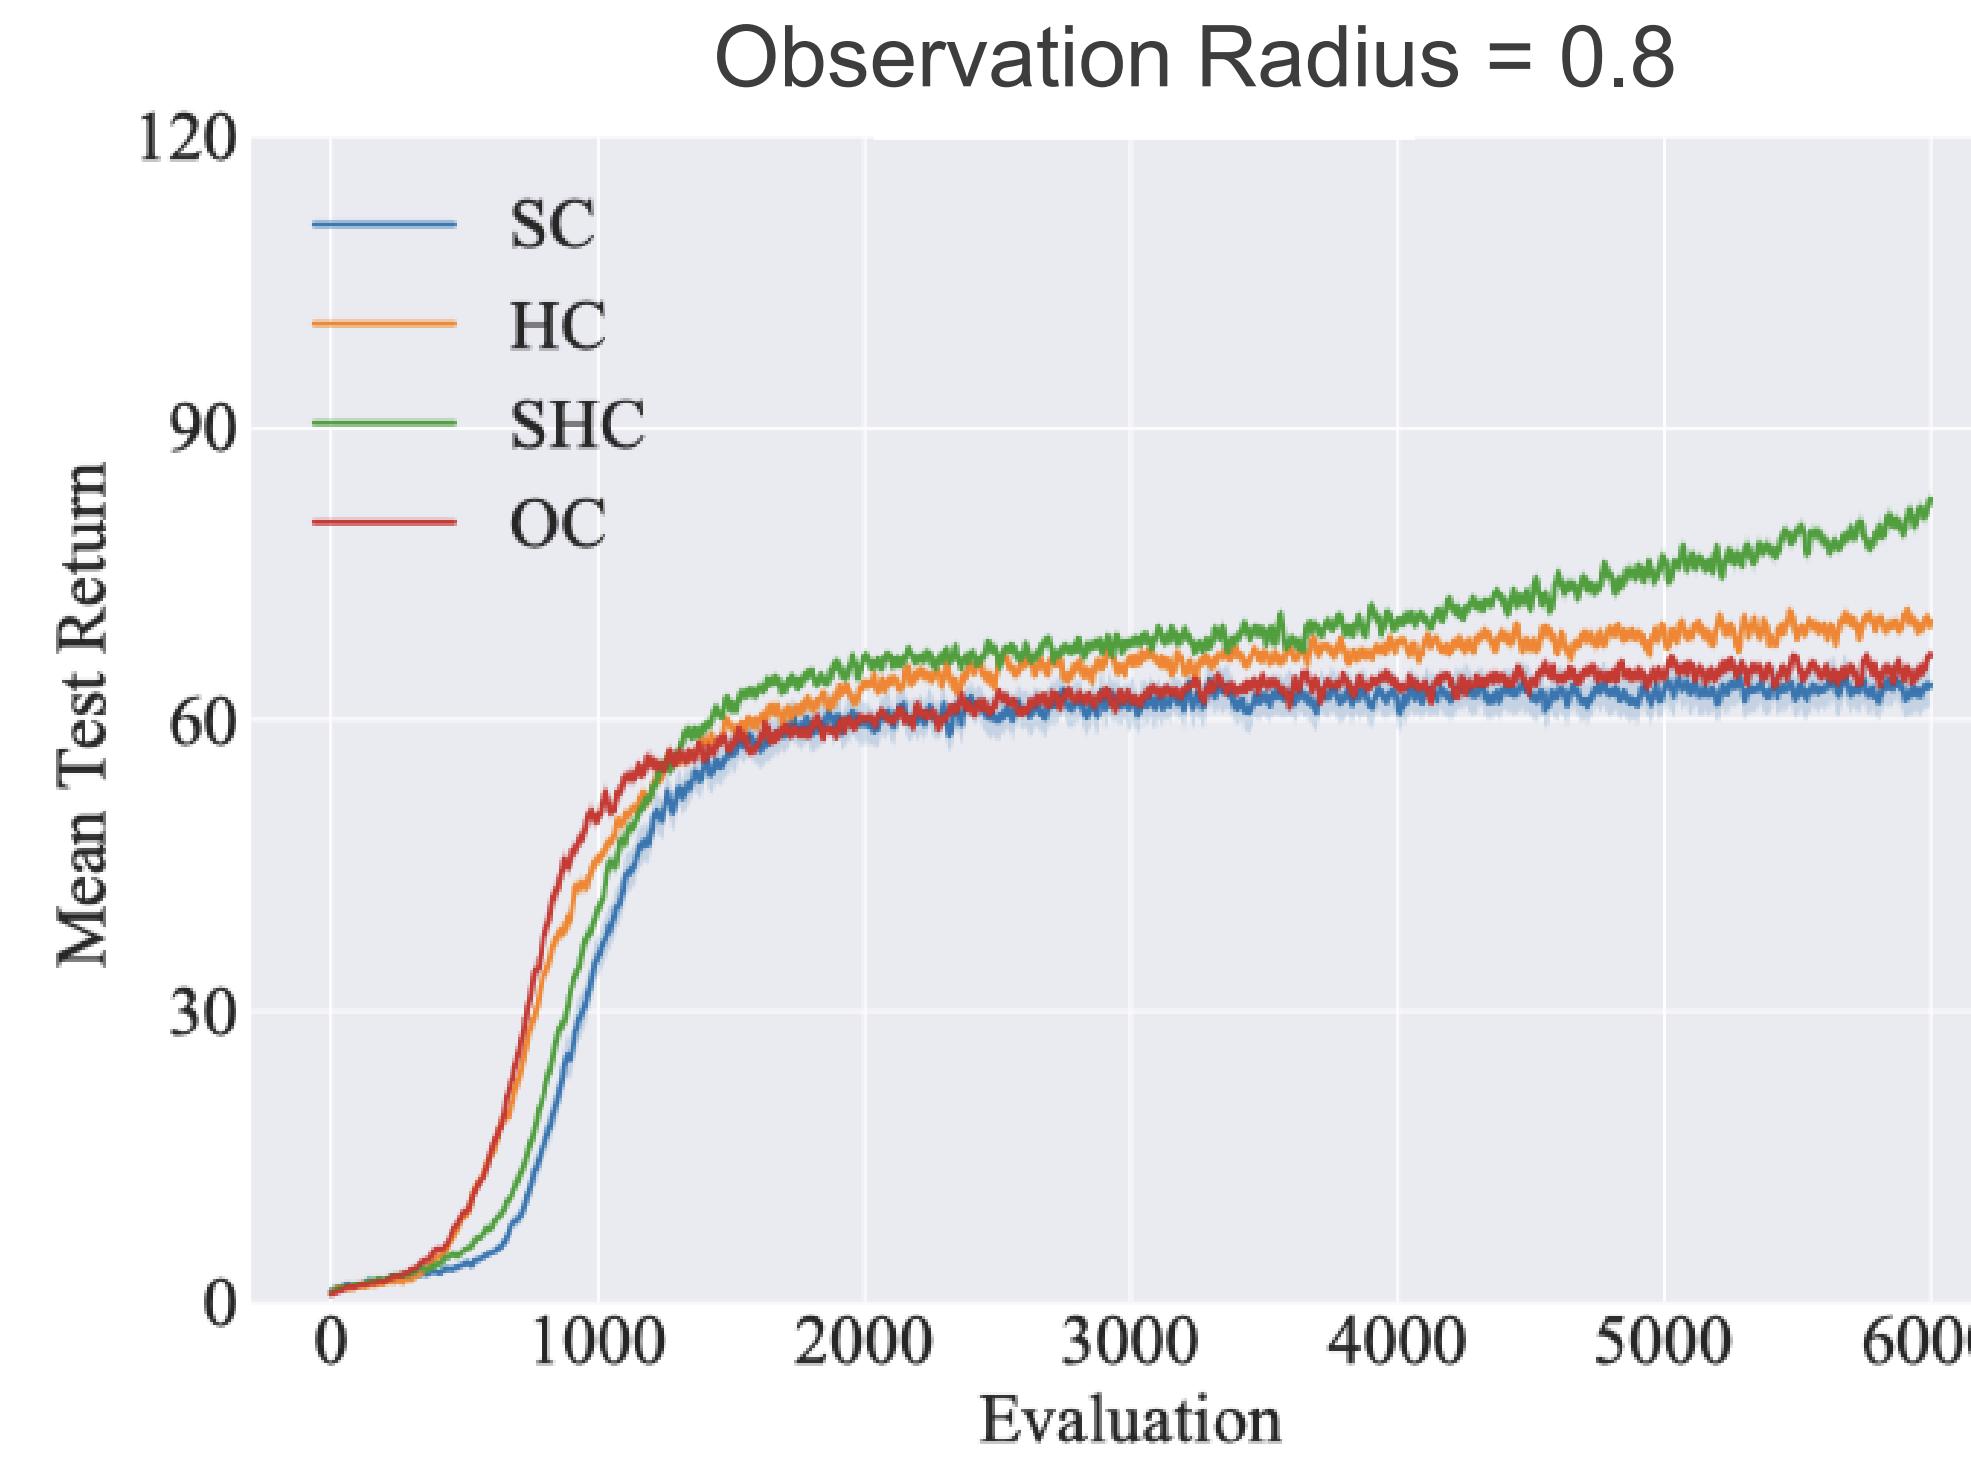


SMAC - StarCraft Multi-Agent Challenge



Partially Observable Particle Environments

Predator and Prey



Takeaways

Benchmark problems

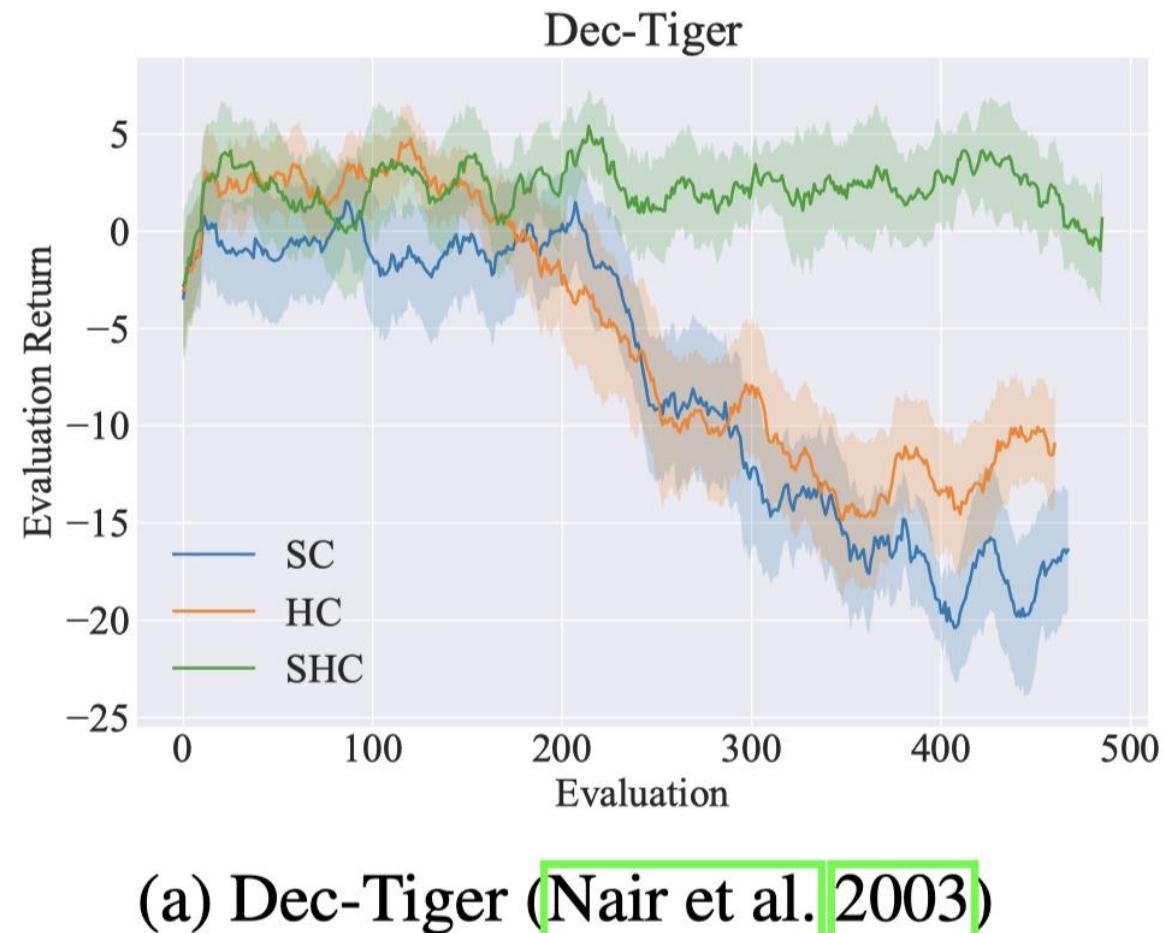
- We need harder, more partially observable problems

Methods to use

- Decentralized critics and (centralized) state-history-based often work the best
- MAPPO paper had a similar result
- Not really clear why

CTDE

- What is the best way to perform centralized training for decentralized execution (that's both principled and performs well)?



Other CTDE methods

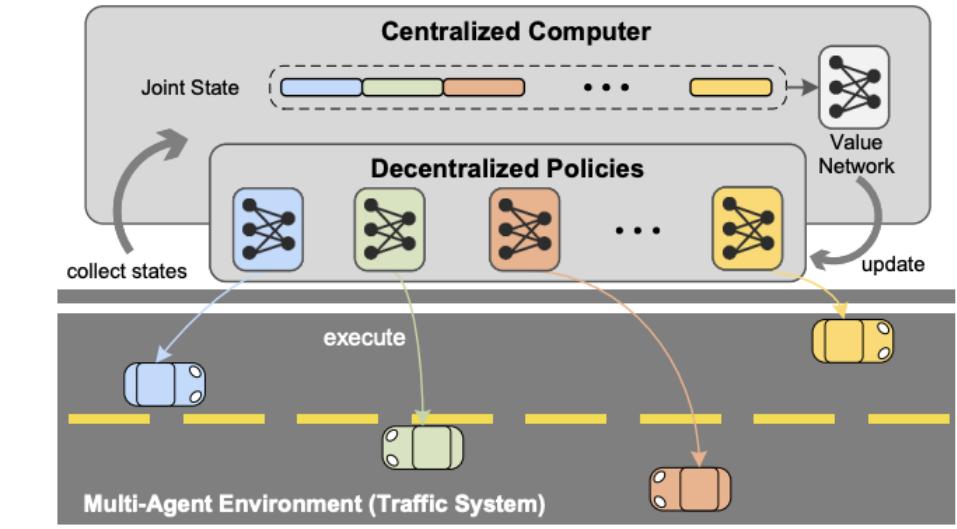
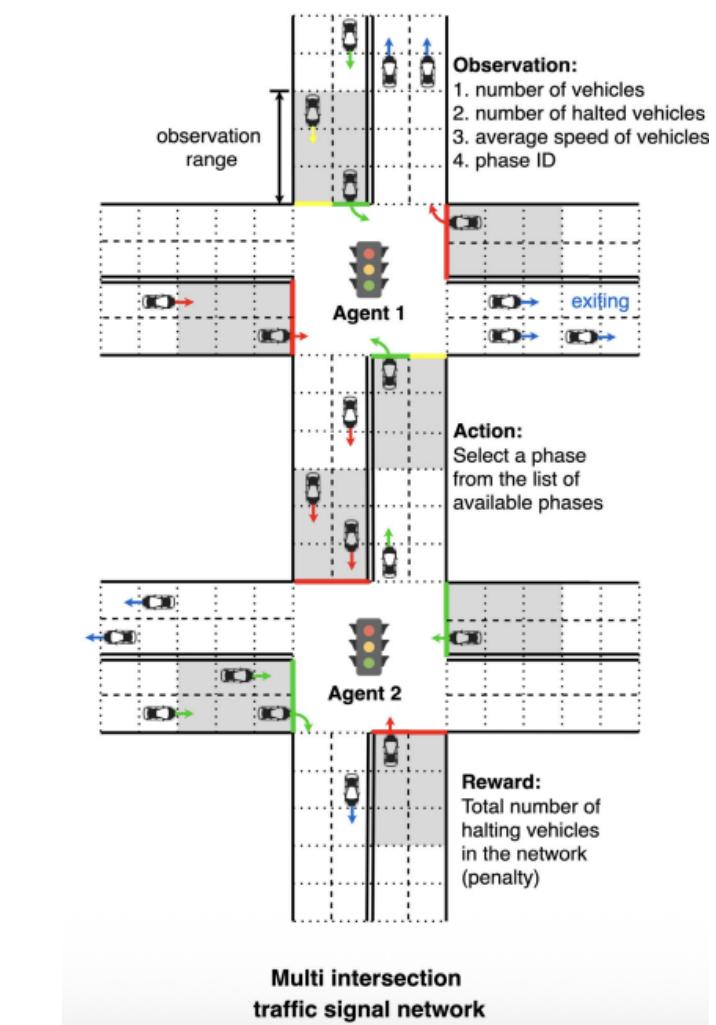
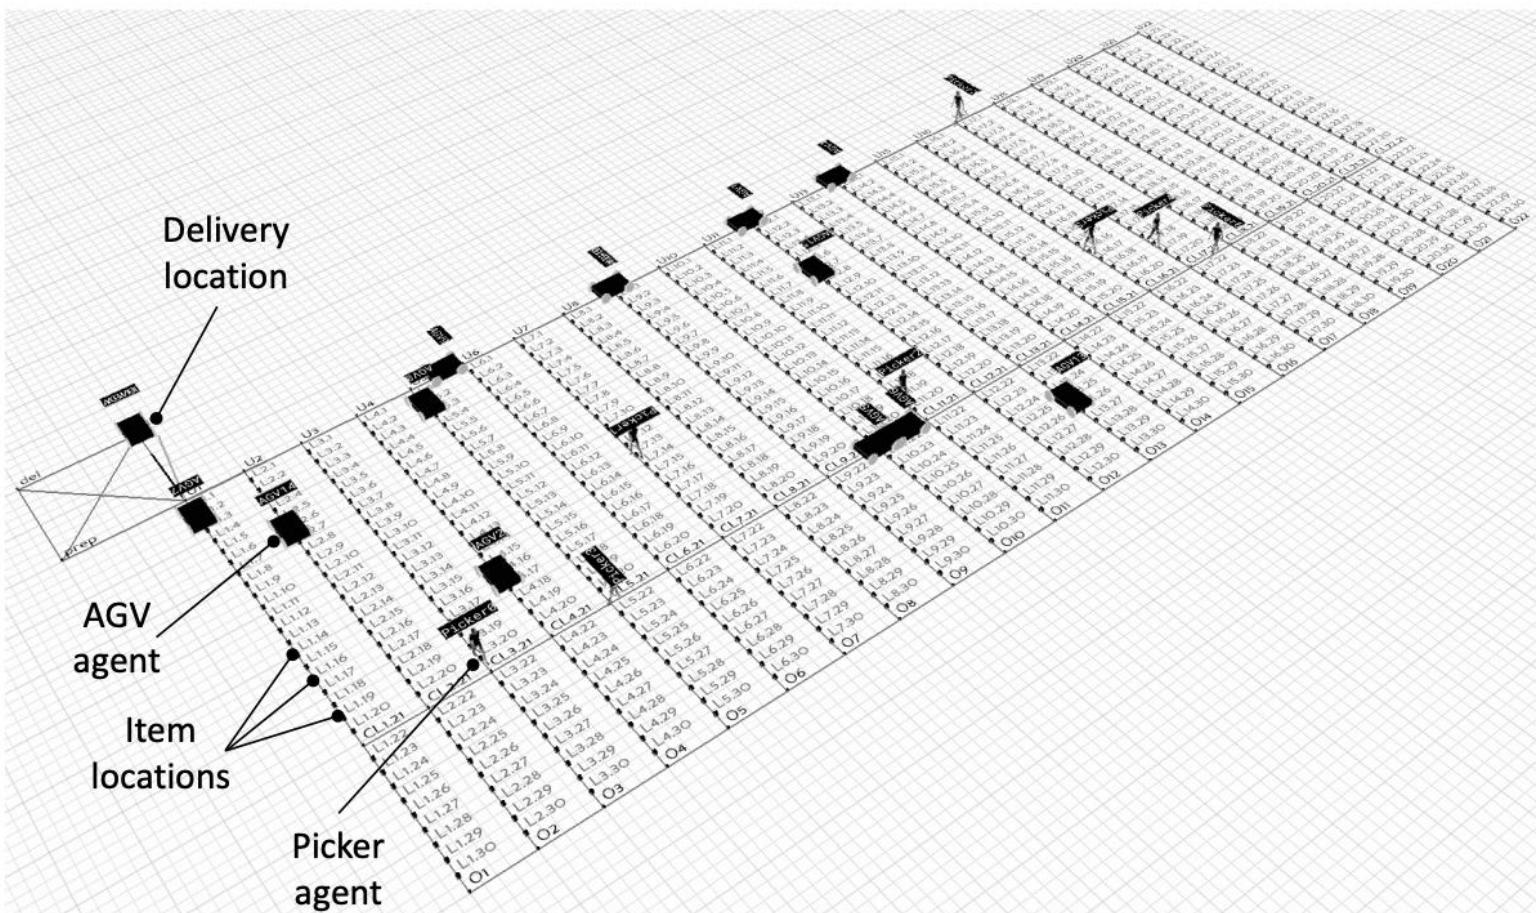
- Many other extensions and approaches:
 - E.g., FACMAC: Use a factored critic (doesn't need IGM) (Peng et al., 2021)
- **Parameter Sharing**
- **Alternating learning**
 - (Banerjee et al., 2012, Su et al., 2024)
 - Sequential agent updates as in HATRPO and HAPPO (Kuba et al. 2022)
- **Other agent modeling**, e.g., LOLA (Foerster et al. 2018a)

Other topics

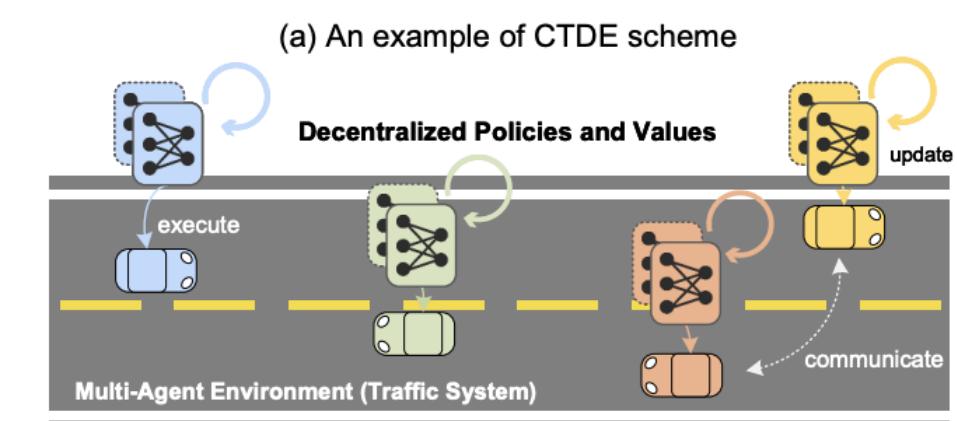
Many other topics in (cooperative) MARL that we don't have time to cover

- Communication ([Zhu et al., 2024](#))
- Ad hoc teamwork ([Mirsky et al., 2022](#)),
- Model-based methods ([Wang et al., 2022](#))
- Exploration, offline methods, model-based methods, hierarchical methods, role decomposition, multi-task approaches, etc.

Applications



(Bokade et al., 2023)



(b) An example of DTDE scheme

- Video games (e.g.,AlphaStar (Vinyals et al., 2019)
 - Centralized MARL for a team
 - Warehouse robots (Krnjaic et al. 2024)
 - Hierarchical CTDE approach

- Traffic signal control (e.g., survey by Wei et al. 2021)
 - Autonomous vehicle control (e.g., survey by Zhang et al. 2024)
 - Power systems, etc!

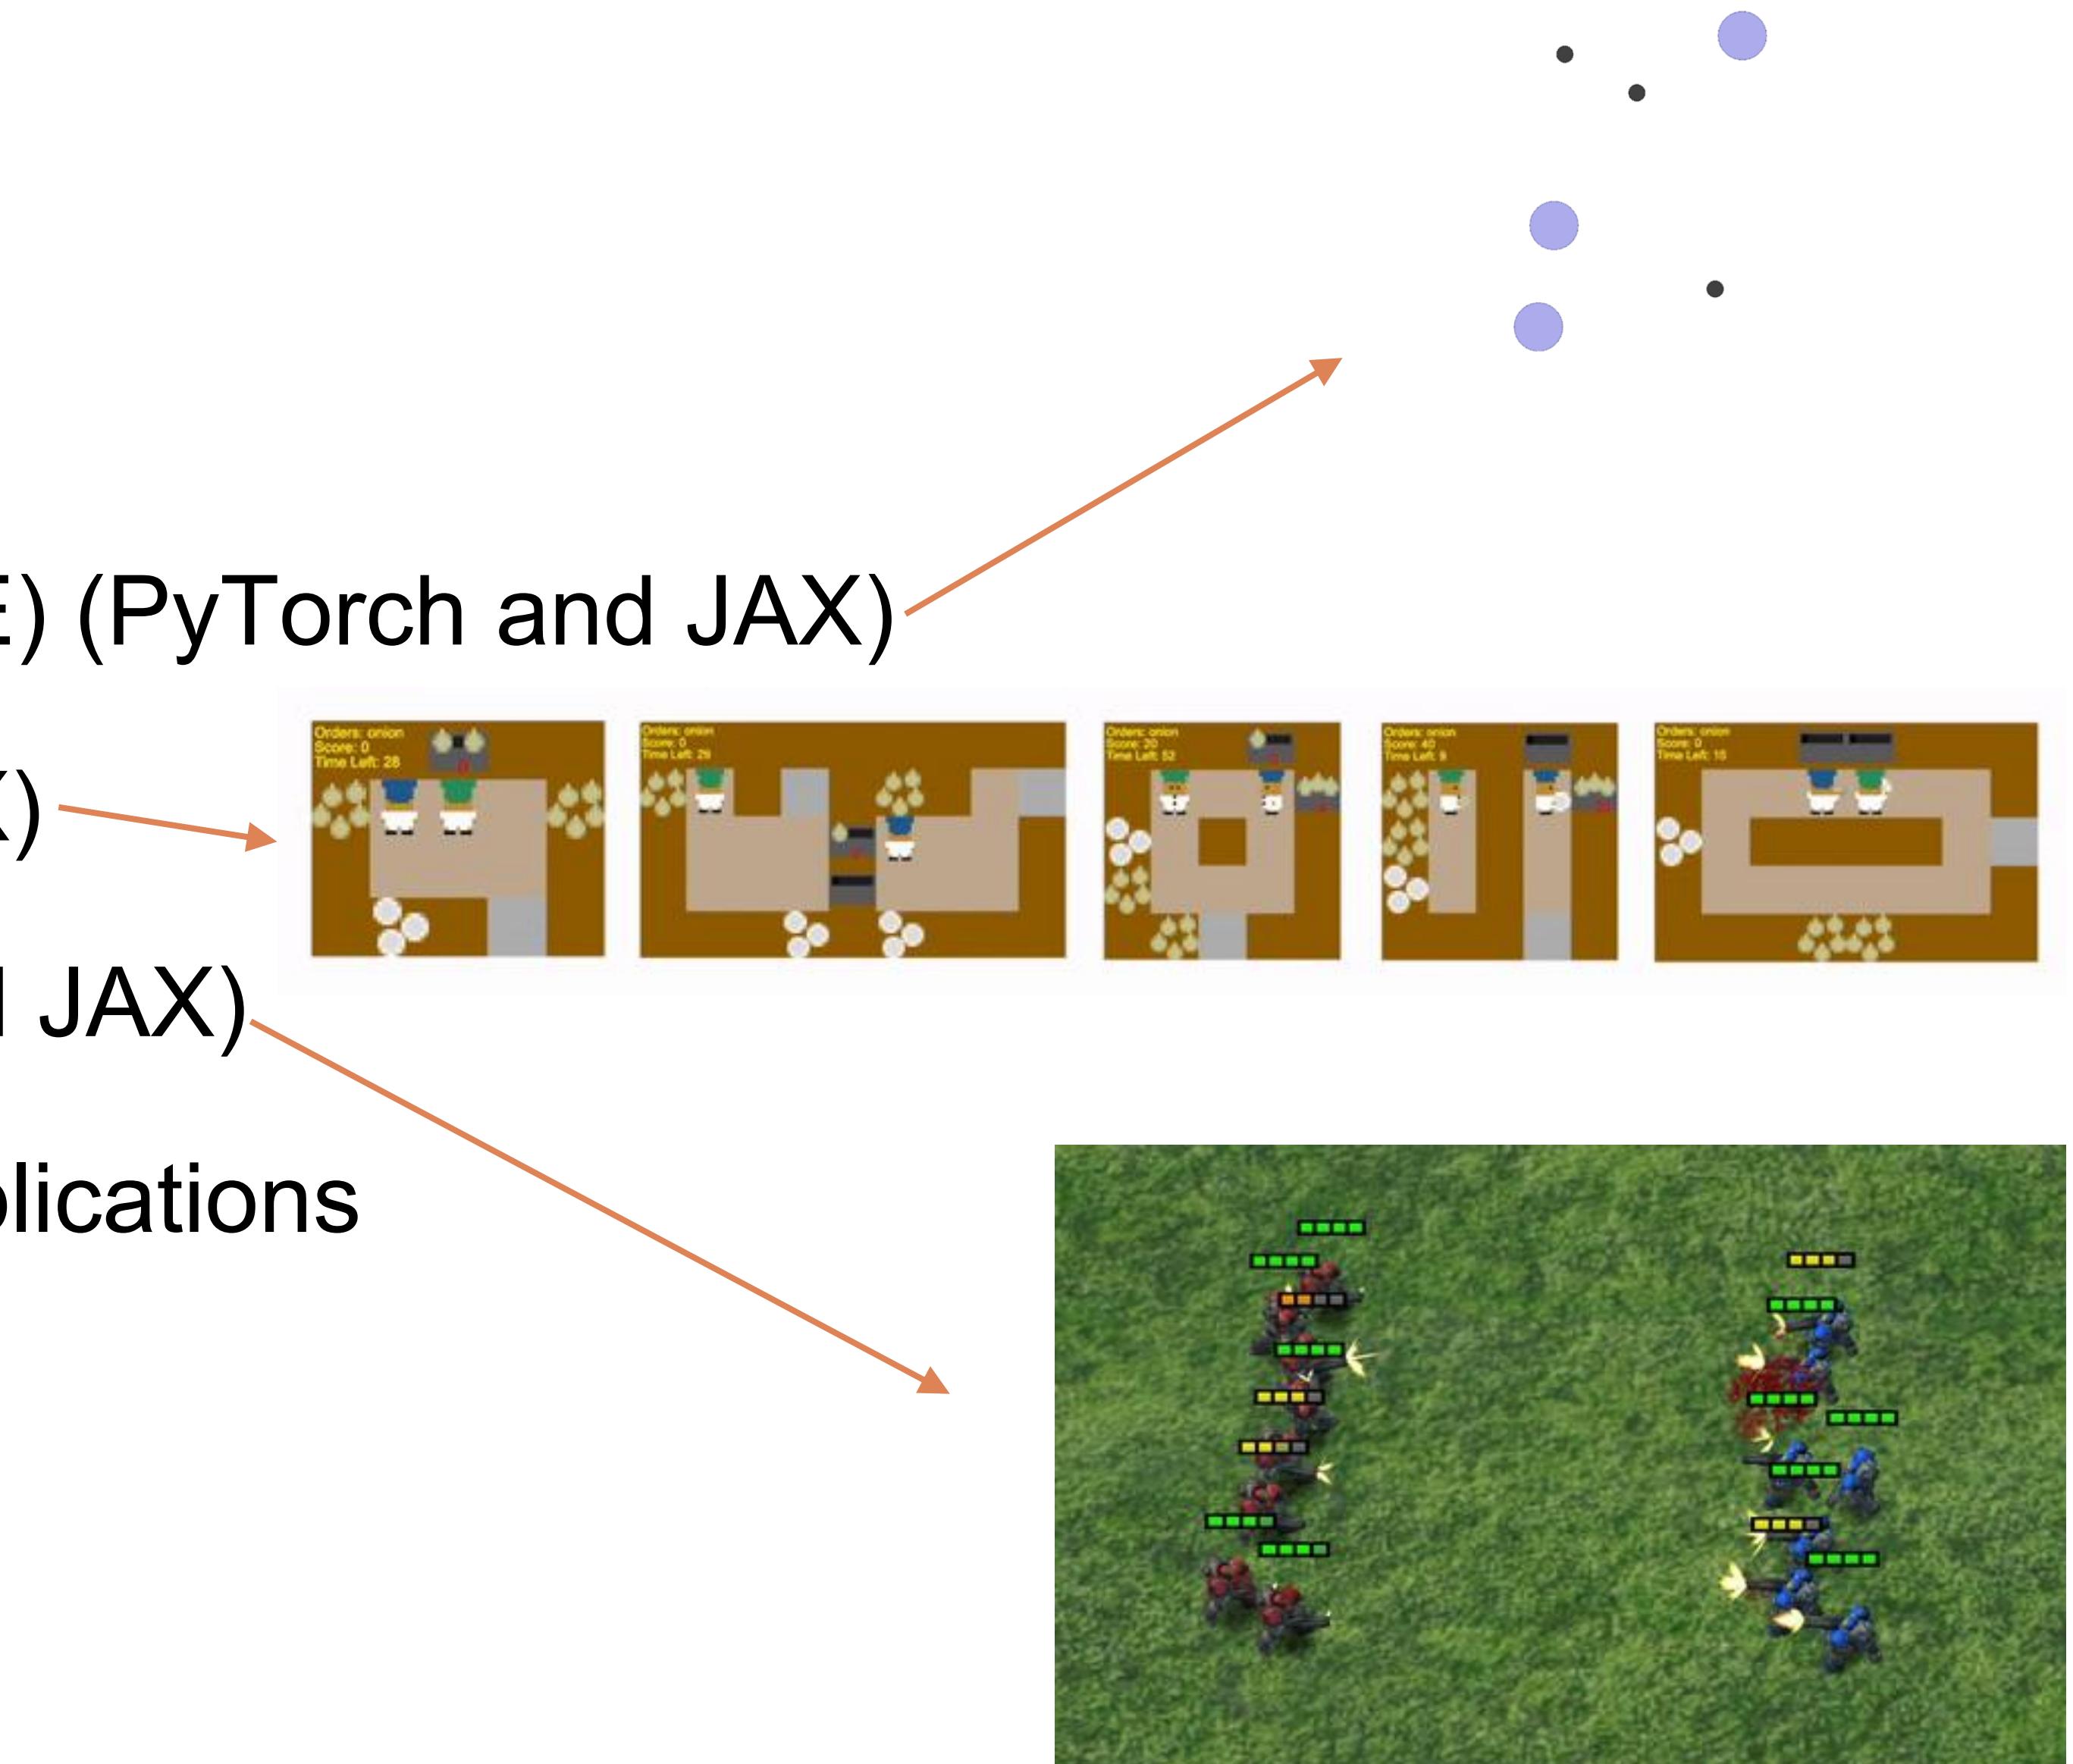
Multi-agent RL with macro-actions

[Xiao, Hoffman, Xia and Amato – ICRA20](#)



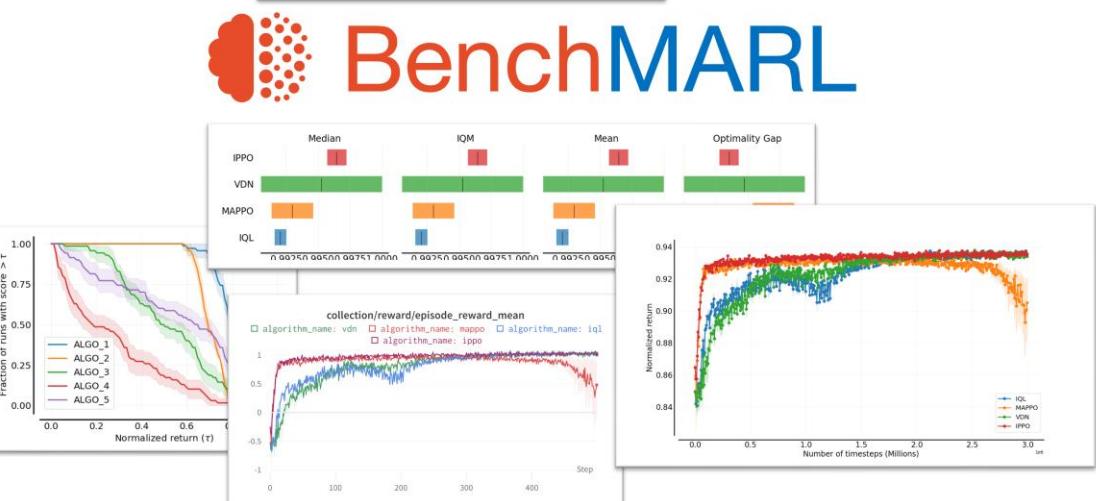
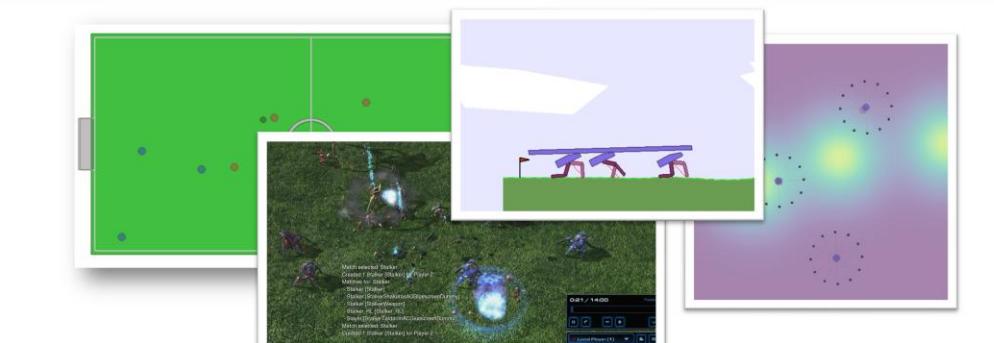
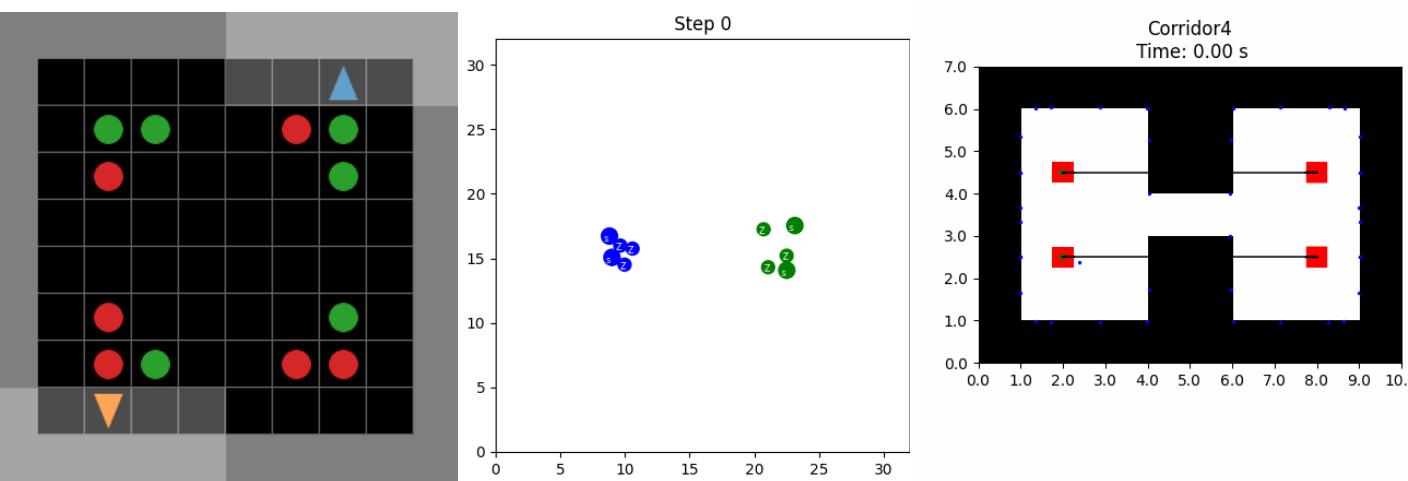
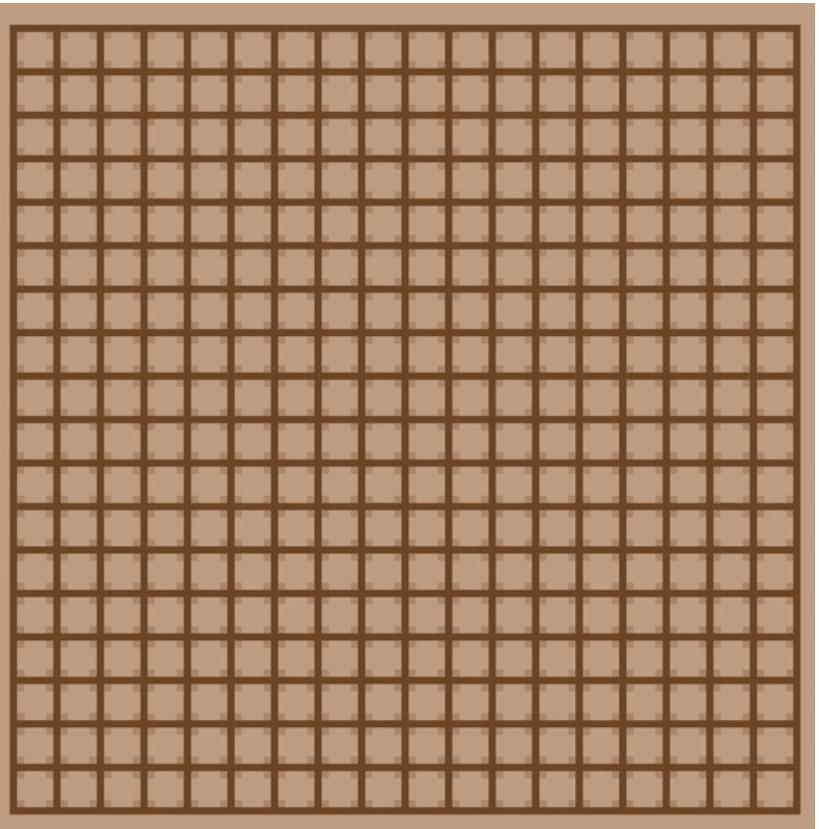
Benchmarks

- Standard domains:
 - Multi-agent Particle Envs (MPE) (PyTorch and JAX)
 - Overcooked (PyTorch and JAX)
 - SMAC v1 and v2 (PyTorch and JAX)
- Many many more inspired by applications



Environments and code

- PettingZoo
 - Multi-agent version of gym
 - Interface and some environments
 - <https://pettingzoo.farama.org/>
- JAXMARL
 - Efficient (JAX-based) baseline methods and environments
 - <https://github.com/FLAIROx/JaxMARL/tree/main/jaxmarl/environments/smax>
- BenchMARL
 - PyTorch baseline methods and environments
 - <https://github.com/facebookresearch/BenchMARL>
- Several more...

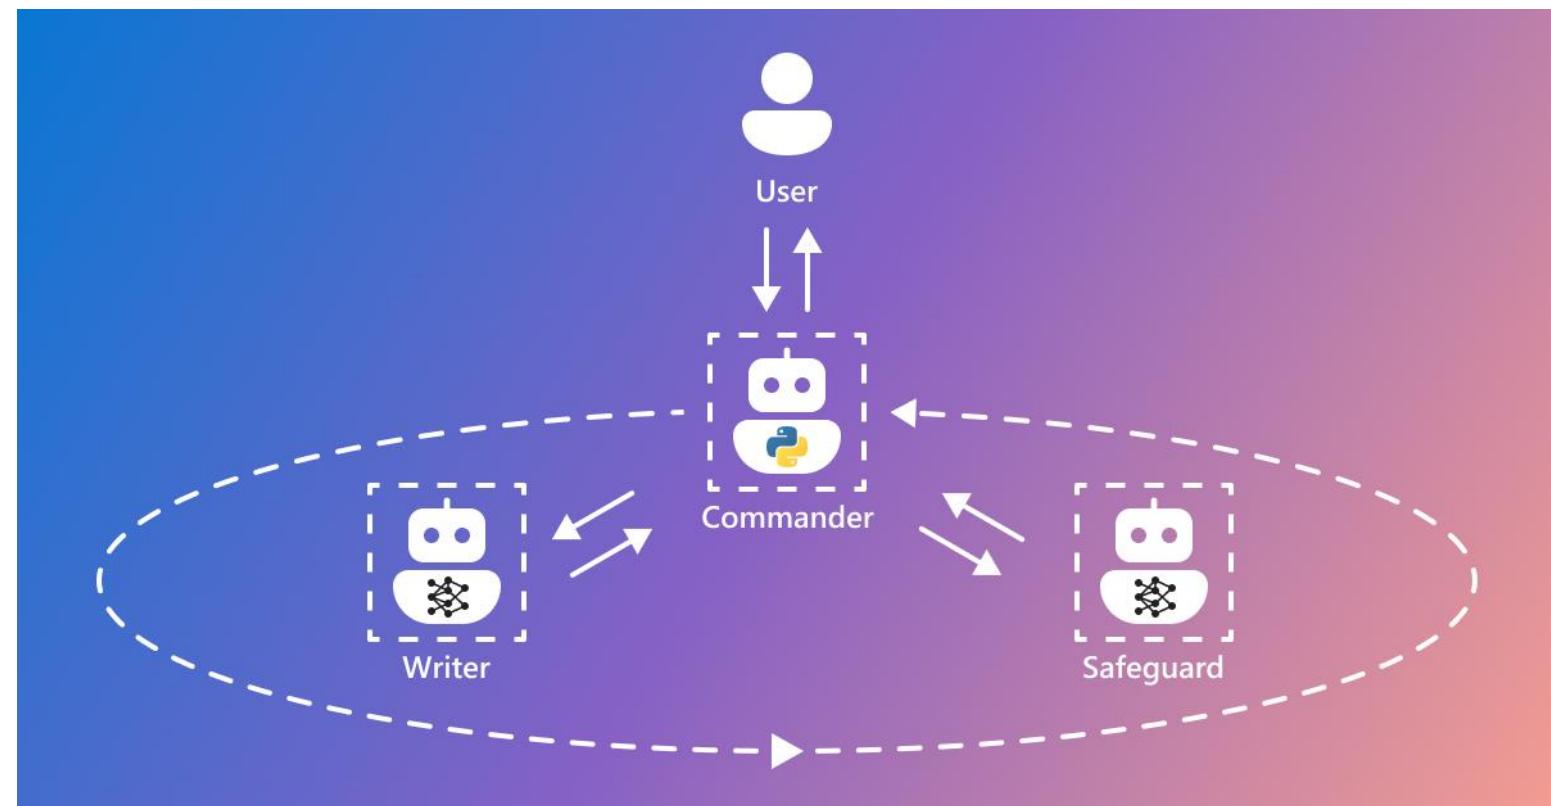


MARL and LLMs

- RL is widely used for LLMs
- MARL is *not* currently used for multi-agent LLMs (to best of my knowledge)
- There is no reason it couldn't be
- Open questions
 - Use cases
 - Control scheme
 - MARLHF
 - Training
- Benefits: specialization, robustness, scalability/performance
- Disconnect between academia and industry



<https://developer.nvidia.com/blog/introduction-to-lm-agents/>



<https://www.microsoft.com/en-us/research/project/autogen/>

Conclusion

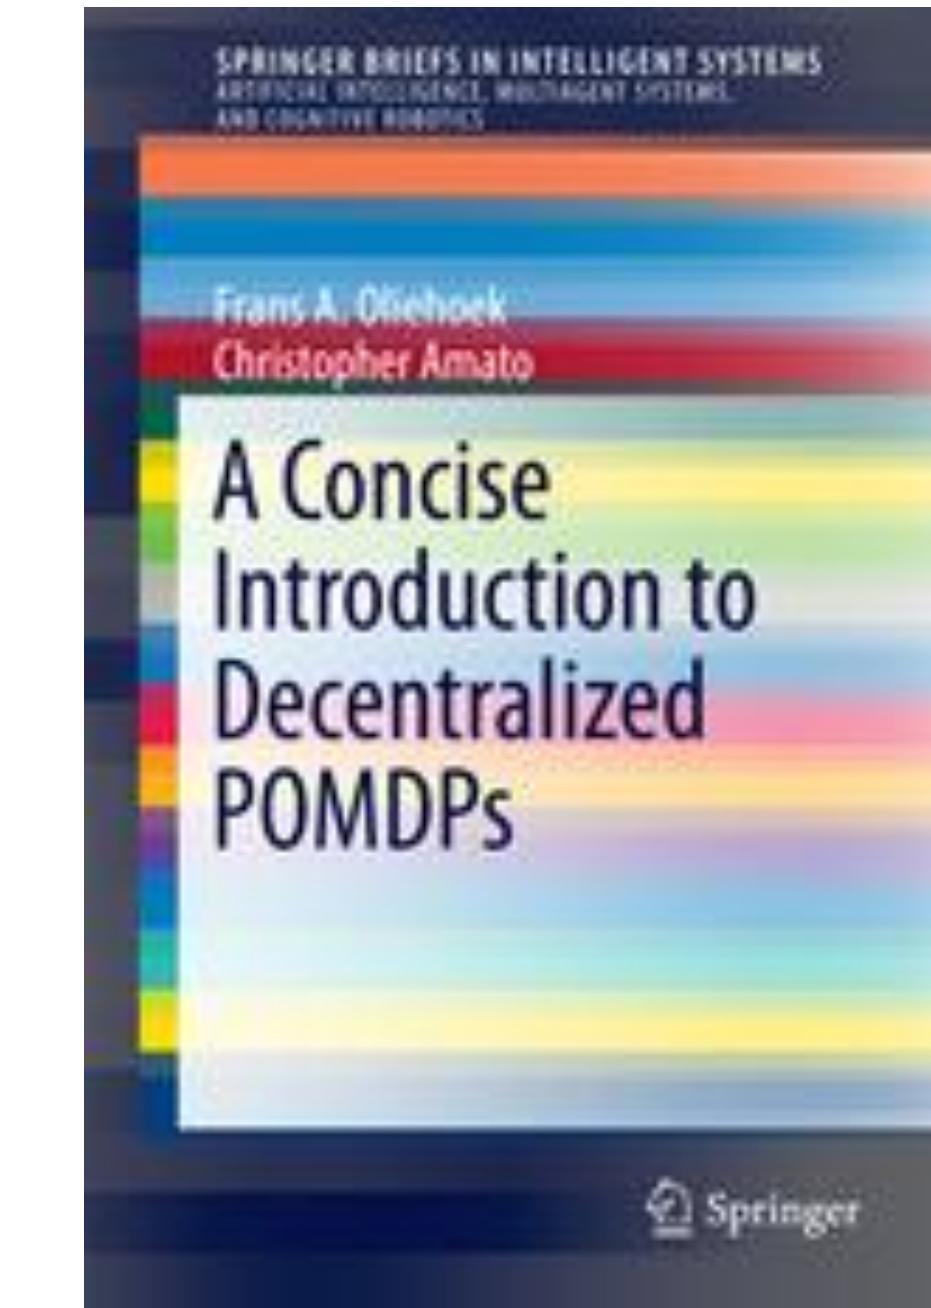
- Cooperative multi-agent reinforcement learning is a very general setting that fits with lots of applications
- A lot of work cooperative MARL
 - Centralized training and execution
 - Decentralized training and execution
 - Centralized training for decentralized execution (CTDE)
- Academia and industry are working on improved methods to improve scalability and performance

Conclusion

- Many open questions
 - MARL for LLM agents
 - Very scalable MARL
 - Optimal MARL
 - How to best do CTDE
 - Multiagent approaches to ML (e.g., GANs, decentralized methods)

Our resources

- Dec-POMDP book
 - Background on models and planning methods
- Book draft (An Initial Introduction to Cooperative Multi-Agent Reinforcement Learning):
<https://arxiv.org/abs/2405.06161>



- Let us know what you think and what should be changed/added for the final version!
- Slides will be available
 - <https://www.khoury.northeastern.edu/home/camato/tutorials.html>

Contents

	CONTENTS
1	Introduction
1.1	Overview
1.2	The cooperative MARL problem: The Dec-POMDP
1.3	Background on (single-agent) reinforcement learning
1.3.1	Value-based methods
1.3.2	Policy gradient methods
2	Centralized training and execution (CTE)
2.1	CTE overview
2.2	Centralized models
2.3	Centralized solutions
2.4	Improving scalability
3	Decentralized training and execution (DTE)
3.1	DTE overview
3.2	Decentralized, value-based methods
3.2.1	Independent Q-learning (IQL)
3.2.2	Improving the performance of IQL
3.2.3	Deep extensions, issues, and fixes
3.3	Decentralized policy gradient methods
3.3.1	Decentralized REINFORCE
3.3.2	Independent actor critic (IAC)
3.3.3	Other decentralized policy gradient methods
3.4	Other topics
4	Centralized training for decentralized execution (CTDE)
4.1	CTDE overview
4.2	Value function factorization methods
4.2.1	VDN
4.2.2	QMIX
4.2.3	Weighted QMIX
4.2.4	QTRAN
4.2.5	QPLEX
4.2.6	The use of state in factorization methods
4.3	Centralized critic methods
4.3.1	Preliminaries